# Chapter 4. The `class` File Format

This chapter describes the Java Virtual Machine `class` file format. Each `class` file contains the definition of a single class or interface. Although a class or interface need not have an external representation literally contained in a file (for instance, because the class is generated by a class loader), we will colloquially refer to any valid representation of a class or interface as being in the `class` file format.

A `class` file consists of a stream of 8-bit bytes. All 16-bit, 32-bit, and 64-bit quantities are constructed by reading in two, four, and eight consecutive 8-bit bytes, respectively. Multibyte data items are always stored in big-endian order, where the high bytes come first. In the Java SE platform, this format is supported by interfaces `java.io.DataInput` and `java.io.DataOutput` and classes such as `java.io.DataInputStream` and `java.io.DataOutputStream`.

This chapter defines its own set of data types representing `class` file data: The types u1, u2, and u4 represent an unsigned one-, two-, or four-byte quantity, respectively. In the Java SE platform, these types may be read by methods such as `readUnsignedByte`, `readUnsignedShort`, and `readInt` of the interface `java.io.DataInput`.

This chapter presents the `class` file format using pseudostructures written in a C-like structure notation. To avoid confusion with the fields of classes and class instances, etc., the contents of the structures describing the `class` file format are referred to as *items*. Successive items are stored in the `class` file sequentially, without padding or alignment.

*Tables*, consisting of zero or more variable-sized items, are used in several `class` file structures. Although we use C-like array syntax to refer to table items, the fact that tables are streams of varying-sized structures means that it is not possible to translate a table index directly to a byte offset into the table.

Where we refer to a data structure as an array, it consists of zero or more contiguous fixed-sized items and can be indexed like an array.

Reference to an ASCII character in this chapter should be interpreted to mean the Unicode code point corresponding to the ASCII character.

## 4.1. The `ClassFile` Structure

A `class` file consists of a single `ClassFile` structure:

```
ClassFile {
    u4              magic;
    u2              minor_version;
    u2              major_version;
    u2              constant_pool_count;
    cp_info         constant_pool[constant_pool_count-1];
    u2              access_flags;
    u2              this_class;
    u2              super_class;
    u2              interfaces_count;
    u2              interfaces[interfaces_count];
    u2              fields_count;
    field_info      fields[fields_count];
    u2              methods_count;
    method_info     methods[methods_count];
    u2              attributes_count;
    attribute_info  attributes[attributes_count];
}
```

The items in the `ClassFile` structure are as follows:

magic

The `magic` item supplies the magic number identifying the `class` file format; it has the value 0xCAFEBABE.

minor_version, major_version

> The values of the `minor_version` and `major_version` items are the minor and major version numbers of this `class` file. Together, a major and a minor version number determine the version of the `class` file format. If a `class` file has major version number M and minor version number m, we denote the version of its `class` file format as M.m. Thus, `class` file format versions may be ordered lexicographically, for example, 1.5 < 2.0 < 2.1.

> A Java Virtual Machine implementation can support a `class` file format of version v if and only if v lies in some contiguous range Mi.0 ≤ v ≤ Mj.m. The release level of the Java SE platform to which a Java Virtual Machine implementation conforms is responsible for determining the range.

> *Oracle's Java Virtual Machine implementation in JDK release 1.0.2 supports `class` file format versions 45.0 through 45.3 inclusive. JDK releases 1.1.\* support `class` file format versions in the range 45.0 through 45.65535 inclusive. For k ≥ 2, JDK release 1.k supports `class` file format versions in the range 45.0 through 44+k.0 inclusive.*

constant_pool_count

> The value of the `constant_pool_count` item is equal to the number of entries in the `constant_pool` table plus one. A `constant_pool` index is considered valid if it is greater than zero and less than `constant_pool_count`, with the exception for constants of type `long` and `double` noted in §4.4.5.

constant_pool[]

> The `constant_pool` is a table of structures (§4.4) representing various string constants, class and interface names, field names, and other constants that are referred to within the `ClassFile` structure and its substructures. The format of each `constant_pool` table entry is indicated by its first "tag" byte.

> The `constant_pool` table is indexed from 1 to `constant_pool_count`-1.

access_flags

> The value of the `access_flags` item is a mask of flags used to denote access permissions to and properties of this class or interface. The interpretation of each flag, when set, is as shown in Table 4.1.

**Table 4.1. Class access and property modifiers**

| Flag Name | Value | Interpretation |
|---|---|---|
| ACC_PUBLIC | 0x0001 | Declared `public`; may be accessed from outside its package. |
| ACC_FINAL | 0x0010 | Declared `final`; no subclasses allowed. |
| ACC_SUPER | 0x0020 | Treat superclass methods specially when invoked by the *invokespecial* instruction. |
| ACC_INTERFACE | 0x0200 | Is an interface, not a class. |
| ACC_ABSTRACT | 0x0400 | Declared `abstract`; must not be instantiated. |
| ACC_SYNTHETIC | 0x1000 | Declared synthetic; not present in the source code. |
| ACC_ANNOTATION | 0x2000 | Declared as an annotation type. |
| ACC_ENUM | 0x4000 | Declared as an `enum` type. |

A class may be marked with the ACC_SYNTHETIC flag to indicate that it was generated by a compiler and does not appear in source code.

The ACC_ENUM flag indicates that this class or its superclass is declared as an enumerated type.

An interface is distinguished by its ACC_INTERFACE flag being set. If its ACC_INTERFACE flag is not set, this `class` file defines a class, not an interface.

If the ACC_INTERFACE flag of this `class` file is set, its ACC_ABSTRACT flag must also be set (JLS §9.1.1.1). Such a `class` file must not have its ACC_FINAL, ACC_SUPER or ACC_ENUM flags set.

An annotation type must have its ACC_ANNOTATION flag set. If the ACC_ANNOTATION flag is set, the ACC_INTERFACE flag must be set as well. If the ACC_INTERFACE flag of this `class` file is not set, it may have any of the other flags in *Table 4.1* set, except the ACC_ANNOTATION flag. However, such a `class` file cannot have both its ACC_FINAL and ACC_ABSTRACT flags set (JLS §8.1.1.2).

The ACC_SUPER flag indicates which of two alternative semantics is to be expressed by

The ACC_SUPER flag indicates which of two alternative semantics is to be expressed by the *invokespecial* instruction (*§invokespecial*) if it appears in this class. Compilers to the instruction set of the Java Virtual Machine should set the ACC_SUPER flag.

*The ACC_SUPER flag exists for backward compatibility with code compiled by older compilers for the Java programming language. In Oracle's JDK prior to release 1.0.2, the compiler generated ClassFile access_flags in which the flag now representing ACC_SUPER had no assigned meaning, and Oracle's Java Virtual Machine implementation ignored the flag if it was set.*

All bits of the `access_flags` item not assigned in Table 4.1 are reserved for future use. They should be set to zero in generated `class` files and should be ignored by Java Virtual Machine implementations.

this_class

The value of the `this_class` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a CONSTANT_Class_info structure (§4.4.1) representing the class or interface defined by this `class` file.

super_class

For a class, the value of the `super_class` item either must be zero or must be a valid index into the `constant_pool` table. If the value of the `super_class` item is nonzero, the `constant_pool` entry at that index must be a CONSTANT_Class_info structure (§4.4.1) representing the direct superclass of the class defined by this `class` file. Neither the direct superclass nor any of its superclasses may have the ACC_FINAL flag set in the `access_flags` item of its ClassFile structure.

If the value of the `super_class` item is zero, then this `class` file must represent the class Object, the only class or interface without a direct superclass.

For an interface, the value of the `super_class` item must always be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a CONSTANT_Class_info structure representing the class Object.

interfaces_count

The value of the `interfaces_count` item gives the number of direct superinterfaces of this class or interface type.

interfaces[]

Each value in the `interfaces` array must be a valid index into the `constant_pool` table. The `constant_pool` entry at each value of `interfaces[`$i$`]`, where $0 \leq i <$ `interfaces_count`, must be a `CONSTANT_Class_info` structure ([§4.4.1](#)) representing an interface that is a direct superinterface of this class or interface type, in the left-to-right order given in the source for the type.

fields_count

The value of the `fields_count` item gives the number of `field_info` structures in the `fields` table. The `field_info` structures ([§4.5](#)) represent all fields, both class variables and instance variables, declared by this class or interface type.

fields[]

Each value in the `fields` table must be a `field_info` ([§4.5](#)) structure giving a complete description of a field in this class or interface. The `fields` table includes only those fields that are declared by this class or interface. It does not include items representing fields that are inherited from superclasses or superinterfaces.

methods_count

The value of the `methods_count` item gives the number of `method_info` structures in the `methods` table.

methods[]

Each value in the `methods` table must be a `method_info` ([§4.6](#)) structure giving a complete description of a method in this class or interface. If neither of the `ACC_NATIVE` and `ACC_ABSTRACT` flags are set in the `access_flags` item of a `method_info` structure, the Java Virtual Machine instructions implementing the method are also supplied.

The `method_info` structures represent all methods declared by this class or interface type, including instance methods, class methods, instance initialization methods ([§2.9](#)), and any class or interface initialization method ([§2.9](#)). The `methods` table does not include items representing methods that are inherited from superclasses or superinterfaces.

attributes_count

> The value of the `attributes_count` item gives the number of attributes ([§4.7](#)) in the `attributes` table of this class.

attributes[]

> Each value of the `attributes` table must be an `attribute_info` ([§4.7](#)) structure.
>
> The attributes defined by this specification as appearing in the `attributes` table of a `ClassFile` structure are the `InnerClasses` ([§4.7.6](#)), `EnclosingMethod` ([§4.7.7](#)), `Synthetic` ([§4.7.8](#)), `Signature` ([§4.7.9](#)), `SourceFile` ([§4.7.10](#)), `SourceDebugExtension` ([§4.7.11](#)), `Deprecated` ([§4.7.15](#)), `RuntimeVisibleAnnotations` ([§4.7.16](#)), `RuntimeInvisibleAnnotations` ([§4.7.17](#)), and `BootstrapMethods` ([§4.7.21](#)) attributes.
>
> If a Java Virtual Machine implementation recognizes `class` files whose version number is 49.0 or above, it must recognize and correctly read `Signature` ([§4.7.9](#)), `RuntimeVisibleAnnotations` ([§4.7.16](#)), and `RuntimeInvisibleAnnotations` ([§4.7.17](#)) attributes found in the `attributes` table of a `ClassFile` structure of a `class` file whose version number is 49.0 or above.
>
> If a Java Virtual Machine implementation recognizes `class` files whose version number is 51.0 or above, it must recognize and correctly read `BootstrapMethods` ([§4.7.21](#)) attributes found in the `attributes` table of a `ClassFile` structure of a `class` file whose version number is 51.0 or above.
>
> A Java Virtual Machine implementation is required to silently ignore any or all attributes in the `attributes` table of a `ClassFile` structure that it does not recognize. Attributes not defined in this specification are not allowed to affect the semantics of the `class` file, but only to provide additional descriptive information ([§4.7.1](#)).

# 4.2. The Internal Form of Names

### 4.2.1. Binary Class and Interface Names

Class and interface names that appear in `class` file structures are always represented in a fully qualified form known as *binary names* (JLS §13.1). Such names are always represented as `CONSTANT_Utf8_info` structures ([§4.4.7](#)) and thus may be drawn, where not further constrained, from the entire Unicode codespace. Class and interface names are referenced from those `CONSTANT_NameAndType_info` structures ([§4.4.6](#)) which have such names as part of their descriptor ([§4.3](#)), and from all `CONSTANT_Class_info` structures ([§4.4.1](#)).

For historical reasons, the syntax of binary names that appear in `class` file structures differs from the syntax of binary names documented in JLS §13.1. In this internal form, the ASCII periods (`.`) that normally separate the identifiers which make up the binary name are replaced by ASCII forward slashes (`/`). The identifiers themselves must be unqualified names ([§4.2.2](#)).

*For example, the normal binary name of class `Thread` is `java.lang.Thread`. In the internal form used in descriptors in the `class` file format, a reference to the name of class `Thread` is implemented using a `CONSTANT_Utf8_info` structure representing the string `java/lang/Thread`.*

### 4.2.2. Unqualified Names

Names of methods, fields, and local variables are stored as *unqualified names*. An unqualified name must not contain any of the ASCII characters `.` `;` `[` `/` (that is, period or semicolon or left square bracket or forward slash).

Method names are further constrained so that, with the exception of the special method names `<init>` and `<clinit>` ([§2.9](#)), they must not contain the ASCII characters `<` or `>` (that is, left angle bracket or right angle bracket).

*Note that a field name or interface method name may be `<init>` or `<clinit>`, but no method invocation instruction may reference `<clinit>` and only the invokespecial instruction ([§invokespecial](#)) may reference `<init>`.*

## 4.3. Descriptors and Signatures

A *descriptor* is a string representing the type of a field or method. Descriptors are represented in the `class` file format using modified UTF-8 strings ([§4.4.7](#)) and thus may be drawn, where not further constrained, from the entire Unicode codespace.

A *signature* is a string representing the generic type of a field or method, or generic type

information for a class declaration.

### 4.3.1. Grammar Notation

Descriptors and signatures are specified using a grammar. This grammar is a set of productions that describe how sequences of characters can form syntactically correct descriptors of various types. Terminal symbols of the grammar are shown in `bold fixed-width` font. Nonterminal symbols are shown in *italic* type. The definition of a nonterminal is introduced by the name of the nonterminal being defined, followed by a colon. One or more alternative right-hand sides for the nonterminal then follow on succeeding lines. For example, the production:

```
FieldType:
    BaseType
    ObjectType
    ArrayType
```

states that a *FieldType* may represent either a *BaseType*, an *ObjectType* or an *ArrayType*.

A nonterminal symbol on the right-hand side of a production that is followed by an asterisk (*) represents zero or more possibly different values produced from that nonterminal, appended without any intervening space. Similarly, a nonterminal symbol on the right-hand side of a production that is followed by an plus sign (+) represents one or more possibly different values produced from that nonterminal, appended without any intervening space. The production:

```
MethodDescriptor:
    ( ParameterDescriptor* ) ReturnDescriptor
```

states that a *MethodDescriptor* represents a left parenthesis, followed by zero or more *ParameterDescriptor* values, followed by a right parenthesis, followed by a *ReturnDescriptor*.

### 4.3.2. Field Descriptors

A *field descriptor* represents the type of a class, instance, or local variable. It is a series of

characters generated by the grammar:

```
FieldDescriptor:
    FieldType

FieldType:
    BaseType
    ObjectType
    ArrayType

BaseType:
    B
    C
    D
    F
    I
    J
    S
    Z

ObjectType:
    L ClassName ;

ArrayType:
    [ ComponentType

ComponentType:
    FieldType
```

The characters of *BaseType*, the L and ; of *ObjectType*, and the [ of *ArrayType* are all ASCII characters.

The *ClassName* represents a binary class or interface name encoded in internal form (§4.2.1).

The interpretation of field descriptors as types is as shown in Table 4.2.

A field descriptor representing an array type is valid only if it represents a type with 255 or

fewer dimensions.

**Table 4.2. Interpretation of *FieldType* characters**

| *BaseType* Character | Type | Interpretation |
|---|---|---|
| B | `byte` | signed byte |
| C | `char` | Unicode character code point in the Basic Multilingual Plane, encoded with UTF-16 |
| D | `double` | double-precision floating-point value |
| F | `float` | single-precision floating-point value |
| I | `int` | integer |
| J | `long` | long integer |
| L *ClassName ;* | `reference` | an instance of class *ClassName* |
| S | `short` | signed short |
| Z | `boolean` | `true` or `false` |
| [ | `reference` | one array dimension |

*The field descriptor of an instance variable of type `int` is simply `I`.*

*The field descriptor of an instance variable of type `Object` is `Ljava/lang/Object;`.*
*Note that the internal form of the binary name for class `Object` is used.*

*The field descriptor of an instance variable that is a multidimensional double array,*
*`double d[][][]`, is `[[[D`.*

### 4.3.3. Method Descriptors

A *method descriptor* represents the parameters that the method takes and the value that it returns:

```
MethodDescriptor:
```

```
      ( ParameterDescriptor* ) ReturnDescriptor
```

A *parameter descriptor* represents a parameter passed to a method:

```
ParameterDescriptor:
    FieldType
```

A *return descriptor* represents the type of the value returned from a method. It is a series of characters generated by the grammar:

```
ReturnDescriptor:
    FieldType
    VoidDescriptor


VoidDescriptor:
    V
```

The character V indicates that the method returns no value (its return type is void).

A method descriptor is valid only if it represents method parameters with a total length of 255 or less, where that length includes the contribution for this in the case of instance or interface method invocations. The total length is calculated by summing the contributions of the individual parameters, where a parameter of type long or double contributes two units to the length and a parameter of any other type contributes one unit.

> *The method descriptor for the method:*
>
> ```
> Object m(int i, double d, Thread t) {..}
> ```
>
> *is (IDLjava/lang/Thread;)Ljava/lang/Object;. Note that the internal forms of the binary names of Thread and Object are used.*
>
> *The method descriptor for m is the same whether m is a class method or an instance method. Although an instance method is passed this, a reference to the current class instance, in addition to its intended parameters, that fact is not reflected in the method descriptor. The reference to this is passed implicitly by the method invocation*

### 4.3.4. Signatures

Signatures are used to encode Java programming language type information that is not part of the Java Virtual Machine type system, such as generic type and method declarations and parameterized types. See *The Java Language Specification, Java SE 7 Edition* for details about such types.

*This kind of type information is needed to support reflection and debugging, and by a Java compiler.*

In the following, the terminal symbol *Identifier* is used to denote the name of a type, field, local variable, parameter, method, or type variable, as generated by a Java compiler. Such a name must not contain any of the ASCII characters . ; [ / < > : (that is, the characters forbidden in method names (§4.2.2) and also colon) but may contain characters that must not appear in an identifier in the Java programming language (JLS §3.8).

A class signature, defined by the production *ClassSignature*, is used to encode type information about a class declaration. It describes any formal type parameters the class might have, and lists its (possibly parameterized) direct superclass and direct superinterfaces, if any.

```
ClassSignature:
    FormalTypeParameters_opt SuperclassSignature SuperinterfaceSignature*
```

A formal type parameter is described by its name, followed by its class and interface bounds. If the class bound does not specify a type, it is taken to be Object.

```
FormalTypeParameters:
    < FormalTypeParameter+ >

FormalTypeParameter:
    Identifier ClassBound InterfaceBound*

ClassBound:
```

```
        : FieldTypeSignatureopt

InterfaceBound:
    : FieldTypeSignature

SuperclassSignature:
    ClassTypeSignature

SuperinterfaceSignature:
    ClassTypeSignature
```

A field type signature, defined by the production *FieldTypeSignature*, encodes the (possibly parameterized) type for a field, parameter or local variable.

```
FieldTypeSignature:
    ClassTypeSignature
    ArrayTypeSignature
    TypeVariableSignature
```

A class type signature gives complete type information for a class or interface type. The class type signature must be formulated such that it can be reliably mapped to the binary name of the class it denotes by erasing any type arguments and converting each . character in the signature to a $ character.

```
ClassTypeSignature:
    L PackageSpecifier_opt SimpleClassTypeSignature ClassTypeSignatureSuffix* ;

PackageSpecifier:
    Identifier / PackageSpecifier*

SimpleClassTypeSignature:
    Identifier TypeArgumentsopt

ClassTypeSignatureSuffix:
    . SimpleClassTypeSignature

TypeVariableSignature:
```

```
        T Identifier ;

TypeArguments:
    < TypeArgument+ >

TypeArgument:
    WildcardIndicator_opt FieldTypeSignature
    *

WildcardIndicator:
    +
    -

ArrayTypeSignature:
    [ TypeSignature

TypeSignature:
    FieldTypeSignature
    BaseType
```

A method signature, defined by the production *MethodTypeSignature*, encodes the (possibly parameterized) types of the method's formal arguments and of the exceptions it has declared in its `throws` clause, its (possibly parameterized) return type, and any formal type parameters in the method declaration.

```
MethodTypeSignature:
    FormalTypeParameters_opt (TypeSignature*) ReturnType ThrowsSignature*

ReturnType:
    TypeSignature
    VoidDescriptor

ThrowsSignature:
    ^ ClassTypeSignature
    ^ TypeVariableSignature
```

If the `throws` clause of a method or constructor does not involve type variables, the

*ThowsSignature* may be elided from the *MethodTypeSignature*.

A Java compiler must output generic signature information for any class, interface, constructor or member whose generic signature in the Java programming language would include references to type variables or parameterized types.

*The signature and descriptor ([§4.3.3](#)) of a given method or constructor may not correspond exactly, due to compiler-generated artifacts. In particular, the number of TypeSignatures that encode formal arguments in MethodTypeSignature may be less than the number of ParameterDescriptors in MethodDescriptor.*

*Oracle's Java Virtual Machine implementation does not check the well-formedness of the signatures described in this subsection during loading or linking. Instead, these checks are deferred until the signatures are used by reflective methods, as specified in the API of* `Class` *and members of* `java.lang.reflect`. *Future versions of a Java Virtual Machine implementation may be required to perform some or all of these checks during loading or linking.*

## 4.4. The Constant Pool

Java Virtual Machine instructions do not rely on the run-time layout of classes, interfaces, class instances, or arrays. Instead, instructions refer to symbolic information in the `constant_pool` table.

All `constant_pool` table entries have the following general format:

```
cp_info {
    u1 tag;
    u1 info[];
}
```

Each item in the `constant_pool` table must begin with a 1-byte tag indicating the kind of `cp_info` entry. The contents of the `info` array vary with the value of `tag`. The valid tags and their values are listed in [Table 4.3](#). Each tag byte must be followed by two or more bytes giving information about the specific constant. The format of the additional information varies with the tag value.

**Table 4.3. Constant pool tags**

| Constant Type | Value |
|---|---|
| `CONSTANT_Class` | 7 |
| `CONSTANT_Fieldref` | 9 |
| `CONSTANT_Methodref` | 10 |
| `CONSTANT_InterfaceMethodref` | 11 |
| `CONSTANT_String` | 8 |
| `CONSTANT_Integer` | 3 |
| `CONSTANT_Float` | 4 |
| `CONSTANT_Long` | 5 |
| `CONSTANT_Double` | 6 |
| `CONSTANT_NameAndType` | 12 |
| `CONSTANT_Utf8` | 1 |
| `CONSTANT_MethodHandle` | 15 |
| `CONSTANT_MethodType` | 16 |
| `CONSTANT_InvokeDynamic` | 18 |

### 4.4.1. The `CONSTANT_Class_info` Structure

The `CONSTANT_Class_info` structure is used to represent a class or an interface:

```
CONSTANT_Class_info {
    u1 tag;
    u2 name_index;
}
```

The items of the `CONSTANT_Class_info` structure are the following:

tag

        The `tag` item has the value `CONSTANT_Class` (7).

name_index

> The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a valid binary class or interface name encoded in internal form ([§4.2.1](#)).

Because arrays are objects, the opcodes *anewarray* and *multianewarray* can reference array "classes" via `CONSTANT_Class_info` structures in the `constant_pool` table. For such array classes, the name of the class is the descriptor of the array type.

*For example, the class name representing a two-dimensional int array type*

```
int[][]
```

*is*

```
[[I
```

*The class name representing the type array of class `Thread`*

```
Thread[]
```

*is*

```
[Ljava/lang/Thread;
```

An array type descriptor is valid only if it represents 255 or fewer dimensions.

### 4.4.2. The `CONSTANT_Fieldref_info`, `CONSTANT_Methodref_info`, and `CONSTANT_InterfaceMethodref_info` Structures

Fields, methods, and interface methods are represented by similar structures:

```
CONSTANT_Fieldref_info {
    u1 tag;
    u2 class_index;
    u2 name_and_type_index;
}

CONSTANT_Methodref_info {
    u1 tag;
    u2 class_index;
    u2 name_and_type_index;
}

CONSTANT_InterfaceMethodref_info {
    u1 tag;
    u2 class_index;
    u2 name_and_type_index;
}
```

The items of these structures are as follows:

tag

> The tag item of a `CONSTANT_Fieldref_info` structure has the value
> `CONSTANT_Fieldref` (9).
>
> The tag item of a `CONSTANT_Methodref_info` structure has the value
> `CONSTANT_Methodref` (10).
>
> The tag item of a `CONSTANT_InterfaceMethodref_info` structure has the value
> `CONSTANT_InterfaceMethodref` (11).

class_index

> The value of the `class_index` item must be a valid index into the `constant_pool`
> table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info`
> ([§4.4.1](#)) structure representing a class or interface type that has the field or method as

a member.

The `class_index` item of a `CONSTANT_Methodref_info` structure must be a class type, not an interface type.

The `class_index` item of a `CONSTANT_InterfaceMethodref_info` structure must be an interface type.

The `class_index` item of a `CONSTANT_Fieldref_info` structure may be either a class type or an interface type.

`name_and_type_index`

The value of the `name_and_type_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_NameAndType_info` ([§4.4.6](#)) structure. This `constant_pool` entry indicates the name and descriptor of the field or method.

In a `CONSTANT_Fieldref_info`, the indicated descriptor must be a field descriptor ([§4.3.2](#)). Otherwise, the indicated descriptor must be a method descriptor ([§4.3.3](#)).

If the name of the method of a `CONSTANT_Methodref_info` structure begins with a '<' ('\u003c'), then the name must be the special name `<init>`, representing an instance initialization method ([§2.9](#)). The return type of such a method must be `void`.

### 4.4.3. The `CONSTANT_String_info` Structure

The `CONSTANT_String_info` structure is used to represent constant objects of the type `String`:

```
CONSTANT_String_info {
    u1 tag;
    u2 string_index;
}
```

The items of the `CONSTANT_String_info` structure are as follows:

tag

> The `tag` item of the `CONSTANT_String_info` structure has the value
> `CONSTANT_String` (8).

string_index

> The value of the `string_index` item must be a valid index into the `constant_pool`
> table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info`
> ([§4.4.7](#)) structure representing the sequence of Unicode code points to which the
> `String` object is to be initialized.

### 4.4.4. The `CONSTANT_Integer_info` and `CONSTANT_Float_info` Structures

The `CONSTANT_Integer_info` and `CONSTANT_Float_info` structures represent 4-byte
numeric (`int` and `float`) constants:

```
CONSTANT_Integer_info {
    u1 tag;
    u4 bytes;
}

CONSTANT_Float_info {
    u1 tag;
    u4 bytes;
}
```

The items of these structures are as follows:

tag

> The `tag` item of the `CONSTANT_Integer_info` structure has the value
> `CONSTANT_Integer` (3).
>
> The `tag` item of the `CONSTANT_Float_info` structure has the value `CONSTANT_Float`
> (4).

## bytes

The `bytes` item of the `CONSTANT_Integer_info` structure represents the value of the `int` constant. The bytes of the value are stored in big-endian (high byte first) order.

The `bytes` item of the `CONSTANT_Float_info` structure represents the value of the `float` constant in IEEE 754 floating-point single format ([§2.3.2](#)). The bytes of the single format representation are stored in big-endian (high byte first) order.

The value represented by the `CONSTANT_Float_info` structure is determined as follows. The bytes of the value are first converted into an `int` constant *bits*. Then:

- If *bits* is `0x7f800000`, the `float` value will be positive infinity.

- If *bits* is `0xff800000`, the `float` value will be negative infinity.

- If *bits* is in the range `0x7f800001` through `0x7fffffff` or in the range `0xff800001` through `0xffffffff`, the `float` value will be NaN.

- In all other cases, let `s`, `e`, and `m` be three values that might be computed from *bits*:

```
int s = ((bits >> 31) == 0) ? 1 : -1;
int e = ((bits >> 23) & 0xff);
int m = (e == 0) ?
           (bits & 0x7fffff) << 1 :
           (bits & 0x7fffff) | 0x800000;
```

Then the `float` value equals the result of the mathematical expression $s \cdot m \cdot 2^{e-150}$.

### 4.4.5. The `CONSTANT_Long_info` and `CONSTANT_Double_info` Structures

The `CONSTANT_Long_info` and `CONSTANT_Double_info` represent 8-byte numeric (`long` and `double`) constants:

```
CONSTANT_Long_info {
    u1 tag;
    u4 high_bytes;
    u4 low_bytes;
}

CONSTANT_Double_info {
    u1 tag;
    u4 high_bytes;
    u4 low_bytes;
}
```

All 8-byte constants take up two entries in the `constant_pool` table of the `class` file. If a `CONSTANT_Long_info` or `CONSTANT_Double_info` structure is the item in the `constant_pool` table at index *n*, then the next usable item in the pool is located at index *n*+2. The `constant_pool` index *n*+1 must be valid but is considered unusable.

*In retrospect, making 8-byte constants take two constant pool entries was a poor choice.*

The items of these structures are as follows:

tag

> The `tag` item of the `CONSTANT_Long_info` structure has the value `CONSTANT_Long` (5).

> The `tag` item of the `CONSTANT_Double_info` structure has the value `CONSTANT_Double` (6).

high_bytes, low_bytes

> The unsigned `high_bytes` and `low_bytes` items of the `CONSTANT_Long_info` structure together represent the value of the `long` constant

> ```
> ((long) high_bytes << 32) + low_bytes
> ```

> where the bytes of each of `high_bytes` and `low_bytes` are stored in big-endian (high

byte first) order.

The `high_bytes` and `low_bytes` items of the `CONSTANT_Double_info` structure together represent the `double` value in IEEE 754 floating-point double format (§2.3.2). The bytes of each item are stored in big-endian (high byte first) order.

The value represented by the `CONSTANT_Double_info` structure is determined as follows. The `high_bytes` and `low_bytes` items are converted into the `long` constant *bits*, which is equal to

```
((long) high_bytes << 32) + low_bytes
```

Then:

- If *bits* is `0x7ff0000000000000L`, the `double` value will be positive infinity.

- If *bits* is `0xfff0000000000000L`, the `double` value will be negative infinity.

- If *bits* is in the range `0x7ff0000000000001L` through `0x7fffffffffffffffL` or in the range `0xfff0000000000001L` through `0xffffffffffffffffL`, the `double` value will be NaN.

- In all other cases, let `s`, `e`, and `m` be three values that might be computed from *bits*:

```
int s = ((bits >> 63) == 0) ? 1 : -1;
int e = (int)((bits >> 52) & 0x7ffL);
long m = (e == 0) ?
          (bits & 0xfffffffffffffL) << 1 :
          (bits & 0xfffffffffffffL) | 0x10000000000000L;
```

Then the floating-point value equals the `double` value of the mathematical expression $s \cdot m \cdot 2^{e-1075}$.

### 4.4.6. The `CONSTANT_NameAndType_info` Structure

The `CONSTANT_NameAndType_info` structure is used to represent a field or method, without indicating which class or interface type it belongs to:

```
CONSTANT_NameAndType_info {
    u1 tag;
    u2 name_index;
    u2 descriptor_index;
}
```

The items of the `CONSTANT_NameAndType_info` structure are as follows:

tag

> The `tag` item of the `CONSTANT_NameAndType_info` structure has the value `CONSTANT_NameAndType` (12).

name_index

> The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing either the special method name `<init>` ([§2.9](#)) or a valid unqualified name ([§4.2.2](#)) denoting a field or method.

descriptor_index

> The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a valid field descriptor ([§4.3.2](#)) or method descriptor ([§4.3.3](#)).

### 4.4.7. The `CONSTANT_Utf8_info` Structure

The `CONSTANT_Utf8_info` structure is used to represent constant string values:

```
CONSTANT_Utf8_info {
    u1 tag;
```

```
    u2 length;
    u1 bytes[length];
}
```

The items of the `CONSTANT_Utf8_info` structure are the following:

tag

> The `tag` item of the `CONSTANT_Utf8_info` structure has the value `CONSTANT_Utf8` (1).

length

> The value of the `length` item gives the number of bytes in the `bytes` array (not the length of the resulting string). The strings in the `CONSTANT_Utf8_info` structure are not null-terminated.

bytes[]

> The `bytes` array contains the bytes of the string. No byte may have the value `(byte)0` or lie in the range `(byte)0xf0` - `(byte)0xff`.

String content is encoded in modified UTF-8. Modified UTF-8 strings are encoded so that code point sequences that contain only non-null ASCII characters can be represented using only 1 byte per code point, but all code points in the Unicode codespace can be represented.

- Code points in the range '\u0001' to '\u007F' are represented by a single byte:

  **Table 4.4.**

  | 0 | bits 6-0 |
  |---|----------|

  The 7 bits of data in the byte give the value of the code point represented.

- The null code point ('\u0000') and code points in the range '\u0080' to '\u07FF' are represented by a pair of bytes `x` and `y` :

  **Table 4.5.**

**Table 4.6.**

x:

| 1 | 1 | 0 | bits 10-6 |
|---|---|---|-----------|

**Table 4.7.**

y:

| 1 | 0 | bits 5-0 |
|---|---|----------|

The bytes represent the code point with the value:

```
((x & 0x1f) << 6) + (y & 0x3f)
```

- Code points in the range '\u0800' to '\uFFFF' are represented by 3 bytes x, y, and z :

**Table 4.8.**

**Table 4.9.**

x:

| 1 | 1 | 1 | 0 | bits 15-12 |
|---|---|---|---|------------|

**Table 4.10.**

y:

| 1 | 0 | bits 11-6 |
|---|---|-----------|

**Table 4.11.**

z:

| 1 | 0 | bits 5-0 |
|---|---|----------|

The three bytes represent the code point with the value:

```
((x & 0xf) << 12) + ((y & 0x3f) << 6) + (z & 0x3f)
```

- Characters with code points above U+FFFF (so-called *supplementary characters*) are represented by separately encoding the two surrogate code units of their UTF-16 representation. Each of the surrogate code units is represented by three bytes. This means supplementary characters are represented by six bytes, u, v, w, x, y, and z :

**Table 4.12.**

**Table 4.13.**

u:

| 1 | 1 | 1 | 0 | 1 | 1 | 0 | 1 |
|---|---|---|---|---|---|---|---|

**Table 4.14.**

v:

| 1 | 0 | 1 | 0 | (bits 20-16)-1 |
|---|---|---|---|---|

**Table 4.15.**

w:

| 1 | 0 | bits 15-10 |
|---|---|---|

**Table 4.16.**

x:

| 1 | 1 | 1 | 0 | 1 | 1 | 0 | 1 |
|---|---|---|---|---|---|---|---|

**Table 4.17.**

y:

| 1 | 0 | 1 | 1 | bits 9-6 |
|---|---|---|---|---|

**Table 4.18.**

z:

| 1 | 0 | bits 5-0 |
|---|---|----------|

The six bytes represent the code point with the value:

```
0x10000 + ((v & 0x0f) << 16) + ((w & 0x3f) << 10) +
((y & 0x0f) << 6) + (z & 0x3f)
```

The bytes of multibyte characters are stored in the `class` file in big-endian (high byte first) order.

There are two differences between this format and the "standard" UTF-8 format. First, the null character `(char)0` is encoded using the 2-byte format rather than the 1-byte format, so that modified UTF-8 strings never have embedded nulls. Second, only the 1-byte, 2-byte, and 3-byte formats of standard UTF-8 are used. The Java Virtual Machine does not recognize the four-byte format of standard UTF-8; it uses its own two-times-three-byte format instead.

*For more information regarding the standard UTF-8 format, see Section 3.9 Unicode Encoding Forms of The Unicode Standard, Version 6.0.0.*

### 4.4.8. The `CONSTANT_MethodHandle_info` Structure

The `CONSTANT_MethodHandle_info` structure is used to represent a method handle:

```
CONSTANT_MethodHandle_info {
    u1 tag;
    u1 reference_kind;
    u2 reference_index;
}
```

The items of the `CONSTANT_MethodHandle_info` structure are the following:

tag

> The `tag` item of the `CONSTANT_MethodHandle_info` structure has the value `CONSTANT_MethodHandle` (15).

reference_kind

> The value of the `reference_kind` item must be in the range 1 to 9. The value denotes the *kind* of this method handle, which characterizes its bytecode behavior ([§5.4.3.5](#)).

reference_index

> The value of the `reference_index` item must be a valid index into the `constant_pool` table.
>
> If the value of the `reference_kind` item is 1 (`REF_getField`), 2 (`REF_getStatic`), 3 (`REF_putField`), or 4 (`REF_putStatic`), then the `constant_pool` entry at that index must be a `CONSTANT_Fieldref_info` ([§4.4.2](#)) structure representing a field for which a method handle is to be created.
>
> If the value of the `reference_kind` item is 5 (`REF_invokeVirtual`), 6 (`REF_invokeStatic`), 7 (`REF_invokeSpecial`), or 8 (`REF_newInvokeSpecial`), then the `constant_pool` entry at that index must be a `CONSTANT_Methodref_info` structure ([§4.4.2](#)) representing a class's method or constructor ([§2.9](#)) for which a method handle is to be created.
>
> If the value of the `reference_kind` item is 9 (`REF_invokeInterface`), then the `constant_pool` entry at that index must be a `CONSTANT_InterfaceMethodref_info` ([§4.4.2](#)) structure representing an interface's method for which a method handle is to be created.
>
> If the value of the `reference_kind` item is 5 (`REF_invokeVirtual`), 6 (`REF_invokeStatic`), 7 (`REF_invokeSpecial`), or 9 (`REF_invokeInterface`), the name of the method represented by a `CONSTANT_Methodref_info` structure must not be `<init>` or `<clinit>`.
>
> If the value is 8 (`REF_newInvokeSpecial`), the name of the method represented by a

### 4.4.9. The `CONSTANT_MethodType_info` Structure

The `CONSTANT_MethodType_info` structure is used to represent a method type:

```
CONSTANT_MethodType_info {
    u1 tag;
    u2 descriptor_index;
}
```

The items of the `CONSTANT_MethodType_info` structure are as follows:

tag

> The `tag` item of the `CONSTANT_MethodType_info` structure has the value
> `CONSTANT_MethodType` (16).

descriptor_index

> The value of the `descriptor_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a method descriptor ([§4.3.3](#)).

### 4.4.10. The `CONSTANT_InvokeDynamic_info` Structure

The `CONSTANT_InvokeDynamic_info` structure is used by an *invokedynamic* instruction
([*§invokedynamic*](#)) to specify a bootstrap method, the dynamic invocation name, the argument
and return types of the call, and optionally, a sequence of additional constants called *static
arguments* to the bootstrap method.

```
CONSTANT_InvokeDynamic_info {
    u1 tag;
    u2 bootstrap_method_attr_index;
    u2 name_and_type_index;
```

```
    }
```

The items of the `CONSTANT_InvokeDynamic_info` structure are as follows:

tag

> The `tag` item of the `CONSTANT_InvokeDynamic_info` structure has the value `CONSTANT_InvokeDynamic` (18).

bootstrap_method_attr_index

> The value of the `bootstrap_method_attr_index` item must be a valid index into the `bootstrap_methods` array of the bootstrap method table ([§4.7.21](#)) of this `class` file.

name_and_type_index

> The value of the `name_and_type_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_NameAndType_info` ([§4.4.6](#)) structure representing a method name and method descriptor ([§4.3.3](#)).

## 4.5. Fields

Each field is described by a `field_info` structure. No two fields in one `class` file may have the same name and descriptor ([§4.3.2](#)).

The structure has the following format:

```
field_info {
    u2              access_flags;
    u2              name_index;
    u2              descriptor_index;
    u2              attributes_count;
    attribute_info attributes[attributes_count];
}
```

The items of the `field_info` structure are as follows:

access_flags

The value of the `access_flags` item is a mask of flags used to denote access permission to and properties of this field. The interpretation of each flag, when set, is as shown in [Table 4.4](#).

**Table 4.4. Field access and property flags**

| Flag Name | Value | Interpretation |
|---|---|---|
| ACC_PUBLIC | 0x0001 | Declared `public`; may be accessed from outside its package. |
| ACC_PRIVATE | 0x0002 | Declared `private`; usable only within the defining class. |
| ACC_PROTECTED | 0x0004 | Declared `protected`; may be accessed within subclasses. |
| ACC_STATIC | 0x0008 | Declared `static`. |
| ACC_FINAL | 0x0010 | Declared `final`; never directly assigned to after object construction (JLS §17.5). |
| ACC_VOLATILE | 0x0040 | Declared `volatile`; cannot be cached. |
| ACC_TRANSIENT | 0x0080 | Declared `transient`; not written or read by a persistent object manager. |
| ACC_SYNTHETIC | 0x1000 | Declared synthetic; not present in the source code. |
| ACC_ENUM | 0x4000 | Declared as an element of an `enum`. |

A field may be marked with the ACC_SYNTHETIC flag to indicate that it was generated by a compiler and does not appear in source code.

The ACC_ENUM flag indicates that this field is used to hold an element of an enumerated type.

Fields of classes may set any of the flags in [Table 4.4](#). However, a specific field of a class may have at most one of its ACC_PRIVATE, ACC_PROTECTED, and ACC_PUBLIC flags set (JLS §8.3.1) and must not have both its ACC_FINAL and ACC_VOLATILE flags set (JLS §8.3.1.4).

All fields of interfaces must have their `ACC_PUBLIC`, `ACC_STATIC`, and `ACC_FINAL` flags set; they may have their `ACC_SYNTHETIC` flag set and must not have any of the other flags in [Table 4.4](#) set (JLS §9.3).

All bits of the `access_flags` item not assigned in [Table 4.4](#) are reserved for future use. They should be set to zero in generated `class` files and should be ignored by Java Virtual Machine implementations.

name_index

The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure which must represent a valid unqualified name ([§4.2.2](#)) denoting a field.

descriptor_index

The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure that must represent a valid field descriptor ([§4.3.2](#)).

attributes_count

The value of the `attributes_count` item indicates the number of additional attributes ([§4.7](#)) of this field.

attributes[]

Each value of the `attributes` table must be an attribute structure ([§4.7](#)). A field can have any number of attributes associated with it.

The attributes defined by this specification as appearing in the `attributes` table of a `field_info` structure are ConstantValue ([§4.7.2](#)), Synthetic ([§4.7.8](#)), Signature ([§4.7.9](#)), Deprecated ([§4.7.15](#)), RuntimeVisibleAnnotations ([§4.7.16](#)) and RuntimeInvisibleAnnotations ([§4.7.17](#)).

A Java Virtual Machine implementation must recognize and correctly read ConstantValue ([§4.7.2](#)) attributes found in the `attributes` table of a `field_info` structure. If a Java Virtual Machine implementation recognizes `class` files whose

version number is 49.0 or above, it must recognize and correctly read `Signature` (§4.7.9), `RuntimeVisibleAnnotations` (§4.7.16) and `RuntimeInvisibleAnnotations` (§4.7.17) attributes found in the `attributes` table of a `field_info` structure of a `class` file whose version number is 49.0 or above.

A Java Virtual Machine implementation is required to silently ignore any or all attributes that it does not recognize in the `attributes` table of a `field_info` structure. Attributes not defined in this specification are not allowed to affect the semantics of the `class` file, but only to provide additional descriptive information (§4.7.1).

## 4.6. Methods

Each method, including each instance initialization method (§2.9) and the class or interface initialization method (§2.9), is described by a `method_info` structure. No two methods in one `class` file may have the same name and descriptor (§4.3.3).

The structure has the following format:

```
method_info {
    u2              access_flags;
    u2              name_index;
    u2              descriptor_index;
    u2              attributes_count;
    attribute_info attributes[attributes_count];
}
```

The items of the `method_info` structure are as follows:

access_flags

> The value of the `access_flags` item is a mask of flags used to denote access permission to and properties of this method. The interpretation of each flag, when set, is as shown in Table 4.5.

**Table 4.5. Method access and property flags**

| Flag Name | Value | Interpretation |
|---|---|---|

| | | |
|---|---|---|
| ACC_PUBLIC | 0x0001 | Declared `public`; may be accessed from outside its package. |
| ACC_PRIVATE | 0x0002 | Declared `private`; accessible only within the defining class. |
| ACC_PROTECTED | 0x0004 | Declared `protected`; may be accessed within subclasses. |
| ACC_STATIC | 0x0008 | Declared `static`. |
| ACC_FINAL | 0x0010 | Declared `final`; must not be overridden (§5.4.5). |
| ACC_SYNCHRONIZED | 0x0020 | Declared `synchronized`; invocation is wrapped by a monitor use. |
| ACC_BRIDGE | 0x0040 | A bridge method, generated by the compiler. |
| ACC_VARARGS | 0x0080 | Declared with variable number of arguments. |
| ACC_NATIVE | 0x0100 | Declared `native`; implemented in a language other than Java. |
| ACC_ABSTRACT | 0x0400 | Declared `abstract`; no implementation is provided. |
| ACC_STRICT | 0x0800 | Declared `strictfp`; floating-point mode is FP-strict. |
| ACC_SYNTHETIC | 0x1000 | Declared synthetic; not present in the source code. |

The ACC_VARARGS flag indicates that this method takes a variable number of arguments at the source code level. A method declared to take a variable number of arguments must be compiled with the ACC_VARARGS flag set to 1. All other methods must be compiled with the ACC_VARARGS flag set to 0.

The ACC_BRIDGE flag is used to indicate a bridge method generated by a Java compiler.

A method may be marked with the ACC_SYNTHETIC flag to indicate that it was generated by a compiler and does not appear in source code, unless it is one of the methods named in §4.7.8.

Methods of classes may set any of the flags in Table 4.5. However, a specific method of a class may have at most one of its ACC_PRIVATE, ACC_PROTECTED and ACC_PUBLIC flags set (JLS §8.4.3). If a specific method has its ACC_ABSTRACT flag

set, it must not have any of its ACC_FINAL, ACC_NATIVE, ACC_PRIVATE, ACC_STATIC, ACC_STRICT or ACC_SYNCHRONIZED flags set (JLS §8.4.3.1, JLS §8.4.3.3, JLS §8.4.3.4).

All interface methods must have their ACC_ABSTRACT and ACC_PUBLIC flags set; they may have their ACC_VARARGS, ACC_BRIDGE and ACC_SYNTHETIC flags set and must not have any of the other flags in Table 4.5 set (JLS §9.4).

A specific instance initialization method (§2.9) may have at most one of its ACC_PRIVATE, ACC_PROTECTED, and ACC_PUBLIC flags set, and may also have its ACC_STRICT, ACC_VARARGS and ACC_SYNTHETIC flags set, but must not have any of the other flags in Table 4.5 set.

Class and interface initialization methods (§2.9) are called implicitly by the Java Virtual Machine. The value of their access_flags item is ignored except for the setting of the ACC_STRICT flag.

All bits of the access_flags item not assigned in Table 4.5 are reserved for future use. They should be set to zero in generated class files and should be ignored by Java Virtual Machine implementations.

name_index

The value of the name_index item must be a valid index into the constant_pool table. The constant_pool entry at that index must be a CONSTANT_Utf8_info (§4.4.7) structure representing either one of the special method names (§2.9) <init> or <clinit>, or a valid unqualified name (§4.2.2) denoting a method.

descriptor_index

The value of the descriptor_index item must be a valid index into the constant_pool table. The constant_pool entry at that index must be a CONSTANT_Utf8_info (§4.4.7) structure representing a valid method descriptor (§4.3.3).

*A future edition of this specification may require that the last parameter descriptor of the method descriptor is an array type if the ACC_VARARGS flag is set in the access_flags item.*

attributes_count

The value of the `attributes_count` item indicates the number of additional attributes ([§4.7](#)) of this method.

attributes[]

Each value of the `attributes` table must be an attribute structure ([§4.7](#)). A method can have any number of optional attributes associated with it.

The attributes defined by this specification as appearing in the `attributes` table of a `method_info` structure are the Code ([§4.7.3](#)), Exceptions ([§4.7.5](#)), Synthetic ([§4.7.8](#)), Signature ([§4.7.9](#)), Deprecated ([§4.7.15](#)), RuntimeVisibleAnnotations ([§4.7.16](#)), RuntimeInvisibleAnnotations ([§4.7.17](#)), RuntimeVisibleParameterAnnotations ([§4.7.18](#)), RuntimeInvisibleParameterAnnotations ([§4.7.19](#)), and AnnotationDefault ([§4.7.20](#)) attributes.

A Java Virtual Machine implementation must recognize and correctly read Code ([§4.7.3](#)) and Exceptions ([§4.7.5](#)) attributes found in the `attributes` table of a `method_info` structure. If a Java Virtual Machine implementation recognizes `class` files whose version number is 49.0 or above, it must recognize and correctly read Signature ([§4.7.9](#)), RuntimeVisibleAnnotations ([§4.7.16](#)), RuntimeInvisibleAnnotations ([§4.7.17](#)), RuntimeVisibleParameterAnnotations ([§4.7.18](#)), RuntimeInvisibleParameterAnnotations ([§4.7.19](#)) and AnnotationDefault ([§4.7.20](#)) attributes found in the `attributes` table of a `method_info` structure of a `class` file whose version number is 49.0 or above.

A Java Virtual Machine implementation is required to silently ignore any or all attributes in the `attributes` table of a `method_info` structure that it does not recognize. Attributes not defined in this specification are not allowed to affect the semantics of the `class` file, but only to provide additional descriptive information ([§4.7.1](#)).

## 4.7. Attributes

Attributes are used in the `ClassFile`, `field_info`, `method_info`, and `Code_attribute` structures ([§4.1](#), [§4.5](#), [§4.6](#), [§4.7.3](#)) of the `class` file format. All attributes have the following general format:

```
attribute_info {
    u2 attribute_name_index;
    u4 attribute_length;
    u1 info[attribute_length];
}
```

For all attributes, the `attribute_name_index` must be a valid unsigned 16-bit index into the constant pool of the class. The `constant_pool` entry at `attribute_name_index` must be a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing the name of the attribute. The value of the `attribute_length` item indicates the length of the subsequent information in bytes. The length does not include the initial six bytes that contain the `attribute_name_index` and `attribute_length` items.

Certain attributes are predefined as part of the `class` file specification. They are listed in [Table 4.6](#), accompanied by the version of the Java SE platform and the version of the `class` file format in which each first appeared. Within the context of their use in this specification, that is, in the `attributes` tables of the `class` file structures in which they appear, the names of these predefined attributes are reserved. Of the predefined attributes:

- The `ConstantValue`, `Code` and `Exceptions` attributes must be recognized and correctly read by a `class` file reader for correct interpretation of the `class` file by a Java Virtual Machine implementation.

- The `InnerClasses`, `EnclosingMethod` and `Synthetic` attributes must be recognized and correctly read by a `class` file reader in order to properly implement the Java SE platform class libraries ([§2.12](#)).

- The `RuntimeVisibleAnnotations`, `RuntimeInvisibleAnnotations`, `RuntimeVisibleParameterAnnotations`, `RuntimeInvisibleParameterAnnotations` and `AnnotationDefault` attributes must be recognized and correctly read by a `class` file reader in order to properly implement the Java SE platform class libraries ([§2.12](#)), if the `class` file's version number is 49.0 or above and the Java Virtual Machine implementation recognizes `class` files whose version number is 49.0 or above.

- The `Signature` attribute must be recognized and correctly read by a `class` file reader if the `class` file's version number is 49.0 or above and the Java Virtual Machine implementation recognizes `class` files whose version number is 49.0 or above.

- The `StackMapTable` attribute must be recognized and correctly read by a `class` file reader if the `class` file's version number is 50.0 or above and the Java Virtual Machine implementation recognizes `class` files whose version number is 50.0 or above.

- The `BootstrapMethods` attribute must be recognized and correctly read by a `class` file reader if the `class` file's version number is 51.0 or above and the Java Virtual Machine implementation recognizes `class` files whose version number is 51.0 or above.

Use of the remaining predefined attributes is optional; a `class` file reader may use the information they contain, or otherwise must silently ignore those attributes.

**Table 4.6. Predefined `class` file attributes**

| Attribute | Section | Java SE | class file |
|-----------|---------|---------|------------|
| ConstantValue | §4.7.2 | 1.0.2 | 45.3 |
| Code | §4.7.3 | 1.0.2 | 45.3 |
| StackMapTable | §4.7.4 | 6 | 50.0 |
| Exceptions | §4.7.5 | 1.0.2 | 45.3 |
| InnerClasses | §4.7.6 | 1.1 | 45.3 |
| EnclosingMethod | §4.7.7 | 5.0 | 49.0 |
| Synthetic | §4.7.8 | 1.1 | 45.3 |
| Signature | §4.7.9 | 5.0 | 49.0 |
| SourceFile | §4.7.10 | 1.0.2 | 45.3 |
| SourceDebugExtension | §4.7.11 | 5.0 | 49.0 |
| LineNumberTable | §4.7.12 | 1.0.2 | 45.3 |
| LocalVariableTable | §4.7.13 | 1.0.2 | 45.3 |
| LocalVariableTypeTable | §4.7.14 | 5.0 | 49.0 |
| Deprecated | §4.7.15 | 1.1 | 45.3 |
| RuntimeVisibleAnnotations | §4.7.16 | 5.0 | 49.0 |
| RuntimeInvisibleAnnotations | §4.7.17 | 5.0 | 49.0 |
| RuntimeVisibleParameterAnnotations | §4.7.18 | 5.0 | 49.0 |

| | | | |
|---|---|---|---|
| RuntimeInvisibleParameterAnnotations | §4.7.19 | 5.0 | 49.0 |
| AnnotationDefault | §4.7.20 | 5.0 | 49.0 |
| BootstrapMethods | §4.7.21 | 7 | 51.0 |

### 4.7.1. Defining and Naming New Attributes

Compilers are permitted to define and emit `class` files containing new attributes in the `attributes` tables of `class` file structures. Java Virtual Machine implementations are permitted to recognize and use new attributes found in the `attributes` tables of `class` file structures. However, any attribute not defined as part of this Java Virtual Machine specification must not affect the semantics of class or interface types. Java Virtual Machine implementations are required to silently ignore attributes they do not recognize.

For instance, defining a new attribute to support vendor-specific debugging is permitted. Because Java Virtual Machine implementations are required to ignore attributes they do not recognize, `class` files intended for that particular Java Virtual Machine implementation will be usable by other implementations even if those implementations cannot make use of the additional debugging information that the `class` files contain.

Java Virtual Machine implementations are specifically prohibited from throwing an exception or otherwise refusing to use `class` files simply because of the presence of some new attribute. Of course, tools operating on `class` files may not run correctly if given `class` files that do not contain all the attributes they require.

Two attributes that are intended to be distinct, but that happen to use the same attribute name and are of the same length, will conflict on implementations that recognize either attribute. Attributes defined other than in this specification must have names chosen according to the package naming convention described in *The Java Language Specification, Java SE 7 Edition* (JLS §6.1).

Future versions of this specification may define additional attributes.

### 4.7.2. The `ConstantValue` Attribute

The `ConstantValue` attribute is a fixed-length attribute in the `attributes` table of a `field_info` structure ([§4.5](#)). A `ConstantValue` attribute represents the value of a constant field. There can be no more than one `ConstantValue` attribute in the `attributes` table of a given `field_info` structure. If the field is static (that is, the `ACC_STATIC` flag ([Table 4.4](#)) in the `access_flags` item of the `field_info` structure is set) then the constant field represented by the `field_info` structure is assigned the value referenced by its `ConstantValue` attribute as part of the initialization of the class or interface declaring the constant field ([§5.5](#)). This occurs prior to the invocation of the class or interface initialization method ([§2.9](#)) of that class or interface.

If a `field_info` structure representing a non-static field has a `ConstantValue` attribute, then that attribute must silently be ignored. Every Java Virtual Machine implementation must recognize `ConstantValue` attributes.

The `ConstantValue` attribute has the following format:

```
ConstantValue_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 constantvalue_index;
}
```

The items of the `ConstantValue_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "`ConstantValue`".

attribute_length

> The value of the `attribute_length` item of a `ConstantValue_attribute` structure must be 2.

constantvalue_index

> The value of the `constantvalue_index` item must be a valid index into the

constant_pool table. The constant_pool entry at that index gives the constant value represented by this attribute. The constant_pool entry must be of a type appropriate to the field, as shown by Table 4.7.

**Table 4.7. Constant value attribute types**

| Field Type | Entry Type |
|---|---|
| `long` | `CONSTANT_Long` |
| `float` | `CONSTANT_Float` |
| `double` | `CONSTANT_Double` |
| `int`, `short`, `char`, `byte`, `boolean` | `CONSTANT_Integer` |
| `String` | `CONSTANT_String` |

### 4.7.3. The Code Attribute

The Code attribute is a variable-length attribute in the `attributes` table of a `method_info` (§4.6) structure. A Code attribute contains the Java Virtual Machine instructions and auxiliary information for a single method, instance initialization method (§2.9), or class or interface initialization method (§2.9). Every Java Virtual Machine implementation must recognize Code attributes. If the method is either `native` or `abstract`, its `method_info` structure must not have a Code attribute. Otherwise, its `method_info` structure must have exactly one Code attribute.

The Code attribute has the following format:

```
Code_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 max_stack;
    u2 max_locals;
    u4 code_length;
    u1 code[code_length];
    u2 exception_table_length;
```

```
    {   u2 start_pc;
        u2 end_pc;
        u2 handler_pc;
        u2 catch_type;
    } exception_table[exception_table_length];
    u2 attributes_count;
    attribute_info attributes[attributes_count];
}
```

The items of the `Code_attribute` structure are as follows:

attribute_name_index

>	The value of the `attribute_name_index` item must be a valid index into the
>	`constant_pool` table. The `constant_pool` entry at that index must be a
>	`CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "Code".

attribute_length

>	The value of the `attribute_length` item indicates the length of the attribute,
>	excluding the initial six bytes.

max_stack

>	The value of the `max_stack` item gives the maximum depth of the operand stack of this
>	method ([§2.6.2](#)) at any point during execution of the method.

max_locals

>	The value of the `max_locals` item gives the number of local variables in the local
>	variable array allocated upon invocation of this method ([§2.6.1](#)), including the local
>	variables used to pass parameters to the method on its invocation.

>	The greatest local variable index for a value of type `long` or `double` is `max_locals` -
>	2. The greatest local variable index for a value of any other type is `max_locals` - 1.

code_length

>	The value of the `code_length` item gives the number of bytes in the code array for this
>	method. The value of `code_length` must be greater than zero; the code array must not

be empty.

code[]

The code array gives the actual bytes of Java Virtual Machine code that implement the method.

When the code array is read into memory on a byte-addressable machine, if the first byte of the array is aligned on a 4-byte boundary, the *tableswitch* and *lookupswitch* 32-bit offsets will be 4-byte aligned. (Refer to the descriptions of those instructions for more information on the consequences of code array alignment.)

The detailed constraints on the contents of the code array are extensive and are given in a separate section ([§4.9](#)).

exception_table_length

The value of the exception_table_length item gives the number of entries in the exception_table table.

exception_table[]

Each entry in the exception_table array describes one exception handler in the code array. The order of the handlers in the exception_table array is significant ([§2.10](#)).

Each exception_table entry contains the following four items:

start_pc, end_pc

The values of the two items start_pc and end_pc indicate the ranges in the code array at which the exception handler is active. The value of start_pc must be a valid index into the code array of the opcode of an instruction. The value of end_pc either must be a valid index into the code array of the opcode of an instruction or must be equal to code_length, the length of the code array. The value of start_pc must be less than the value of end_pc.

The start_pc is inclusive and end_pc is exclusive; that is, the exception handler must be active while the program counter is within the interval [start_pc, end_pc).

*The fact that end_pc is exclusive is a historical mistake in the design of the Java Virtual Machine: if the Java Virtual Machine code for a method is exactly 65535 bytes long and ends with an instruction that is 1 byte long, then that instruction cannot be protected by an exception handler. A compiler writer can work around this bug by limiting the maximum size of the generated Java Virtual Machine code for any method, instance initialization method, or static initializer (the size of any code array) to 65534 bytes.*

handler_pc

> The value of the `handler_pc` item indicates the start of the exception handler. The value of the item must be a valid index into the `code` array and must be the index of the opcode of an instruction.

catch_type

> If the value of the `catch_type` item is nonzero, it must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info` structure ([§4.4.1](#)) representing a class of exceptions that this exception handler is designated to catch. The exception handler will be called only if the thrown exception is an instance of the given class or one of its subclasses.

> If the value of the `catch_type` item is zero, this exception handler is called for all exceptions. This is used to implement `finally` ([§3.13](#)).

attributes_count

> The value of the `attributes_count` item indicates the number of attributes of the Code attribute.

attributes[]

> Each value of the `attributes` table must be an attribute structure ([§4.7](#)). A Code attribute can have any number of optional attributes associated with it.

> The only attributes defined by this specification as appearing in the `attributes` table of a Code attribute are the `LineNumberTable` ([§4.7.12](#)), `LocalVariableTable` ([§4.7.13](#)), `LocalVariableTypeTable` ([§4.7.14](#)), and `StackMapTable` ([§4.7.4](#)) attributes.

If a Java Virtual Machine implementation recognizes `class` files whose version number is 50.0 or above, it must recognize and correctly read `StackMapTable` ([§4.7.4](#)) attributes found in the `attributes` table of a `Code` attribute of a `class` file whose version number is 50.0 or above.

A Java Virtual Machine implementation is required to silently ignore any or all attributes in the `attributes` table of a `Code` attribute that it does not recognize. Attributes not defined in this specification are not allowed to affect the semantics of the `class` file, but only to provide additional descriptive information ([§4.7.1](#)).

### 4.7.4. The `StackMapTable` Attribute

The `StackMapTable` attribute is a variable-length attribute in the `attributes` table of a `Code` ([§4.7.3](#)) attribute. This attribute is used during the process of verification by type checking ([§4.10.1](#)). A method's `Code` attribute may have at most one `StackMapTable` attribute.

A `StackMapTable` attribute consists of zero or more *stack map frames*. Each stack map frame specifies (either explicitly or implicitly) a bytecode offset, the verification types ([§4.10.1.2](#)) for the local variables, and the verification types for the operand stack.

The type checker deals with and manipulates the expected types of a method's local variables and operand stack. Throughout this section, a *location* refers to either a single local variable or to a single operand stack entry.

We will use the terms *stack map frame* and *type state* interchangeably to describe a mapping from locations in the operand stack and local variables of a method to verification types. We will usually use the term *stack map frame* when such a mapping is provided in the `class` file, and the term *type state* when the mapping is used by the type checker.

In a `class` file whose version number is greater than or equal to 50.0, if a method's `Code` attribute does not have a `StackMapTable` attribute, it has an *implicit stack map attribute*. This implicit stack map attribute is equivalent to a `StackMapTable` attribute with `number_of_entries` equal to zero.

The `StackMapTable` attribute has the following format:

```
StackMapTable_attribute {
```

```
    u2               attribute_name_index;
    u4               attribute_length;
    u2               number_of_entries;
    stack_map_frame entries[number_of_entries];
}
```

The items of the `StackMapTable_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "`StackMapTable`".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute,
> excluding the initial six bytes.

number_of_entries

> The value of the `number_of_entries` item gives the number of `stack_map_frame`
> entries in the `entries` table.

entries

> The `entries` array gives the method's `stack_map_frame` structures.

Each `stack_map_frame` structure specifies the type state at a particular bytecode offset. Each
frame type specifies (explicitly or implicitly) a value, `offset_delta`, that is used to calculate
the actual bytecode offset at which a frame applies. The bytecode offset at which a frame
applies is calculated by adding `offset_delta + 1` to the bytecode offset of the previous
frame, unless the previous frame is the initial frame of the method, in which case the bytecode
offset is `offset_delta`.

*By using an offset delta rather than the actual bytecode offset we ensure, by definition, that stack map
frames are in the correctly sorted order. Furthermore, by consistently using the formula `offset_delta + 1`
for all explicit frames, we guarantee the absence of duplicates.*

We say that an instruction in the bytecode has a corresponding stack map frame if the

instruction starts at offset *i* in the `code` array of a `Code` attribute, and the `Code` attribute has a `StackMapTable` attribute whose entries array has a `stack_map_frame` structure that applies at bytecode offset *i*.

The `stack_map_frame` structure consists of a one-byte tag followed by zero or more bytes, giving more information, depending upon the tag.

A stack map frame may belong to one of several *frame types*:

```
union stack_map_frame {
    same_frame;
    same_locals_1_stack_item_frame;
    same_locals_1_stack_item_frame_extended;
    chop_frame;
    same_frame_extended;
    append_frame;
    full_frame;
}
```

All frame types, even `full_frame`, rely on the previous frame for some of their semantics. This raises the question of what is the very first frame? The initial frame is implicit, and computed from the method descriptor. (See the Prolog predicate `methodInitialStackFrame` ([§4.10.1.6](#)).)

- The frame type `same_frame` is represented by tags in the range [0-63]. If the frame type is `same_frame`, it means the frame has exactly the same locals as the previous stack map frame and that the number of stack items is zero. The `offset_delta` value for the frame is the value of the tag item, `frame_type`.

```
same_frame {
    u1 frame_type = SAME; /* 0-63 */
}
```

- The frame type `same_locals_1_stack_item_frame` is represented by tags in the range [64, 127]. If the frame_type is `same_locals_1_stack_item_frame`, it means the frame has exactly the same locals as the previous stack map frame and that the number of stack

items is 1. The `offset_delta` value for the frame is the value (`frame_type - 64`). There is a `verification_type_info` following the `frame_type` for the one stack item.

```
same_locals_1_stack_item_frame {
    u1 frame_type = SAME_LOCALS_1_STACK_ITEM; /* 64-127 */
    verification_type_info stack[1];
}
```

Tags in the range [128-246] are reserved for future use.

- The frame type `same_locals_1_stack_item_frame_extended` is represented by the tag 247. The frame type `same_locals_1_stack_item_frame_extended` indicates that the frame has exactly the same locals as the previous stack map frame and that the number of stack items is 1. The `offset_delta` value for the frame is given explicitly. There is a `verification_type_info` following the `frame_type` for the one stack item.

```
same_locals_1_stack_item_frame_extended {
    u1 frame_type = SAME_LOCALS_1_STACK_ITEM_EXTENDED; /* 247 */
    u2 offset_delta;
    verification_type_info stack[1];
}
```

- The frame type `chop_frame` is represented by tags in the range [248-250]. If the `frame_type` is `chop_frame`, it means that the operand stack is empty and the current locals are the same as the locals in the previous frame, except that the *k* last locals are absent. The value of *k* is given by the formula `251 - frame_type`.

```
chop_frame {
    u1 frame_type = CHOP; /* 248-250 */
    u2 offset_delta;
}
```

- The frame type `same_frame_extended` is represented by the tag value 251. If the frame

type is `same_frame_extended`, it means the frame has exactly the same locals as the previous stack map frame and that the number of stack items is zero.

```
same_frame_extended {
    u1 frame_type = SAME_FRAME_EXTENDED; /* 251 */
    u2 offset_delta;
}
```

- The frame type `append_frame` is represented by tags in the range [252-254]. If the `frame_type` is `append_frame`, it means that the operand stack is empty and the current locals are the same as the locals in the previous frame, except that *k* additional locals are defined. The value of *k* is given by the formula `frame_type - 251`.

```
append_frame {
    u1 frame_type = APPEND; /* 252-254 */
    u2 offset_delta;
    verification_type_info locals[frame_type - 251];
}
```

The `0`th entry in `locals` represents the type of the first additional local variable. If `locals[M]` represents local variable `N`, then `locals[M+1]` represents local variable `N+1` if `locals[M]` is one of:

- `Top_variable_info`

- `Integer_variable_info`

- `Float_variable_info`

- `Null_variable_info`

- `UninitializedThis_variable_info`

- `Object_variable_info`

- Uninitialized_variable_info

Otherwise `locals[M+1]` represents local variable N+2.

**It is an error if, for any index *i*, `locals[i]` represents a local variable whose index is greater than the maximum number of local variables for the method.**

- The frame type `full_frame` is represented by the tag value 255.

```
full_frame {
    u1 frame_type = FULL_FRAME; /* 255 */
    u2 offset_delta;
    u2 number_of_locals;
    verification_type_info locals[number_of_locals];
    u2 number_of_stack_items;
    verification_type_info stack[number_of_stack_items];
}
```

The 0th entry in `locals` represents the type of local variable 0. If `locals[M]` represents local variable N, then `locals[M+1]` represents local variable N+1 if `locals[M]` is one of:

- `Top_variable_info`

- `Integer_variable_info`

- `Float_variable_info`

- `Null_variable_info`

- `UninitializedThis_variable_info`

- `Object_variable_info`

- `Uninitialized_variable_info`

Otherwise `locals[M+1]` represents local variable N+2.

**It is an error if, for any index *i*, `locals[i]` represents a local variable whose index is greater than the maximum number of local variables for the method.**

The 0th entry in stack represents the type of the bottom of the stack, and subsequent entries represent types of stack elements closer to the top of the operand stack. We shall refer to the bottom element of the stack as stack element 0, and to subsequent elements as stack element 1, 2 etc. If stack[M] represents stack element N, then stack[M+1] represents stack element N+1 if stack[M] is one of:

- Top_variable_info

- Integer_variable_info

- Float_variable_info

- Null_variable_info

- UninitializedThis_variable_info

- Object_variable_info

- Uninitialized_variable_info

Otherwise, stack[M+1] represents stack element N+2.

**It is an error if, for any index *i*, stack[*i*] represents a stack entry whose index is greater than the maximum operand stack size for the method.**

The verification_type_info structure consists of a one-byte tag followed by zero or more bytes, giving more information about the tag. Each verification_type_info structure specifies the verification type of one or two locations.

```
union verification_type_info {
    Top_variable_info;
    Integer_variable_info;
    Float_variable_info;
    Long_variable_info;
    Double_variable_info;
    Null_variable_info;
    UninitializedThis_variable_info;
    Object_variable_info;
    Uninitialized_variable_info;
```

```
 }
```

- The `Top_variable_info` type indicates that the local variable has the verification type
  `top`.

```
Top_variable_info {
    u1 tag = ITEM_Top; /* 0 */
}
```

- The `Integer_variable_info` type indicates that the location contains the verification
  type `int`.

```
Integer_variable_info {
    u1 tag = ITEM_Integer; /* 1 */
}
```

- The `Float_variable_info` type indicates that the location contains the verification type
  `float`.

```
Float_variable_info {
    u1 tag = ITEM_Float; /* 2 */
}
```

- The `Long_variable_info` type indicates that the location contains the verification type
  `long`.

```
Long_variable_info {
    u1 tag = ITEM_Long; /* 4 */
}
```

This structure gives the contents of two locations in the operand stack or in the local variable array.

If the location is a local variable, then:

- It must not be the local variable with the highest index.

- The next higher numbered local variable contains the verification type `top`.

If the location is an operand stack entry, then:

- The current location must not be the topmost location of the operand stack.

- The next location closer to the top of the operand stack contains the verification type `top`.

- The `Double_variable_info` type indicates that the location contains the verification type `double`.

```
Double_variable_info {
    u1 tag = ITEM_Double; /* 3 */
}
```

This structure gives the contents of two locations in the operand stack or in the local variable array.

If the location is a local variable, then:

- It must not be the local variable with the highest index.

- The next higher numbered local variable contains the verification type `top`.

If the location is an operand stack entry, then:

- The current location must not be the topmost location of the operand stack.

- The next location closer to the top of the operand stack contains the verification type `top`.

- The `Null_variable_info` type indicates that location contains the verification type `null`.

```
Null_variable_info {
    u1 tag = ITEM_Null; /* 5 */
}
```

- The `UninitializedThis_variable_info` type indicates that the location contains the verification type `uninitializedThis`.

```
UninitializedThis_variable_info {
    u1 tag = ITEM_UninitializedThis; /* 6 */
}
```

- The `Object_variable_info` type indicates that the location contains an instance of the class represented by the `CONSTANT_Class_info` ([§4.4.1](#)) structure found in the `constant_pool` table at the index given by `cpool_index`.

```
Object_variable_info {
    u1 tag = ITEM_Object; /* 7 */
    u2 cpool_index;
}
```

- The `Uninitialized_variable_info` type indicates that the location contains the verification type `uninitialized(offset)`. The `offset` item indicates the offset, in the code array of the Code attribute ([§4.7.3](#)) that contains this `StackMapTable` attribute, of the *new* instruction ([§*new*](#)) that created the object being stored in the location.

```
Uninitialized_variable_info {
    u1 tag = ITEM_Uninitialized /* 8 */
    u2 offset;
}
```

### 4.7.5. The Exceptions Attribute

The `Exceptions` attribute is a variable-length attribute in the `attributes` table of a `method_info` structure (§4.6). The `Exceptions` attribute indicates which checked exceptions a method may throw. There may be at most one `Exceptions` attribute in each `method_info` structure.

The `Exceptions` attribute has the following format:

```
Exceptions_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 number_of_exceptions;
    u2 exception_index_table[number_of_exceptions];
}
```

The items of the `Exceptions_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be the `CONSTANT_Utf8_info` (§4.4.7) structure representing the string "Exceptions".

attribute_length

> The value of the `attribute_length` item indicates the attribute length, excluding the initial six bytes.

number_of_exceptions

> The value of the `number_of_exceptions` item indicates the number of entries in the `exception_index_table`.

exception_index_table[]

> Each value in the `exception_index_table` array must be a valid index into the

constant_pool table. The constant_pool entry referenced by each table item must be a CONSTANT_Class_info structure ([§4.4.1](#)) representing a class type that this method is declared to throw.

*A method should throw an exception only if at least one of the following three criteria is met:*

- *The exception is an instance of RuntimeException or one of its subclasses.*

- *The exception is an instance of Error or one of its subclasses.*

- *The exception is an instance of one of the exception classes specified in the exception_index_table just described, or one of their subclasses.*

*These requirements are not enforced in the Java Virtual Machine; they are enforced only at compile-time.*

### 4.7.6. The InnerClasses Attribute

The InnerClasses attribute is a variable-length attribute in the attributes table of a ClassFile structure ([§4.1](#)). If the constant pool of a class or interface C contains a CONSTANT_Class_info entry which represents a class or interface that is not a member of a package, then C's ClassFile structure must have exactly one InnerClasses attribute in its attributes table.

The InnerClasses attribute has the following format:

```
InnerClasses_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 number_of_classes;
    {   u2 inner_class_info_index;
        u2 outer_class_info_index;
        u2 inner_name_index;
        u2 inner_class_access_flags;
    } classes[number_of_classes];
}
```

The items of the InnerClasses_attribute structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a CONSTANT_Utf8_info ([§4.4.7](#)) structure representing the string "InnerClasses".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

number_of_classes

> The value of the `number_of_classes` item indicates the number of entries in the `classes` array.

classes[]

> Every CONSTANT_Class_info entry in the `constant_pool` table which represents a class or interface C that is not a package member must have exactly one corresponding entry in the `classes` array.
>
> If a class has members that are classes or interfaces, its `constant_pool` table (and hence its `InnerClasses` attribute) must refer to each such member, even if that member is not otherwise mentioned by the class. These rules imply that a nested class or interface member will have `InnerClasses` information for each enclosing class and for each immediate member.
>
> Each `classes` array entry contains the following four items:
>
> inner_class_info_index
>
> > The value of the `inner_class_info_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a CONSTANT_Class_info structure ([§4.4.1](#)) representing C. The remaining items in the `classes` array entry give information about C.
>
> outer_class_info_index
>
> > If C is not a member of a class or an interface (that is, if C is a top-level class or

interface (JLS §7.6) or a local class (JLS §14.3) or an anonymous class (JLS §15.9.5)), the value of the `outer_class_info_index` item must be zero.

Otherwise, the value of the `outer_class_info_index` item must be a valid index into the `constant_pool` table, and the entry at that index must be a `CONSTANT_Class_info` ([§4.4.1](#)) structure representing the class or interface of which C is a member.

inner_name_index

If C is anonymous (JLS §15.9.5), the value of the `inner_name_index` item must be zero.

Otherwise, the value of the `inner_name_index` item must be a valid index into the `constant_pool` table, and the entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure that represents the original simple name of C, as given in the source code from which this `class` file was compiled.

inner_class_access_flags

The value of the `inner_class_access_flags` item is a mask of flags used to denote access permissions to and properties of class or interface C as declared in the source code from which this `class` file was compiled. It is used by a compiler to recover the original information when source code is not available. The flags are shown in [Table 4.8](#).

**Table 4.8. Nested class access and property flags**

| Flag Name | Value | Interpretation |
|-----------|-------|----------------|
| ACC_PUBLIC | 0x0001 | Marked or implicitly `public` in source. |
| ACC_PRIVATE | 0x0002 | Marked `private` in source. |
| ACC_PROTECTED | 0x0004 | Marked `protected` in source. |
| ACC_STATIC | 0x0008 | Marked or implicitly `static` in source. |
| ACC_FINAL | 0x0010 | Marked `final` in source. |
| ACC_INTERFACE | 0x0200 | Was an `interface` in source. |

| ACC_ABSTRACT | 0x0400 | Marked or implicitly `abstract` in source. |
|---|---|---|
| ACC_SYNTHETIC | 0x1000 | Declared synthetic; not present in the source code. |
| ACC_ANNOTATION | 0x2000 | Declared as an annotation type. |
| ACC_ENUM | 0x4000 | Declared as an enum type. |

All bits of the `inner_class_access_flags` item not assigned in [Table 4.8](#) are reserved for future use. They should be set to zero in generated `class` files and should be ignored by Java Virtual Machine implementations.

If a `class` file has a version number that is greater than or equal to 51.0, and has an `InnerClasses` attribute in its `attributes` table, then for all entries in the `classes` array of the `InnerClasses` attribute, the value of the `outer_class_info_index` item must be zero if the value of the `inner_name_index` item is zero.

*Oracle's Java Virtual Machine implementation does not check the consistency of an `InnerClasses` attribute against a `class` file representing a class or interface referenced by the attribute.*

### 4.7.7. The EnclosingMethod Attribute

The `EnclosingMethod` attribute is an optional fixed-length attribute in the `attributes` table of a `ClassFile` structure ([§4.1](#)). A class must have an `EnclosingMethod` attribute if and only if it is a local class or an anonymous class. A class may have no more than one `EnclosingMethod` attribute.

The `EnclosingMethod` attribute has the following format:

```
EnclosingMethod_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 class_index;
    u2 method_index;
}
```

The items of the `EnclosingMethod_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "EnclosingMethod".

attribute_length

> The value of the `attribute_length` item is four.

class_index

> The value of the `class_index` item must be a valid index into the `constant_pool`
> table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info`
> ([§4.4.1](#)) structure representing the innermost class that encloses the declaration of the
> current class.

method_index

> If the current class is not immediately enclosed by a method or constructor, then the
> value of the `method_index` item must be zero.
>
> Otherwise, the value of the `method_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_NameAndType_info` structure ([§4.4.6](#)) representing the name and type of a
> method in the class referenced by the `class_index` attribute above.
>
> *It is the responsibility of a Java compiler to ensure that the method identified via the `method_index`
> is indeed the closest lexically enclosing method of the class that contains this `EnclosingMethod`
> attribute.*

## 4.7.8. The Synthetic Attribute

The `Synthetic` attribute is a fixed-length attribute in the `attributes` table of a `ClassFile`,
`field_info`, or `method_info` structure ([§4.1](#), [§4.5](#), [§4.6](#)). A class member that does not
appear in the source code must be marked using a `Synthetic` attribute, or else it must have
its `ACC_SYNTHETIC` flag set. The only exceptions to this requirement are compiler-generated

methods which are not considered implementation artifacts, namely the instance initialization method representing a default constructor of the Java programming language ([§2.9](#)), the class initialization method ([§2.9](#)), and the `Enum.values()` and `Enum.valueOf()` methods.

*The `Synthetic` attribute was introduced in JDK release 1.1 to support nested classes and interfaces.*

The `Synthetic` attribute has the following format:

```
Synthetic_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
}
```

The items of the `Synthetic_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "`Synthetic`".

attribute_length

> The value of the `attribute_length` item is zero.


## 4.7.9. The `Signature` Attribute

The `Signature` attribute is an optional fixed-length attribute in the `attributes` table of a `ClassFile`, `field_info`, or `method_info` structure ([§4.1](#), [§4.5](#), [§4.6](#)). The `Signature` attribute records generic signature information for any class, interface, constructor or member whose generic signature in the Java programming language would include references to type variables or parameterized types.

The `Signature` attribute has the following format:

```
Signature_attribute {
    u2 attribute_name_index;
```

```
    u4 attribute_length;
    u2 signature_index;
}
```

The items of the `Signature_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "`Signature`".

attribute_length

> The value of the `attribute_length` item of a `Signature_attribute` structure must
> be 2.

signature_index

> The value of the `signature_index` item must be a valid index into the
> `constant_pool` table. The constant pool entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a class signature ([§4.3.4](#)) if this
> `Signature` attribute is an attribute of a `ClassFile` structure; a method signature if this
> `Signature` attribute is an attribute of a `method_info` structure; or a field type
> signature otherwise.

### 4.7.10. The SourceFile Attribute

The `SourceFile` attribute is an optional fixed-length attribute in the `attributes` table of a
`ClassFile` structure ([§4.1](#)). There can be no more than one `SourceFile` attribute in the
`attributes` table of a given `ClassFile` structure.

The `SourceFile` attribute has the following format:

```
SourceFile_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
```

```
        u2 sourcefile_index;
}
```

The items of the `SourceFile_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "SourceFile".

attribute_length

> The value of the `attribute_length` item of a `SourceFile_attribute` structure
> must be 2.

sourcefile_index

> The value of the `sourcefile_index` item must be a valid index into the
> `constant_pool` table. The constant pool entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a string.

> The string referenced by the `sourcefile_index` item will be interpreted as indicating
> the name of the source file from which this `class` file was compiled. It will not be
> interpreted as indicating the name of a directory containing the file or an absolute path
> name for the file; such platform-specific additional information must be supplied by the
> run-time interpreter or development tool at the time the file name is actually used.

### 4.7.11. The SourceDebugExtension Attribute

The `SourceDebugExtension` attribute is an optional attribute in the `attributes` table of a
`ClassFile` structure ([§4.1](#)). There can be no more than one `SourceDebugExtension`
attribute in the `attributes` table of a given `ClassFile` structure.

The `SourceDebugExtension` attribute has the following format:

```
SourceDebugExtension_attribute {
```

```
    u2 attribute_name_index;
    u4 attribute_length;
    u1 debug_extension[attribute_length];
}
```

The items of the `SourceDebugExtension_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "SourceDebugExtension".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

> The value of the `attribute_length` item is thus the number of bytes in the `debug_extension[]` item.

debug_extension[]

> The `debug_extension` array holds extended debugging information which has no semantic effect on the Java Virtual Machine. The information is represented using a modified UTF-8 string ([§4.4.7](#)) with no terminating zero byte.

> *Note that the `debug_extension` array may denote a string longer than that which can be represented with an instance of class `String`.*

### 4.7.12. The `LineNumberTable` Attribute

The `LineNumberTable` attribute is an optional variable-length attribute in the `attributes` table of a Code ([§4.7.3](#)) attribute. It may be used by debuggers to determine which part of the Java Virtual Machine code array corresponds to a given line number in the original source file.

If `LineNumberTable` attributes are present in the `attributes` table of a given Code attribute,

then they may appear in any order. Furthermore, multiple `LineNumberTable` attributes may together represent a given line of a source file; that is, `LineNumberTable` attributes need not be one-to-one with source lines.

The `LineNumberTable` attribute has the following format:

```
LineNumberTable_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 line_number_table_length;
    {   u2 start_pc;
        u2 line_number;
    } line_number_table[line_number_table_length];
}
```

The items of the `LineNumberTable_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "LineNumberTable".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

line_number_table_length

> The value of the `line_number_table_length` item indicates the number of entries in the `line_number_table` array.

line_number_table[]

> Each entry in the `line_number_table` array indicates that the line number in the original source file changes at a given point in the `code` array. Each `line_number_table` entry must contain the following two items:

start_pc

> The value of the `start_pc` item must indicate the index into the `code` array at which the code for a new line in the original source file begins.

> The value of `start_pc` must be less than the value of the `code_length` item of the `Code` attribute of which this `LineNumberTable` is an attribute.

line_number

> The value of the `line_number` item must give the corresponding line number in the original source file.

### 4.7.13. The `LocalVariableTable` Attribute

The `LocalVariableTable` attribute is an optional variable-length attribute in the `attributes` table of a Code ([§4.7.3](#)) attribute. It may be used by debuggers to determine the value of a given local variable during the execution of a method.

If `LocalVariableTable` attributes are present in the `attributes` table of a given Code attribute, then they may appear in any order. There may be no more than one `LocalVariableTable` attribute per local variable in the Code attribute.

The `LocalVariableTable` attribute has the following format:

```
LocalVariableTable_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 local_variable_table_length;
    {   u2 start_pc;
        u2 length;
        u2 name_index;
        u2 descriptor_index;
        u2 index;
    } local_variable_table[local_variable_table_length];
}
```

The items of the `LocalVariableTable_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string `"LocalVariableTable"`.

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

local_variable_table_length

> The value of the `local_variable_table_length` item indicates the number of entries in the `local_variable_table` array.

local_variable_table[]

> Each entry in the `local_variable_table` array indicates a range of code array offsets within which a local variable has a value. It also indicates the index into the local variable array of the current frame at which that local variable can be found. Each entry must contain the following five items:

> start_pc, length

>> The given local variable must have a value at indices into the `code` array in the interval [`start_pc`, `start_pc + length`), that is, between `start_pc` inclusive and `start_pc + length` exclusive.

>> The value of `start_pc` must be a valid index into the `code` array of this Code attribute and must be the index of the opcode of an instruction.

>> The value of `start_pc + length` must either be a valid index into the `code` array of this Code attribute and be the index of the opcode of an instruction, or it must be the first index beyond the end of that `code` array.

> name_index

The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a valid unqualified name ([§4.2.2](#)) denoting a local variable.

descriptor_index

The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing a field descriptor ([§4.3.2](#)) encoding the type of a local variable in the `source program`.

index

The given local variable must be at `index` in the local variable array of the current frame.

If the local variable at `index` is of type `double` or `long`, it occupies both `index` and `index + 1`.

### 4.7.14. The `LocalVariableTypeTable` Attribute

The `LocalVariableTypeTable` attribute is an optional variable-length attribute in the `attributes` table of a Code ([§4.7.3](#)) attribute. It may be used by debuggers to determine the value of a given local variable during the execution of a `method`.

If `LocalVariableTypeTable` attributes are present in the `attributes` table of a given Code attribute, then they may appear in any order. There may be no more than one `LocalVariableTypeTable` attribute per local variable in the Code attribute.

The `LocalVariableTypeTable` attribute differs from the `LocalVariableTable` attribute in that it provides signature information rather than descriptor information. This difference is only significant for variables whose type is a generic reference type. Such variables will appear in both tables, while variables of other types will appear only in `LocalVariableTable`.

The `LocalVariableTypeTable` attribute has the following format:

```
LocalVariableTypeTable_attribute {
```

```
    u2 attribute_name_index;
    u4 attribute_length;
    u2 local_variable_type_table_length;
    {   u2 start_pc;
        u2 length;
        u2 name_index;
        u2 signature_index;
        u2 index;
    } local_variable_type_table[local_variable_type_table_length];
}
```

The items of the `LocalVariableTypeTable_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "LocalVariableTypeTable".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

local_variable_type_table_length

> The value of the `local_variable_type_table_length` item indicates the number of entries in the `local_variable_type_table` array.

local_variable_type_table[]

> Each entry in the `local_variable_type_table` array indicates a range of code array offsets within which a local variable has a value. It also indicates the index into the local variable array of the current frame at which that local variable can be found. Each entry must contain the following five items:

start_pc, length

> The given local variable must have a value at indices into the code array in the

interval [`start_pc`, `start_pc + length`), that is, between `start_pc` inclusive and `start_pc + length` exclusive.

The value of `start_pc` must be a valid index into the `code` array of this `Code` attribute and must be the index of the opcode of an instruction.

The value of `start_pc + length` must either be a valid index into the `code` array of this `Code` attribute and be the index of the opcode of an instruction, or it must be the first index beyond the end of that `code` array.

name_index

The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing a valid unqualified name ([§4.2.2](#)) denoting a local variable.

signature_index

The value of the `signature_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing a field type signature ([§4.3.4](#)) encoding the type of a local variable in the source program.

index

The given local variable must be at `index` in the local variable array of the current frame.

If the local variable at `index` is of type `double` or `long`, it occupies both `index` and `index + 1`.

## 4.7.15. The Deprecated Attribute

The `Deprecated` attribute is an optional fixed-length attribute in the `attributes` table of a `ClassFile`, `field_info`, or `method_info` structure ([§4.1](#), [§4.5](#), [§4.6](#)). A class, interface, method, or field may be marked using a `Deprecated` attribute to indicate that the class, interface, method, or field has been superseded.

A run-time interpreter or tool that reads the `class` file format, such as a compiler, can use this marking to advise the user that a superceded class, interface, method, or field is being referred to. The presence of a `Deprecated` attribute does not alter the semantics of a class or interface.

The `Deprecated` attribute has the following format:

```
Deprecated_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
}
```

The items of the `Deprecated_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "Deprecated".

attribute_length

> The value of the `attribute_length` item is zero.

### 4.7.16. The `RuntimeVisibleAnnotations` attribute

The `RuntimeVisibleAnnotations` attribute is a variable-length attribute in the `attributes` table of a `ClassFile`, `field_info`, or `method_info` structure ([§4.1](#), [§4.5](#), [§4.6](#)). The `RuntimeVisibleAnnotations` attribute records run-time-visible Java programming language annotations on the corresponding class, field, or method.

Each `ClassFile`, `field_info`, and `method_info` structure may contain at most one `RuntimeVisibleAnnotations` attribute, which records all the run-time-visible Java programming language annotations on the corresponding program element. The Java Virtual Machine must make these annotations available so they can be returned by the appropriate reflective APIs.

The `RuntimeVisibleAnnotations` attribute has the following format:

```
RuntimeVisibleAnnotations_attribute {
    u2          attribute_name_index;
    u4          attribute_length;
    u2          num_annotations;
    annotation annotations[num_annotations];
}
```

The items of the `RuntimeVisibleAnnotations_attribute` structure are as follows:

attribute_name_index

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string "RuntimeVisibleAnnotations".

attribute_length

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

The value of the `attribute_length` item is thus dependent on the number of run-time-visible annotations represented by the structure, and their values.

num_annotations

The value of the `num_annotations` item gives the number of run-time-visible annotations represented by the structure.

*Note that a maximum of 65535 run-time-visible Java programming language annotations may be directly attached to a program element.*

annotations

Each value of the `annotations` table represents a single run-time-visible annotation on a program element. The annotation structure has the following format:

```
annotation {
    u2 type_index;
    u2 num_element_value_pairs;
    {   u2              element_name_index;
        element_value value;
    } element_value_pairs[num_element_value_pairs];
}
```

The items of the annotation structure are as follows:

type_index

> The value of the type_index item must be a valid index into the
> constant_pool table. The constant_pool entry at that index must be a
> CONSTANT_Utf8_info ([§4.4.7](#)) structure representing a field descriptor
> representing the annotation type corresponding to the annotation represented
> by this annotation structure.

num_element_value_pairs

> The value of the num_element_value_pairs item gives the number of
> element-value pairs of the annotation represented by this annotation
> structure.
>
> *Note that a maximum of 65535 element-value pairs may be contained in a single annotation.*

element_value_pairs

> Each value of the element_value_pairs table represents a single element-
> value pair in the annotation represented by this annotation structure. Each
> element_value_pairs entry contains the following two items:

element_name_index

> > The value of the element_name_index item must be a valid index into
> > the constant_pool table. The constant_pool entry at that index must
> > be a CONSTANT_Utf8_info structure ([§4.4.7](#)) representing a valid field
> > descriptor ([§4.3.2](#)) that denotes the name of the annotation type element

represented by this `element_value_pairs` entry.

value

The value of the `value` item represents the value of the element-value pair represented by this `element_value_pairs` entry.

### 4.7.16.1. The `element_value` structure

The `element_value` structure is a discriminated union representing the value of an element-value pair. It is used to represent element values in all attributes that describe annotations (`RuntimeVisibleAnnotations`, `RuntimeInvisibleAnnotations`, `RuntimeVisibleParameterAnnotations`, and `RuntimeInvisibleParameterAnnotations`).

The `element_value` structure has the following format:

```
element_value {
    u1 tag;
    union {
        u2 const_value_index;

        {   u2 type_name_index;
            u2 const_name_index;
        } enum_const_value;

        u2 class_info_index;

        annotation annotation_value;

        {   u2              num_values;
            element_value values[num_values];
        } array_value;
    } value;
}
```

The items of the `element_value` structure are as follows:

tag

> The `tag` item indicates the type of this annotation element-value pair.
>
> The letters B, C, D, F, I, J, S, and Z indicate a primitive type. These letters are interpreted as if they were field descriptors (§4.3.2).
>
> The other legal values for `tag` are listed with their interpretations in Table 4.9.

**Table 4.9. Interpretation of additional tag values**

| tag Value | Element Type |
|---|---|
| s | String |
| e | enum constant |
| c | class |
| @ | annotation type |
| [ | array |

value

> The `value` item represents the value of this annotation element. This item is a union. The `tag` item, above, determines which item of the union is to be used:

> const_value_index

>> The `const_value_index` item is used if the `tag` item is one of B, C, D, F, I, J, S, Z, or s.
>>
>> The value of the `const_value_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be of the correct entry type for the field type designated by the `tag` item, as specified in Table 4.9.

> enum_const_value

>> The `enum_const_value` item is used if the `tag` item is e.

The `enum_const_value` item consists of the following two items:

type_name_index

> The value of the `type_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be
> a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing a valid field
> descriptor ([§4.3.2](#)) that denotes the internal form of the binary name
> ([§4.2.1](#)) of the type of the enum constant represented by this
> `element_value` structure.

const_name_index

> The value of the `const_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be
> a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing the simple
> name of the enum constant represented by this `element_value`
> structure.

class_info_index

> The `class_info_index` item is used if the `tag` item is c.
>
> The `class_info_index` item must be a valid index into the `constant_pool`
> table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info`
> ([§4.4.7](#)) structure representing the return descriptor ([§4.3.3](#)) of the type that is
> reified by the class represented by this `element_value` structure.
>
> *For example, V for Void.class, Ljava/lang/Object; for Object, etc.*

annotation_value

> The `annotation_value` item is used if the `tag` item is @.
>
> The `element_value` structure represents a "nested" annotation.

array_value

> The `array_value` item is used if the `tag` item is [.

The `array_value` item consists of the following two items:

num_values

> The value of the `num_values` item gives the number of elements in the array-typed value represented by this `element_value` structure.
>
> *Note that a maximum of 65535 elements are permitted in an array-typed element value.*

values

> Each value of the `values` table gives the value of an element of the array-typed value represented by this `element_value` structure.

### 4.7.17. The `RuntimeInvisibleAnnotations` attribute

The `RuntimeInvisibleAnnotations` attribute is similar to the `RuntimeVisibleAnnotations` attribute, except that the annotations represented by a `RuntimeInvisibleAnnotations` attribute must not be made available for return by reflective APIs, unless the Java Virtual Machine has been instructed to retain these annotations via some implementation-specific mechanism such as a command line flag. In the absence of such instructions, the Java Virtual Machine ignores this attribute.

The `RuntimeInvisibleAnnotations` attribute is a variable-length attribute in the `attributes` table of a `ClassFile`, `field_info`, or `method_info` structure ([§4.1](#), [§4.5](#), [§4.6](#)). The `RuntimeInvisibleAnnotations` attribute records run-time-invisible Java programming language annotations on the corresponding class, method, or field.

Each `ClassFile`, `field_info`, and `method_info` structure may contain at most one `RuntimeInvisibleAnnotations` attribute, which records all the run-time-invisible Java programming language annotations on the corresponding program element.

The `RuntimeInvisibleAnnotations` attribute has the following format:

```
RuntimeInvisibleAnnotations_attribute {
    u2          attribute_name_index;
    u4          attribute_length;
```

```
    u2          num_annotations;
    annotation annotations[num_annotations];
}
```

The items of the `RuntimeInvisibleAnnotations_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the
> `constant_pool` table. The `constant_pool` entry at that index must be a
> `CONSTANT_Utf8_info` ([§4.4.7](#)) structure representing the string
> "RuntimeInvisibleAnnotations".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute,
> excluding the initial six bytes.

> The value of the `attribute_length` item is thus dependent on the number of run-
> time-invisible annotations represented by the structure, and their values.

num_annotations

> The value of the `num_annotations` item gives the number of run-time-invisible
> annotations represented by the structure.

> *Note that a maximum of 65535 run-time-invisible Java programming language annotations may be*
> *directly attached to a program element.*

annotations

> Each value of the `annotations` table represents a single run-time-invisible annotation
> on a program element.

### 4.7.18. The `RuntimeVisibleParameterAnnotations` attribute

The `RuntimeVisibleParameterAnnotations` attribute is a variable-length attribute in the
`attributes` table of the `method_info` structure ([§4.6](#)). The
`RuntimeVisibleParameterAnnotations` attribute records run-time-visible Java

programming language annotations on the parameters of the corresponding method.

Each `method_info` structure may contain at most one `RuntimeVisibleParameterAnnotations` attribute, which records all the run-time-visible Java programming language annotations on the parameters of the corresponding method. The Java Virtual Machine must make these annotations available so they can be returned by the appropriate reflective APIs.

The `RuntimeVisibleParameterAnnotations` attribute has the following format:

```
RuntimeVisibleParameterAnnotations_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u1 num_parameters;
    {   u2          num_annotations;
        annotation annotations[num_annotations];
    } parameter_annotations[num_parameters];
}
```

The items of the `RuntimeVisibleParameterAnnotations_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing the string "RuntimeVisibleParameterAnnotations".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

> The value of the `attribute_length` item is thus dependent on the number of parameters, the number of run-time-visible annotations on each parameter, and their values.

num_parameters

The value of the `num_parameters` item gives the number of parameters of the method represented by the `method_info` structure on which the annotation occurs. (This duplicates information that could be extracted from the method descriptor ([§4.3.3](#)).)

parameter_annotations

Each value of the `parameter_annotations` table represents all of the run-time-visible annotations on a single parameter. The sequence of values in the table corresponds to the `sequence` of parameters in the method descriptor. Each `parameter_annotations` entry contains the following two items:

num_annotations

The value of the `num_annotations` item indicates the number of run-time-visible annotations on the parameter corresponding to the `sequence` number of this `parameter_annotations` element.

annotations

Each value of the `annotations` table represents a single run-time-visible annotation on the parameter corresponding to the `sequence` number of this `parameter_annotations` element.

### 4.7.19. The `RuntimeInvisibleParameterAnnotations` attribute

The `RuntimeInvisibleParameterAnnotations` attribute is similar to the `RuntimeVisibleParameterAnnotations` attribute, except that the annotations represented by a `RuntimeInvisibleParameterAnnotations` attribute must not be made available for return by reflective APIs, unless the Java Virtual Machine has specifically been instructed to retain these annotations via some implementation-specific mechanism such as a command line flag. In the absence of such instructions, the Java Virtual Machine ignores this attribute.

The `RuntimeInvisibleParameterAnnotations` attribute is a variable-length attribute in the `attributes` table of a `method_info` structure ([§4.6](#)). The `RuntimeInvisibleParameterAnnotations` attribute records run-time-invisible Java programming language annotations on the parameters of the corresponding method.

Each `method_info` structure may contain at most one

`RuntimeInvisibleParameterAnnotations` attribute, which records all the run-time-invisible Java programming language annotations on the parameters of the corresponding method.

The `RuntimeInvisibleParameterAnnotations` attribute has the following format:

```
RuntimeInvisibleParameterAnnotations_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u1 num_parameters;
    {   u2          num_annotations;
        annotation annotations[num_annotations];
    } parameter_annotations[num_parameters];
}
```

The items of the `RuntimeInvisibleParameterAnnotations_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing the string "RuntimeInvisibleParameterAnnotations".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

> The value of the `attribute_length` item is thus dependent on the number of parameters, the number of run-time-invisible annotations on each parameter, and their values.

num_parameters

> The value of the `num_parameters` item gives the number of parameters of the method represented by the `method_info` structure on which the annotation occurs. (This duplicates information that could be extracted from the method descriptor ([§4.3.3](#)).)

parameter_annotations

> Each value of the `parameter_annotations` table represents all of the run-time-invisible annotations on a single parameter. The sequence of values in the table corresponds to the sequence of parameters in the method descriptor. Each `parameter_annotations` entry contains the following two items:

> num_annotations

> > The value of the `num_annotations` item indicates the number of run-time-invisible annotations on the parameter corresponding to the sequence number of this `parameter_annotations` element.

> annotations

> > Each value of the `annotations` table represents a single run-time-invisible annotation on the parameter corresponding to the sequence number of this `parameter_annotations` element.

### 4.7.20. The `AnnotationDefault` attribute

The `AnnotationDefault` attribute is a variable-length attribute in the `attributes` table of certain `method_info` structures ([§4.6](#)), namely those representing elements of annotation types. The `AnnotationDefault` attribute records the default value for the element represented by the `method_info` structure.

Each `method_info` structure representing an element of an annotation type may contain at most one `AnnotationDefault` attribute. The Java Virtual Machine must make this default value available so it can be applied by appropriate reflective APIs.

The `AnnotationDefault` attribute has the following format:

```
AnnotationDefault_attribute {
    u2            attribute_name_index;
    u4            attribute_length;
    element_value default_value;
}
```

The items of the `AnnotationDefault_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing the string "AnnotationDefault".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.
>
> The value of the `attribute_length` item is thus dependent on the default value.

default_value

> The `default_value` item represents the default value of the annotation type element whose default value is represented by this `AnnotationDefault` attribute.

### 4.7.21. The `BootstrapMethods` attribute

The `BootstrapMethods` attribute is a variable-length attribute in the `attributes` table of a `ClassFile` structure ([§4.1](#)). The `BootstrapMethods` attribute records bootstrap method specifiers referenced by *invokedynamic* instructions ([§*invokedynamic*](#)).

There must be exactly one `BootstrapMethods` attribute in the `attributes` table of a given `ClassFile` structure if the `constant_pool` table of the `ClassFile` structure has at least one `CONSTANT_InvokeDynamic_info` entry ([§4.4.10](#)). There can be no more than one `BootstrapMethods` attribute in the `attributes` table of a given `ClassFile` structure.

The `BootstrapMethods` attribute has the following format:

```
BootstrapMethods_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
```

```
      u2 num_bootstrap_methods;
      {   u2 bootstrap_method_ref;
          u2 num_bootstrap_arguments;
          u2 bootstrap_arguments[num_bootstrap_arguments];
      } bootstrap_methods[num_bootstrap_methods];
}
```

The items of the `BootstrapMethods_attribute` structure are as follows:

attribute_name_index

> The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure ([§4.4.7](#)) representing the string "BootstrapMethods".

attribute_length

> The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

> The value of the `attribute_length` item is thus dependent on the number of *invokedynamic* instructions in this `ClassFile` structure.

num_bootstrap_methods

> The value of the `num_bootstrap_methods` item determines the number of bootstrap method specifiers in the `bootstrap_methods` array.

bootstrap_methods[]

> Each entry in the `bootstrap_methods` array contains an index to a `CONSTANT_MethodHandle_info` structure ([§4.4.8](#)) which specifies a bootstrap method, and a sequence (perhaps empty) of indexes to *static arguments* for the bootstrap method.

> Each `bootstrap_methods` entry must contain the following three items:

> bootstrap_method_ref

The value of the `bootstrap_method_ref` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_MethodHandle_info` structure ([§4.4.8](#)).

*The `reference_kind` item of the `CONSTANT_MethodHandle_info` structure should have the value 6 (`REF_invokeStatic`) or 8 (`REF_newInvokeSpecial`) ([§5.4.3.5](#)) or else invocation of the bootstrap method handle during call site specifier resolution for an invokedynamic instruction will complete abruptly.*

num_bootstrap_arguments

The value of the `num_bootstrap_arguments` item gives the number of items in the `bootstrap_arguments` array.

bootstrap_arguments

Each entry in the `bootstrap_arguments` array must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_String_info`, `CONSTANT_Class_info`, `CONSTANT_Integer_info`, `CONSTANT_Long_info`, `CONSTANT_Float_info`, `CONSTANT_Double_info`, `CONSTANT_MethodHandle_info`, or `CONSTANT_MethodType_info` structure ([§4.4.3](#), [§4.4.1](#), [§4.4.4](#), [§4.4.5](#)), [§4.4.8](#), [§4.4.9](#)).

# 4.8. Format Checking

When a prospective `class` file is loaded ([§5.3](#)) by the Java Virtual Machine, the Java Virtual Machine first ensures that the file has the basic format of a `class` file ([§4.1](#)). This process is known as *format checking*. The first four bytes must contain the right magic number. All recognized attributes must be of the proper length. The `class` file must not be truncated or have extra bytes at the end. The constant pool must not contain any superficially unrecognizable information.

This check for basic `class` file integrity is necessary for any interpretation of the `class` file contents.

Format checking is distinct from bytecode verification. Both are part of the verification process. Historically, format checking has been confused with bytecode verification, because both are a

form of integrity check.

# 4.9. Constraints on Java Virtual Machine code

The Java Virtual Machine code for a method, instance initialization method, or class or interface initialization method ([§2.9](#)) is stored in the `code` array of the Code attribute of a `method_info` structure of a `class` file ([§4.6](#), [§4.7.3](#)). This section describes the constraints associated with the contents of the `Code_attribute` structure.

## 4.9.1. Static Constraints

The *static constraints* on a `class` file are those defining the well-formedness of the file. With the exception of the static constraints on the Java Virtual Machine code of the `class` file, these constraints have been given in the previous sections. The static constraints on the Java Virtual Machine code in a `class` file specify how Java Virtual Machine instructions must be laid out in the `code` array and what the operands of individual instructions must be.

The static constraints on the instructions in the `code` array are as follows:

- The `code` array must not be empty, so the `code_length` item cannot have the value 0.

- The value of the `code_length` item must be less than 65536.

- The opcode of the first instruction in the `code` array begins at index 0.

- Only instances of the instructions documented in [§6.5](#) may appear in the `code` array. Instances of instructions using the reserved opcodes ([§6.2](#)) or any opcodes not documented in this specification must not appear in the `code` array.

- If the `class` file version number is 51.0 or above, then neither the *jsr* opcode or the *jsr_w* opcode may appear in the `code` array.

- For each instruction in the `code` array except the last, the index of the opcode of the next instruction equals the index of the opcode of the current instruction plus the length of that instruction, including all its operands.

  The *wide* instruction is treated like any other instruction for these purposes; the opcode

specifying the operation that a *wide* instruction is to modify is treated as one of the operands of that *wide* instruction. That opcode must never be directly reachable by the computation.

- The last byte of the last instruction in the code array must be the byte at index `code_length - 1`.

The static constraints on the operands of instructions in the code array are as follows:

- The target of each jump and branch instruction (*jsr*, *jsr_w*, *goto*, *goto_w*, *ifeq*, *ifne*, *ifle*, *iflt*, *ifge*, *ifgt*, *ifnull*, *ifnonnull*, *if_icmpeq*, *if_icmpne*, *if_icmple*, *if_icmplt*, *if_icmpge*, *if_icmpgt*, *if_acmpeq*, *if_acmpne*) must be the opcode of an instruction within this method.

  The target of a jump or branch instruction must never be the opcode used to specify the operation to be modified by a *wide* instruction; a jump or branch target may be the *wide* instruction itself.

- Each target, including the default, of each *tableswitch* instruction must be the opcode of an instruction within this method.

  Each *tableswitch* instruction must have a number of entries in its jump table that is consistent with the value of its *low* and *high* jump table operands, and its *low* value must be less than or equal to its *high* value.

  No target of a *tableswitch* instruction may be the opcode used to specify the operation to be modified by a *wide* instruction; a *tableswitch* target may be a *wide* instruction itself.

- Each target, including the default, of each *lookupswitch* instruction must be the opcode of an instruction within this method.

  Each *lookupswitch* instruction must have a number of *match-offset* pairs that is consistent with the value of its *npairs* operand. The *match-offset* pairs must be sorted in increasing numerical order by signed match value.

  No target of a *lookupswitch* instruction may be the opcode used to specify the operation to be modified by a *wide* instruction; a *lookupswitch* target may be a *wide* instruction itself.

- The operand of each *ldc* instruction and each *ldc_w* instruction must be a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type:

- CONSTANT_Integer, CONSTANT_Float, or CONSTANT_String if the class file version number is less than 49.0.

- CONSTANT_Integer, CONSTANT_Float, CONSTANT_String, or CONSTANT_Class if the class file version number is 49.0 or 50.0.

- CONSTANT_Integer, CONSTANT_Float, CONSTANT_String, CONSTANT_Class, CONSTANT_MethodType, or CONSTANT_MethodHandle if the class file version number is 51.0.

- The operands of each *ldc2_w* instruction must represent a valid index into the constant_pool table. The constant pool entry referenced by that index must be of type CONSTANT_Long or CONSTANT_Double.

  In addition, the subsequent constant pool index must also be a valid index into the constant pool, and the constant pool entry at that index must not be used.

- The operands of each *getfield*, *putfield*, *getstatic*, and *putstatic* instruction must represent a valid index into the constant_pool table. The constant pool entry referenced by that index must be of type CONSTANT_Fieldref.

- The *indexbyte* operands of each *invokevirtual*, *invokespecial*, and *invokestatic* instruction must represent a valid index into the constant_pool table. The constant pool entry referenced by that index must be of type CONSTANT_Methodref.

- The *indexbyte* operands of each *invokedynamic* instruction must represent a valid index into the constant_pool table. The constant pool entry referenced by that index must be of type CONSTANT_InvokeDynamic.

  The third and fourth operand bytes of each *invokedynamic* instruction must have the value zero.

- Only the *invokespecial* instruction is allowed to invoke an instance initialization method ([§2.9](#)).

  No other method whose name begins with the character '<' ('\u003c') may be called by the method invocation instructions. In particular, the class or interface initialization method specially named <clinit> is never called explicitly from Java Virtual Machine instructions, but only implicitly by the Java Virtual Machine itself.

- The *indexbyte* operands of each *invokeinterface* instruction must represent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_InterfaceMethodref`.

  The value of the *count* operand of each *invokeinterface* instruction must reflect the number of local variables necessary to store the arguments to be passed to the interface method, as implied by the descriptor of the `CONSTANT_NameAndType_info` structure referenced by the `CONSTANT_InterfaceMethodref` constant pool entry.

  The fourth operand byte of each *invokeinterface* instruction must have the value zero.

- The operands of each *instanceof*, *checkcast*, *new*, and *anewarray* instruction and the *indexbyte* operands of each *multianewarray* instruction must represent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_Class`.

- No *anewarray* instruction may be used to create an array of more than 255 dimensions.

- No *new* instruction may reference a `CONSTANT_Class` `constant_pool` table entry representing an array class. The *new* instruction cannot be used to create an array.

- A *multianewarray* instruction must be used only to create an array of a type that has at least as many dimensions as the value of its *dimensions* operand. That is, while a *multianewarray* instruction is not required to create all of the dimensions of the array type referenced by its *indexbyte* operands, it must not attempt to create more dimensions than are in the array type.

  The *dimensions* operand of each *multianewarray* instruction must not be zero.

- The *atype* operand of each *newarray* instruction must take one of the values T_BOOLEAN (4), T_CHAR (5), T_FLOAT (6), T_DOUBLE (7), T_BYTE (8), T_SHORT (9), T_INT (10), or T_LONG (11).

- The *index* operand of each *iload*, *fload*, *aload*, *istore*, *fstore*, *astore*, *iinc*, and *ret* instruction must be a non-negative integer no greater than `max_locals - 1`.

- The implicit index of each *iload_<n>*, *fload_<n>*, *aload_<n>*, *istore_<n>*, *fstore_<n>*, and *astore_<n>* instruction must be no greater than the value of `max_locals - 1`.

- The *index* operand of each *lload*, *dload*, *lstore*, and *dstore* instruction must be no greater

than the value of `max_locals - 2`.

- The implicit index of each *lload_<n>*, *dload_<n>*, *lstore_<n>*, and *dstore_<n>* instruction must be no greater than the value of `max_locals - 2`.

- The *indexbyte* operands of each *wide* instruction modifying an *iload*, *fload*, *aload*, *istore*, *fstore*, *astore*, *ret*, or *iinc* instruction must represent a non-negative integer no greater than `max_locals - 1`.

  The *indexbyte* operands of each *wide* instruction modifying an *lload*, *dload*, *lstore*, or *dstore* instruction must represent a non-negative integer no greater than `max_locals - 2`.

## 4.9.2. Structural Constraints

The structural constraints on the code array specify constraints on relationships between Java Virtual Machine instructions. The structural constraints are as follows:

- Each instruction must only be executed with the appropriate type and number of arguments in the operand stack and local variable array, regardless of the execution path that leads to its invocation.

  An instruction operating on values of type `int` is also permitted to operate on values of type `boolean`, `byte`, `char`, and `short`. (As noted in [§2.3.4](#) and [§2.11.1](#), the Java Virtual Machine internally converts values of types `boolean`, `byte`, `char`, and `short` to type `int`.)

- If an instruction can be executed along several different execution paths, the operand stack must have the same depth ([§2.6.2](#)) prior to the execution of the instruction, regardless of the path taken.

- At no point during execution can the order of the local variable pair holding a value of type `long` or `double` be reversed or the pair split up.

  At no point can the local variables of such a pair be operated on individually.

- No local variable (or local variable pair, in the case of a value of type `long` or `double`) can be accessed before it is assigned a value.

- At no point during execution can the operand stack grow to a depth ([§2.6.2](#)) greater than

that implied by the `max_stack` item.

- At no point during execution can more values be popped from the operand stack than it contains.

- Each *invokespecial* instruction must name an instance initialization method ([§2.9](#)), a method in the current class, or a method in a superclass of the current class.

  If an *invokespecial* instruction names an instance initialization method from a class that is not the current class or a superclass, and the target reference on the operand stack is a class instance created by an earlier *new* instruction, then *invokespecial* must name an instance initialization method from the class of that class instance.

- When the instance initialization method ([§2.9](#)) is invoked, an uninitialized class instance must be in an appropriate position on the operand stack.

  An instance initialization method must never be invoked on an initialized class instance.

- When any instance method is invoked or when any instance variable is accessed, the class instance that contains the instance method or instance variable must already be initialized.

- There must never be an uninitialized class instance on the operand stack or in a local variable at the target of a backwards branch unless the special type of the uninitialized class instance at the branch instruction is merged with itself at the target of the branch ([§4.10.2.4](#)).

- There must never be an uninitialized class instance in a local variable in code protected by an exception handler ([§4.10.2.4](#)).

- There must never be an uninitialized class instance on the operand stack or in a local variable when a *jsr* or *jsr_w* instruction is executed.

- Each instance initialization method ([§2.9](#)), except for the instance initialization method derived from the constructor of class `Object`, must call either another instance initialization method of `this` or an instance initialization method of its direct superclass `super` before its instance members are accessed.

  However, instance fields of `this` that are declared in the current class may be assigned before calling any instance initialization method.

- The arguments to each method invocation must be method invocation compatible (JLS §5.3) with the method descriptor ([§4.3.3](#)).

- The type of every class instance that is the target of a method invocation instruction must be assignment compatible (JLS §5.2) with the class or interface type specified in the instruction.

  In addition, the type of the target of an *invokespecial* instruction must be assignment compatible with the current class, unless an instance initialization method is being invoked.

- Each return instruction must match its method's return type:

  - If the method returns a `boolean`, `byte`, `char`, `short`, or `int`, only the *ireturn* instruction may be used.

  - If the method returns a `float`, `long`, or `double`, only an *freturn*, *lreturn*, or *dreturn* instruction, respectively, may be used.

  - If the method returns a `reference` type, it must do so using an *areturn* instruction, and the type of the returned value must be assignment compatible (JLS §5.2) with the return descriptor ([§4.3.3](#)) of the method.

  - All instance initialization methods, class or interface initialization methods, and methods declared to return `void` must use only the *return* instruction.

- If *getfield* or *putfield* is used to access a `protected` field declared in a superclass that is a member of a different run-time package than the current class, then the type of the class instance being accessed must be the same as or a subclass of the current class.

  If *invokevirtual* or *invokespecial* is used to access a `protected` method declared in a superclass that is a member of a different run-time package than the current class, then the type of the class instance being accessed must be the same as or a subclass of the current class.

- The type of every class instance accessed by a *getfield* instruction or modified by a *putfield* instruction must be assignment compatible (JLS §5.2) with the class type specified in the instruction.

- The type of every value stored by a *putfield* or *putstatic* instruction must be compatible with the descriptor of the field ([§4.3.2](#)) of the class instance or class being stored into:

- If the descriptor type is `boolean`, `byte`, `char`, `short`, or `int`, then the value must be an `int`.

- If the descriptor type is `float`, `long`, or `double`, then the value must be a `float`, `long`, or `double`, respectively.

- If the descriptor type is a `reference` type, then the value must be of a type that is assignment compatible (JLS §5.2) with the descriptor type.

- The type of every value stored into an array by an *aastore* instruction must be a `reference` type.

  The component type of the array being stored into by the *aastore* instruction must also be a `reference` type.

- Each *athrow* instruction must throw only values that are instances of class `Throwable` or of subclasses of `Throwable`.

  Each class mentioned in a `catch_type` item of a `method`'s exception table must be `Throwable` or a subclass of `Throwable`.

- Execution never falls off the bottom of the `code` array.

- No return address (a value of type `returnAddress`) may be loaded from a local variable.

- The instruction following each *jsr* or *jsr_w* instruction may be returned to only by a single *ret* instruction.

- No *jsr* or *jsr_w* instruction that is returned to may be used to recursively call a subroutine if that subroutine is already present in the subroutine call chain. (Subroutines can be nested when using `try-finally` constructs from within a `finally` clause.)

- Each instance of type `returnAddress` can be returned to at most once.

  If a *ret* instruction returns to a point in the subroutine call chain above the *ret* instruction corresponding to a given instance of type `returnAddress`, then that instance can never be used as a return address.

# 4.10. Verification of `class` Files

Even though a compiler for the Java programming language must only produce `class` files that satisfy all the static and structural constraints in the previous sections, the Java Virtual Machine has no guarantee that any file it is asked to load was generated by that compiler or is properly formed. Applications such as web browsers do not download source code, which they then compile; these applications download already-compiled `class` files. The browser needs to determine whether the `class` file was produced by a trustworthy compiler or by an adversary attempting to exploit the Java Virtual Machine.

*An additional problem with compile-time checking is version skew. A user may have successfully compiled a class, say `PurchaseStockOptions`, to be a subclass of `TradingClass`. But the definition of `TradingClass` might have changed since the time the class was compiled in a way that is not compatible with pre-existing binaries. Methods might have been deleted or had their return types or modifiers changed. Fields might have changed types or changed from instance variables to class variables. The access modifiers of a method or variable may have changed from `public` to `private`. For a discussion of these issues, see Chapter 13, "Binary Compatibility," in The Java Language Specification, Java SE 7 Edition.*

Because of these potential problems, the Java Virtual Machine needs to verify for itself that the desired constraints are satisfied by the `class` files it attempts to incorporate. A Java Virtual Machine implementation verifies that each `class` file satisfies the necessary constraints at linking time ([§5.4](#)).

Linking-time verification enhances the performance of the interpreter. Expensive checks that would otherwise have to be performed to verify constraints at run time for each interpreted instruction can be eliminated. The Java Virtual Machine can assume that these checks have already been performed. For example, the Java Virtual Machine will already know the following:

- There are no operand stack overflows or underflows.

- All local variable uses and stores are valid.

- The arguments to all the Java Virtual Machine instructions are of valid types.

The verifier also performs verification that can be done without looking at the `code` array of the Code attribute ([§4.7.3](#)). The checks performed include the following:

- Ensuring that `final` classes are not subclassed and that `final` methods are not overridden ([§5.4.5](#)).

- Checking that every class (except `Object`) has a direct superclass.

- Ensuring that the constant pool satisfies the documented static constraints; for example, that each `CONSTANT_Class_info` structure in the constant pool contains in its `name_index` item a valid constant pool index for a `CONSTANT_Utf8_info` structure.

- Checking that all field references and method references in the constant pool have valid names, valid classes, and a valid type descriptor.

Note that these checks do not ensure that the given field or method actually exists in the given class, nor do they check that the type descriptors given refer to real classes. They ensure only that these items are well formed. More detailed checking is performed when the bytecodes themselves are verified, and during resolution.

There are two strategies that Java Virtual Machine implementations may use for verification:

- Verification by type checking must be used to verify `class` files whose version number is greater than or equal to 50.0.

- Verification by type inference must be supported by all Java Virtual Machine implementations, except those conforming to the Java ME CLDC and Java Card profiles, in order to verify `class` files whose version number is less than 50.0.

  Verification on Java Virtual Machine implementations supporting the Java ME CLDC and Java Card profiles is governed by their respective specifications.

### 4.10.1. Verification by Type Checking

A `class` file whose version number is greater than or equal to 50.0 ([§4.1](#)) must be verified using the type checking rules given in this section.

If, and only if, a `class` file's version number equals 50.0, then if the type checking fails, a Java Virtual Machine implementation may choose to attempt to perform verification by type inference ([§4.10.2](#)).

*This is a pragmatic adjustment, designed to ease the transition to the new verification discipline. Many tools that manipulate `class` files may alter the bytecodes of a method in a manner that requires adjustment of the method's stack map frames. If a tool does not make the necessary adjustments to the stack map frames, type checking may fail even though the bytecode is in principle valid (and would consequently verify under the*

*old type inference scheme). To allow implementors time to adapt their tools, Java Virtual Machine implementations may fall back to the older verification discipline, but only for a limited time.*

*In cases where type checking fails but type inference is invoked and succeeds, a certain performance penalty is expected. Such a penalty is unavoidable. It also should serve as a signal to tool vendors that their output needs to be adjusted, and provides vendors with additional incentive to make these adjustments.*

*In summary, failover to verification by type inference supports both the gradual addition of stack map frames to the Java SE platform (if they are not present in a version 50.0 `class` file, failover is allowed) and the gradual removal of the jsr and jsr_w instructions from the Java SE platform (if they are present in a version 50.0 `class` file, failover is allowed).*

If a Java Virtual Machine implementation ever attempts to perform verification by type inference on version 50.0 class files, it must do so in all cases where verification by type checking fails.

*This means that a Java Virtual Machine implementation cannot choose to resort to type inference in once case and not in another. It must either reject `class` files that do not verify via type checking, or else consistently failover to the type inferencing verifier whenever type checking fails.*

The type checker enforces type rules that are specified by means of Prolog clauses. English language text is used to describe the type rules in an informal way, while the Prolog clauses provide a formal specification.

The type checker requires a list of stack map frames for each method with a Code attribute ([§4.7.3](#)). A list of stack map frames is given by the `StackMapTable` attribute ([§4.7.4](#)) of a Code attribute. The intent is that a stack map frame must appear at the beginning of each basic block in a method. The stack map frame specifies the verification type of each operand stack entry and of each local variable at the start of each basic block. The type checker reads the stack map frames for each method with a Code attribute and uses these maps to generate a proof of the type safety of the instructions in the Code attribute.

A class is type safe if all its methods are type safe, and it does not subclass a `final` class.

```
classIsTypeSafe(Class) :-
    classClassName(Class, Name),
    classDefiningLoader(Class, L),
    superclassChain(Name, L, Chain),
    Chain \= [],
    classSuperClassName(Class, SuperclassName),
    loadedClass(SuperclassName, L, Superclass),
```

```
        classIsNotFinal(Superclass),
        classMethods(Class, Methods),
        checklist(methodIsTypeSafe(Class), Methods).
```

```
classIsTypeSafe(Class) :-
        classClassName(Class, 'java/lang/Object'),
        classDefiningLoader(Class, L),
        isBootstrapLoader(L),
        classMethods(Class, Methods),
        checklist(methodIsTypeSafe(Class), Methods).
```

The Prolog predicate `classIsTypeSafe` assumes that `Class` is a Prolog term representing a binary class that has been successfully parsed and loaded. This specification does not mandate the precise structure of this term, but does require that certain predicates be defined upon it.

*For example, we assume a predicate `classMethods(Class, Methods)` that, given a term representing a class as described above as its first argument, binds its second argument to a list comprising all the methods of the class, represented in a convenient form described later.*

**Iff the predicate `classIsTypeSafe` is not true, the type checker must throw the exception `VerifyError` to indicate that the `class` file is malformed. Otherwise, the `class` file has type checked successfully and bytecode verification has completed successfully.**

The rest of this section explains the process of type checking in detail:

- First, we give Prolog predicates for core Java Virtual Machine artifacts like classes and methods ([§4.10.1.1](#)).

- Second, we specify the type system known to the type checker ([§4.10.1.2](#)).

- Third, we specify the Prolog representation of instructions and stack map frames ([§4.10.1.3](#), [§4.10.1.4](#)).

- Fourth, we specify how a method is type checked, for methods without code ([§4.10.1.5](#)) and methods with code ([§4.10.1.6](#)).

- Fifth, we discuss type checking issues common to all load and store instructions

([§4.10.1.7](#)), and also issues of access to `protected` members ([§4.10.1.8](#)).

- Finally, we specify the rules to type check each instruction ([§4.10.1.9](#)).

### 4.10.1.1. Accessors for Java Virtual Machine Artifacts

We stipulate the existence of 22 Prolog predicates ("accessors") that have certain expected behavior but whose formal definitions are not given in this specification.

classClassName(Class, ClassName)

       Extracts the name, `ClassName`, of the class `Class`.

classIsInterface(Class)

       True iff the class, `Class`, is an interface.

classIsNotFinal(Class)

       True iff the class, `Class`, is not a `final` class.

classSuperClassName(Class, SuperClassName)

       Extracts the name, `SuperClassName`, of the superclass of class `Class`.

classInterfaces(Class, Interfaces)

       Extracts a list, `Interfaces`, of the direct superinterfaces of the class `Class`.

classMethods(Class, Methods)

       Extracts a list, `Methods`, of the methods declared in the class `Class`.

classAttributes(Class, Attributes)

       Extracts a list, `Attributes`, of the attributes of the class `Class`.

       Each attribute is represented as a functor application of the form `attribute(AttributeName, AttributeContents)`, where `AttributeName` is the name of the attribute. The format of the attribute's contents is unspecified.

classDefiningLoader(Class, Loader)

> Extracts the defining class loader, `Loader`, of the class `Class`.

isBootstrapLoader(Loader)

> True iff the class loader `Loader` is the bootstrap class loader.

loadedClass(Name, InitiatingLoader, ClassDefinition)

> True iff there exists a class named `Name` whose representation (in accordance with this specification) when loaded by the class loader `InitiatingLoader` is `ClassDefinition`.

methodName(Method, Name)

> Extracts the name, `Name`, of the method `Method`.

methodAccessFlags(Method, AccessFlags)

> Extracts the access flags, `AccessFlags`, of the method `Method`.

methodDescriptor(Method, Descriptor)

> Extracts the descriptor, `Descriptor`, of the method `Method`.

methodAttributes(Method, Attributes)

> Extracts a list, `Attributes`, of the attributes of the method `Method`.

isNotFinal(Method, Class)

> True iff `Method` in class `Class` is not `final`.

isProtected(MemberClass, MemberName, MemberDescriptor)

> True iff there is a member named `MemberName` with descriptor `MemberDescriptor` in the class `MemberClass` and it is `protected`.

isNotProtected(MemberClass, MemberName, MemberDescriptor)

> True iff there is a member named `MemberName` with descriptor `MemberDescriptor` in

the class `MemberClass` and it is not `protected`.

parseFieldDescriptor(Descriptor, Type)

> Converts a field descriptor, `Descriptor`, into the corresponding verification type Type ([§4.10.1.2](#)).

parseMethodDescriptor(Descriptor, ArgTypeList, ReturnType)

> Converts a method descriptor, `Descriptor`, into a list of verification types, `ArgTypeList`, corresponding to the method argument types, and a verification type, `ReturnType`, corresponding to the return type.

parseCodeAttribute(Class, Method, FrameSize, MaxStack, ParsedCode, Handlers, StackMap)

> Extracts the instruction stream, `ParsedCode`, of the method `Method` in `Class`, as well as the maximum operand stack size, `MaxStack`, the maximal number of local variables, `FrameSize`, the exception handlers, `Handlers`, and the stack map `StackMap`.

> The representation of the instruction stream and stack map attribute must be as specified in [§4.10.1.3](#) and [§4.10.1.4](#).

samePackageName(Class1, Class2)

> True iff the package names of `Class1` and `Class2` are the same.

differentPackageName(Class1, Class2)

> True iff the package names of `Class1` and `Class2` are different.

When type checking a method's body, it is convenient to access information about the method. For this purpose, we define an *environment*, a six-tuple consisting of:

- a class

- a method

- the declared return type of the method

- the instructions in a method

- the maximal size of the operand stack

- a list of exception handlers

We specify accessors to extract information from the environment.

```prolog
allInstructions(Environment, Instructions) :-
    Environment = environment(_Class, _Method, _ReturnType,
                             Instructions, _, _).

exceptionHandlers(Environment, Handlers) :-
    Environment = environment(_Class, _Method, _ReturnType,
                             _Instructions, _, Handlers).

maxOperandStackLength(Environment, MaxStack) :-
    Environment = environment(_Class, _Method, _ReturnType,
                             _Instructions, MaxStack, _Handlers).

thisClass(Environment, class(ClassName, L)) :-
    Environment = environment(Class, _Method, _ReturnType,
                             _Instructions, _, _),
    classDefiningLoader(Class, L),
    classClassName(Class, ClassName).

thisMethodReturnType(Environment, ReturnType) :-
    Environment = environment(_Class, _Method, ReturnType,
                             _Instructions, _, _).
```

We specify additional predicates to extract higher-level information from the environment.

```prolog
offsetStackFrame(Environment, Offset, StackFrame) :-
    allInstructions(Environment, Instructions),
    member(stackMap(Offset, StackFrame), Instructions).

currentClassLoader(Environment, Loader) :-
    thisClass(Environment, class(_, Loader)).
```

Finally, we specify a general predicate used throughout the type rules:

```
notMember(_, []).
notMember(X, [A | More]) :- X \= A, notMember(X, More).
```

*The principle guiding the determination as to which accessors are stipulated and which are fully specified is that we do not want to over-specify the representation of the* class *file. Providing specific accessors to the* Class *or* Method *term would force us to completely specify the format for a Prolog term representing the* class *file.*

### 4.10.1.2. Verification Type System

The type checker enforces a type system based upon a hierarchy of *verification types*, illustrated below.

Most verification types have a direct correspondence with the types represented by field descriptors ([§4.3.2](#)) in [Table 4.2](#). The only exceptions are the field descriptors B, C, S, and Z, all of which correspond to the verification type `int`.

```
Verification type hierarchy:

                             top
             _____/_____
            /                                        \
           /                                          \
        oneWord                                     twoWord
       /   |   \                                    /      \
      /    |    \                                  /        \
    int  float  reference                        long       double
                   /     \
                  /       _____
                 /                    \
                /                      \
         uninitialized                Object
           /        \                      \
          /          \                      \
```

```
uninitializedThis  uninitialized(offset)      +-----------------+
                                              |  Java reference |
                                              |  type hierarchy |
                                              +-----------------+
                                                      |
                                                      |
                                                    null
```

Subtyping is reflexive.

```
isAssignable(X, X).
```

The verification types which are not reference types in the Java programming language have subtype rules of the form:

```
isAssignable(v, X) :- isAssignable(the_direct_supertype_of_v, X).
```

That is, v is a subtype of X if the direct supertype of v is a subtype of X. The rules are:

```
isAssignable(oneWord, top).
isAssignable(twoWord, top).

isAssignable(int, X)    :- isAssignable(oneWord, X).
isAssignable(float, X)  :- isAssignable(oneWord, X).
isAssignable(long, X)   :- isAssignable(twoWord, X).
isAssignable(double, X) :- isAssignable(twoWord, X).

isAssignable(reference, X)   :- isAssignable(oneWord, X).
isAssignable(class(_, _), X) :- isAssignable(reference, X).
isAssignable(arrayOf(_), X)  :- isAssignable(reference, X).

isAssignable(uninitialized, X)     :- isAssignable(reference, X).
isAssignable(uninitializedThis, X) :- isAssignable(uninitialized, X).
isAssignable(uninitialized(_), X)  :- isAssignable(uninitialized, X).

isAssignable(null, class(_, _)).
```

```
isAssignable(null, arrayOf(_)).
isAssignable(null, X) :- isAssignable(class('java/lang/Object', BL), X),
                         isBootstrapLoader(BL).
```

*These subtype rules are not necessarily the most obvious formulation of subtyping. There is a clear split between subtyping rules for reference types in the Java programming language, and rules for the remaining verification types. The split allows us to state general subtyping relations between Java programming language reference types and other verification types. These relations hold independently of a Java reference type's position in the type hierarchy, and help to prevent excessive class loading by a Java Virtual Machine implementation. For example, we do not want to start climbing the Java superclass hierarchy in response to a query of the form* `class(foo, L) <: twoWord`.

*We also have a rule that says subtyping is reflexive, so together these rules cover most verification types that are not reference types in the Java programming language.*

Subtype rules for the reference types in the Java programming language are specified recursively with `isJavaAssignable`.

```
isAssignable(class(X, Lx), class(Y, Ly)) :-
    isJavaAssignable(class(X, Lx), class(Y, Ly)).

isAssignable(arrayOf(X), class(Y, L)) :-
    isJavaAssignable(arrayOf(X), class(Y, L)).

isAssignable(arrayOf(X), arrayOf(Y)) :-
    isJavaAssignable(arrayOf(X), arrayOf(Y)).
```

For assignments, interfaces are treated like `Object`.

```
isJavaAssignable(class(_, _), class(To, L)) :-
    loadedClass(To, L, ToClass),
    classIsInterface(ToClass).

isJavaAssignable(From, To) :-
    isJavaSubclassOf(From, To).
```

Array types are subtypes of `Object`. The intent is also that array types are subtypes of

Cloneable and `java.io.Serializable`.

```
isJavaAssignable(arrayOf(_), class('java/lang/Object', BL)) :-
    isBootstrapLoader(BL).

isJavaAssignable(arrayOf(_), X) :-
    isArrayInterface(X).

isArrayInterface(class('java/lang/Cloneable', BL)) :-
    isBootstrapLoader(BL).

isArrayInterface(class('java/io/Serializable', BL)) :-
    isBootstrapLoader(BL).
```

Subtyping between arrays of primitive type is the identity relation.

```
isJavaAssignable(arrayOf(X), arrayOf(Y)) :-
    atom(X),
    atom(Y),
    X = Y.
```

Subtyping between arrays of reference type is covariant.

```
isJavaAssignable(arrayOf(X), arrayOf(Y)) :-
    compound(X), compound(Y), isJavaAssignable(X, Y).
```

Subclassing is reflexive.

```
isJavaSubclassOf(class(SubclassName, L), class(SubclassName, L)).
```

```
isJavaSubclassOf(class(SubclassName, LSub), class(SuperclassName, LSuper)) :-
    superclassChain(SubclassName, LSub, Chain),
    member(class(SuperclassName, L), Chain),
    loadedClass(SuperclassName, L, Sup),
```

```
        loadedClass(SuperclassName, LSuper, Sup).

superclassChain(ClassName, L, [class(SuperclassName, Ls) | Rest]) :-
    loadedClass(ClassName, L, Class),
    classSuperClassName(Class, SuperclassName),
    classDefiningLoader(Class, Ls),
    superclassChain(SuperclassName, Ls, Rest).

superclassChain('java/lang/Object', L, []) :-
    loadedClass('java/lang/Object', L, Class),
    classDefiningLoader(Class, BL),
    isBootstrapLoader(BL).
```

### 4.10.1.3. Instruction Representation

Individual bytecode instructions are represented in Prolog as terms whose functor is the name of the instruction and whose arguments are its parsed operands.

*For example, an aload instruction is represented as the term* `aload(N)`, *which includes the index* `N` *that is the operand of the instruction.*

The instructions as a whole are represented as a list of terms of the form:

```
instruction(Offset, AnInstruction)
```

*For example,* `instruction(21, aload(1))`.

The order of instructions in this list must be the same as in the `class` file.

A few instructions have operands that are constant pool entries representing fields, methods, and dynamic call sites. In the constant pool, a field is represented by a `CONSTANT_Fieldref_info` structure, a method is represented by a `CONSTANT_InterfaceMethodref_info` structure (for an interface's method) or a `CONSTANT_Methodref_info` structure (for a class's method), and a dynamic call site is represented by a `CONSTANT_InvokeDynamic_info` structure ([§4.4.2](#), [§4.4.10](#)). Such structures are represented as functor applications of the form:

- `field(FieldClassName, FieldName, FieldDescriptor)` for a field, where `FieldClassName` is the name of the class referenced by the `class_index` item in the `CONSTANT_Fieldref_info` structure, and `FieldName` and `FieldDescriptor` correspond to the name and field descriptor referenced by the `name_and_type_index` item of the `CONSTANT_Fieldref_info` structure.

- `imethod(MethodIntfName, MethodName, MethodDescriptor)` for an interface's method, where `MethodIntfName` is the name of the interface referenced by the `class_index` item of the `CONSTANT_InterfaceMethodref_info` structure, and `MethodName` and `MethodDescriptor` correspond to the name and method descriptor referenced by the `name_and_type_index` item of the `CONSTANT_InterfaceMethodref_info` structure;

- `method(MethodClassName, MethodName, MethodDescriptor)` for a class's method, where `MethodClassName` is the name of the class referenced by the `class_index` item of the `CONSTANT_Methodref_info` structure, and `MethodName` and `MethodDescriptor` correspond to the name and method descriptor referenced by the `name_and_type_index` item of the `CONSTANT_Methodref_info` structure; and

- `dmethod(CallSiteName, MethodDescriptor)` for a dynamic call site, where `CallSiteName` and `MethodDescriptor` correspond to the name and method descriptor referenced by the `name_and_type_index` item of the `CONSTANT_InvokeDynamic_info` structure.

For clarity, we assume that field and method descriptors ([§4.3.2](#)) are mapped into more readable names: the leading `L` and trailing `;` are dropped from class names, and the *BaseType* characters used for primitive types are mapped to the names of those types.

*For example, a getfield instruction whose operand was an index into the constant pool that refers to a field* `foo` *of type* `F` *in class* `Bar` *would be represented as* `getfield(field('Bar', 'foo', 'F'))`.

Constant pool entries that refer to constant values, such as `CONSTANT_String`, `CONSTANT_Integer`, `CONSTANT_Float`, `CONSTANT_Long`, `CONSTANT_Double`, and `CONSTANT_Class`, are encoded via the functors whose names are `string`, `int`, `float`, `long`, `double`, and `classConstant` respectively.

*For example, an ldc instruction for loading the integer 91 would be encoded as* `ldc(int(91))`.

### 4.10.1.4. Stack Map Frame Representation

Stack map frames are represented in Prolog as a list of terms of the form:

```
stackMap(Offset, TypeState)
```

where `Offset` is an integer indicating the offset of the instruction the frame map applies to, and `TypeState` is the expected incoming type state (§4.7.4) for that instruction.

The order of stack map frames in this list must be the same as in the `class` file.

`TypeState` has the form:

```
frame(Locals, OperandStack, Flags)
```

where:

- `Locals` is a list of verification types, such that the Nth element of the list (with 0-based indexing) represents the type of local variable N.

  If any local variable in `Locals` has the type `uninitializedThis`, then `Flags` has the single element `flagThisUninit`, otherwise it is an empty list.

- `OperandStack` is a list of types, such that the first element represents the type of the top of the operand stack, and the elements below the top follow in the appropriate order.

  Types of size 2 (`long` and `double`) are represented by two entries, with the first entry being `top` and the second one being the type itself.

  *For example, a stack with a `double`, an `int`, and a `long` would be represented as `[top, double, int, top, long]`.*

  Reference types other than array types are represented using the functor `class`. `class(N, L)` represents the class whose binary name is `N` as loaded by the loader `L`. Note that `L` is an initiating loader (§5.3) of the class represented by `class(N, L)` and may, or may not, be the class's defining loader.

  Array types are represented by applying the functor `arrayOf` to an argument denoting the

component type of the array.

The verification type `uninitialized(offset)` is represented by applying the functor `uninitialized` to an argument representing the numerical value of the `offset`.

Other verification types are represented in prolog as atoms whose name denotes the verification type in question.

*The class `Object` would be represented as `class('java/lang/Object', BL)`, where BL is the bootstrap loader.*

*The types `int[]` and `Object[]` would be represented by `arrayOf(int)` and `arrayOf(class('java/lang/Object', BL))` respectively.*

- `Flags` is a list which may either be empty or have the single element `flagThisUninit`.

  This flag is used in constructors, to mark type states where initialization of this has not yet been completed. In such type states, it is illegal to return from the method.

Subtyping ([§4.10.1.2](#)) is extended pointwise to type states.

The local variable array of a method has a fixed length by construction (see `methodInitialStackFrame` in [§4.10.1.6](#)) while the operand stack grows and shrinks. Therefore, we require an explicit check on the length of the operand stacks whose assignability is desired.

```
frameIsAssignable(frame(Locals1, StackMap1, Flags1),
                  frame(Locals2, StackMap2, Flags2)) :-
    length(StackMap1, StackMapLength),
    length(StackMap2, StackMapLength),
    maplist(isAssignable, Locals1, Locals2),
    maplist(isAssignable, StackMap1, StackMap2),
    subset(Flags1, Flags2).
```

The length of the operand stack must not exceed the declared maximum stack length.

```
operandStackHasLegalLength(Environment, OperandStack) :-
    length(OperandStack, Length),
    maxOperandStackLength(Environment, MaxStack),
```

```
        Length =< MaxStack.
```

Certain array instructions (§*aaload*, §*arraylength*, §*baload*, §*bastore*) peek at the types of values on the operand stack in order to check they are array types. The following clause accesses the I'th element of the operand stack from a type state.

```
nth1OperandStackIs(I, frame(_Locals, OperandStack, _Flags), Element) :-
    nth1(I, OperandStack, Element).
```

Manipulation of the operand stack by load and store instructions (§4.10.1.7) is complicated by the fact that some types occupy two entries on the stack. The predicates given below take this into account, allowing the rest of the specification to abstract from this issue.

Pop a list of types off the stack.

```
canPop(frame(Locals, OperandStack, Flags), Types,
       frame(Locals, PoppedOperandStack, Flags)) :-
    popMatchingList(OperandStack, Types, PoppedOperandStack).

popMatchingList(OperandStack, [], OperandStack).
popMatchingList(OperandStack, [P | Rest], NewOperandStack) :-
    popMatchingType(OperandStack, P, TempOperandStack, _ActualType),
    popMatchingList(TempOperandStack, Rest, NewOperandStack).
```

Pop an individual type off the stack. More precisely, if the logical top of the stack is some subtype of the specified type, Type, then pop it. If a type occupies two stack slots, the logical top of stack type is really the type just below the top, and the top of stack is the unusable type `top`.

```
popMatchingType([ActualType | OperandStack],
                Type, OperandStack, ActualType) :-
    sizeOf(Type, 1),
    isAssignable(ActualType, Type).

popMatchingType([top, ActualType | OperandStack],
                Type, OperandStack, ActualType) :-
```

```
        sizeOf(Type, 2),
        isAssignable(ActualType, Type).

sizeOf(X, 2) :- isAssignable(X, twoWord).
sizeOf(X, 1) :- isAssignable(X, oneWord).
sizeOf(top, 1).
```

Push a logical type onto the stack. The exact behavior varies with the size of the type. If the pushed type is of size 1, we just push it onto the stack. If the pushed type is of size 2, we push it, and then push `top`.

```
pushOperandStack(OperandStack, 'void', OperandStack).
pushOperandStack(OperandStack, Type, [Type | OperandStack]) :-
    sizeOf(Type, 1).
pushOperandStack(OperandStack, Type, [top, Type | OperandStack]) :-
    sizeOf(Type, 2).
```

Push a list of types onto the stack if there is space.

```
canSafelyPush(Environment, InputOperandStack, Type, OutputOperandStack) :-
    pushOperandStack(InputOperandStack, Type, OutputOperandStack),
    operandStackHasLegalLength(Environment, OutputOperandStack).

canSafelyPushList(Environment, InputOperandStack, Types,
                  OutputOperandStack) :-
    canPushList(InputOperandStack, Types, OutputOperandStack),
    operandStackHasLegalLength(Environment, OutputOperandStack).

canPushList(InputOperandStack, [], InputOperandStack).
canPushList(InputOperandStack, [Type | Rest], OutputOperandStack) :-
    pushOperandStack(InputOperandStack, Type, InterimOperandStack),
    canPushList(InterimOperandStack, Rest, OutputOperandStack).
```

Manipulation of the operand stack by the *dup* instructions is specified entirely in terms of the *category* of types for values on the stack ([§2.11.1](#)).

Category 1 types occupy a single stack slot. Popping a logical type of category 1, Type, off the stack is possible if the top of the stack is Type and Type is not top (otherwise it could denote the upper half of a category 2 type). The result is the incoming stack, with the top slot popped off.

```
popCategory1([Type | Rest], Type, Rest) :-
    Type \= top,
    sizeOf(Type, 1).
```

Category 2 types occupy two stack slots. Popping a logical type of category 2, Type, off the stack is possible if the top of the stack is type top, and the slot directly below it is Type. The result is the incoming stack, with the top 2 slots popped off.

```
popCategory2([top, Type | Rest], Type, Rest) :-
    sizeOf(Type, 2).
```

Most of the type rules for individual instructions ([§4.10.1.9](#)) depend on the notion of a valid *type transition*. A type transition is *valid* if one can pop a list of expected types off the incoming type state's operand stack and replace them with an expected result type, resulting in a new valid type state. In particular, the size of the operand stack in the new type state must not exceed its maximum declared size.

```
validTypeTransition(Environment, ExpectedTypesOnStack, ResultType,
                    frame(Locals, InputOperandStack, Flags),
                    frame(Locals, NextOperandStack, Flags)) :-
    popMatchingList(InputOperandStack, ExpectedTypesOnStack,
                    InterimOperandStack),
    pushOperandStack(InterimOperandStack, ResultType, NextOperandStack),
    operandStackHasLegalLength(Environment, NextOperandStack).
```

### 4.10.1.5. Type Checking Abstract and Native Methods

Abstract methods and native methods are considered to be type safe if they do not override a final method.

```
methodIsTypeSafe(Class, Method) :-
    doesNotOverrideFinalMethod(Class, Method),
    methodAccessFlags(Method, AccessFlags),
    member(abstract, AccessFlags).

methodIsTypeSafe(Class, Method) :-
    doesNotOverrideFinalMethod(Class, Method),
    methodAccessFlags(Method, AccessFlags),
    member(native, AccessFlags).

doesNotOverrideFinalMethod(class('java/lang/Object', L), Method) :-
    isBootstrapLoader(L).

doesNotOverrideFinalMethod(Class, Method) :-
    classSuperClassName(Class, SuperclassName),
    classDefiningLoader(Class, L),
    loadedClass(SuperclassName, L, Superclass),
    classMethods(Superclass, MethodList),
    finalMethodNotOverridden(Method, Superclass, MethodList).

finalMethodNotOverridden(Method, Superclass, MethodList) :-
    methodName(Method, Name),
    methodDescriptor(Method, Descriptor),
    member(method(_, Name, Descriptor), MethodList),
    isNotFinal(Method, Superclass).

finalMethodNotOverridden(Method, Superclass, MethodList) :-
    methodName(Method, Name),
    methodDescriptor(Method, Descriptor),
    notMember(method(_, Name, Descriptor), MethodList),
    doesNotOverrideFinalMethod(Superclass, Method).
```

### 4.10.1.6. Type Checking Methods with Code

Non-abstract, non-native methods are type correct if they have code and the code is type

correct.

```
methodIsTypeSafe(Class, Method) :-
    doesNotOverrideFinalMethod(Class, Method),
    methodAccessFlags(Method, AccessFlags),
    methodAttributes(Method, Attributes),
    notMember(native, AccessFlags),
    notMember(abstract, AccessFlags),
    member(attribute('Code', _), Attributes),
    methodWithCodeIsTypeSafe(Class, Method).
```

A method with code is type safe if it is possible to merge the code and the stack map frames into a single stream such that each stack map frame precedes the instruction it corresponds to, and the merged stream is type correct. The method's exception handlers, if any, must also be legal.

```
methodWithCodeIsTypeSafe(Class, Method) :-
    parseCodeAttribute(Class, Method, FrameSize, MaxStack,
                       ParsedCode, Handlers, StackMap),
    mergeStackMapAndCode(StackMap, ParsedCode, MergedCode),
    methodInitialStackFrame(Class, Method, FrameSize, StackFrame, ReturnType),
    Environment = environment(Class, Method, ReturnType, MergedCode,
                              MaxStack, Handlers),
    handlersAreLegal(Environment),
    mergedCodeIsTypeSafe(Environment, MergedCode, StackFrame).
```

Let us consider exception handlers first.

An exception handler is represented by a functor application of the form:

```
handler(Start, End, Target, ClassName)
```

whose arguments are, respectively, the start and end of the range of instructions covered by the handler, the first instruction of the handler code, and the name of the exception class that this handler is designed to handle.

An exception handler is *legal* if its start (`Start`) is less than its end (`End`), there exists an instruction whose offset is equal to `Start`, there exists an instruction whose offset equals `End`, and the handler's exception class is assignable to the class `Throwable`. The exception class of a handler is `Throwable` if the handler's class entry is `0`, otherwise it is the class named in the handler.

```
handlersAreLegal(Environment) :-
    exceptionHandlers(Environment, Handlers),
    checklist(handlerIsLegal(Environment), Handlers).

handlerIsLegal(Environment, Handler) :-
    Handler = handler(Start, End, Target, _),
    Start < End,
    allInstructions(Environment, Instructions),
    member(instruction(Start, _), Instructions),
    offsetStackFrame(Environment, Target, _),
    instructionsIncludeEnd(Instructions, End),
    currentClassLoader(Environment, CurrentLoader),
    handlerExceptionClass(Handler, ExceptionClass, CurrentLoader),
    isBootstrapLoader(BL),
    isAssignable(ExceptionClass, class('java/lang/Throwable', BL)).

instructionsIncludeEnd(Instructions, End) :-
    member(instruction(End, _), Instructions).
instructionsIncludeEnd(Instructions, End) :-
    member(endOfCode(End), Instructions).

handlerExceptionClass(handler(_, _, _, 0),
                      class('java/lang/Throwable', BL), _) :-
    isBootstrapLoader(BL).

handlerExceptionClass(handler(_, _, _, Name),
                      class(Name, L), L) :-
    Name \= 0.
```

Let us now turn to the stream of instructions and stack map frames.

Merging instructions and stack map frames into a single stream involves four cases:

- Merging an empty `StackMap` and a list of instructions yields the original list of instructions.

```
mergeStackMapAndCode([], CodeList, CodeList).
```

- Given a list of stack map frames beginning with the type state for the instruction at `Offset`, and a list of instructions beginning at `Offset`, the merged list is the head of the stack frame list, followed by the head of the instruction list, followed by the merge of the tails of the two lists.

```
mergeStackMapAndCode([stackMap(Offset, Map) | RestMap],
                     [instruction(Offset, Parse) | RestCode],
                     [stackMap(Offset, Map),
                         instruction(Offset, Parse) | RestMerge]) :-
    mergeStackMapAndCode(RestMap, RestCode, RestMerge).
```

- Otherwise, given a list of stack frames beginning with the type state for the instruction at `OffsetM`, and a list of instructions beginning at `OffsetP`, then, if `OffsetP < OffsetM`, the merged list consists of the head of the instruction list, followed by the merge of the stack frame list and the tail of the instruction list.

```
mergeStackMapAndCode([stackMap(OffsetM, Map) | RestMap],
                     [instruction(OffsetP, Parse) | RestCode],
                     [instruction(OffsetP, Parse) | RestMerge]) :-
    OffsetP < OffsetM,
    mergeStackMapAndCode([stackMap(OffsetM, Map) | RestMap],
                         RestCode, RestMerge).
```

- Otherwise, the merge of the two lists is undefined. Since the instruction list has monotonically increasing offsets, the merge of the two lists is not defined unless every stack map frame offset has a corresponding instruction offset and the stack map frames are in monotonically increasing order.

To determine if the merged stream for a method is type correct, we first infer the method's initial type state.

The initial type state of a method consists of an empty operand stack and local variable types derived from the type of `this` and the arguments, as well as the appropriate flag, depending on whether this is an `<init>` method.

```
methodInitialStackFrame(Class, Method, FrameSize, frame(Locals, [], Flags),
                        ReturnType):-
    methodDescriptor(Method, Descriptor),
    parseMethodDescriptor(Descriptor, RawArgs, ReturnType),
    expandTypeList(RawArgs, Args),
    methodInitialThisType(Class, Method, ThisList),
    flags(ThisList, Flags),
    append(ThisList, Args, ThisArgs),
    expandToLength(ThisArgs, FrameSize, top, Locals).
```

Given a list of types, the following clause produces a list where every type of size 2 has been substituted by two entries: one for itself, and one `top` entry. The result then corresponds to the representation of the list as 32-bit words in the Java Virtual Machine.

```
expandTypeList([], []).
expandTypeList([Item | List], [Item | Result]) :-
    sizeOf(Item, 1),
    expandTypeList(List, Result).
expandTypeList([Item | List], [Item, top | Result]) :-
    sizeOf(Item, 2),
    expandTypeList(List, Result).
```

```
flags([uninitializedThis], [flagThisUninit]).
flags(X, []) :- X \= [uninitializedThis].

expandToLength(List, Size, _Filler, List) :-
    length(List, Size).
expandToLength(List, Size, Filler, Result) :-
    length(List, ListLength),
    ListLength < Size,
    Delta is Size - ListLength,
```

```
        length(Extra, Delta),
        checklist(=(Filler), Extra),
        append(List, Extra, Result).
```

For the initial type state of an instance method, we compute the type of `this` and put it in a list. The type of `this` in the `<init>` method of `Object` is `Object`; in other `<init>` methods, the type of `this` is `uninitializedThis`; otherwise, the type of `this` in an instance method is `class(N, L)` where `N` is the name of the class containing the method and `L` is its defining class loader.

For the initial type state of a static method, `this` is irrelevant, so the list is empty.

```
methodInitialThisType(_Class, Method, []) :-
    methodAccessFlags(Method, AccessFlags),
    member(static, AccessFlags),
    methodName(Method, MethodName),
    MethodName \= '<init>'.

methodInitialThisType(Class, Method, [This]) :-
    methodAccessFlags(Method, AccessFlags),\
    notMember(static, AccessFlags),\
    instanceMethodInitialThisType(Class, Method, This).

instanceMethodInitialThisType(Class, Method, class('java/lang/Object', L)) :-
    methodName(Method, '<init>'),
    classDefiningLoader(Class, L),
    isBootstrapLoader(L),
    classClassName(Class, 'java/lang/Object').

instanceMethodInitialThisType(Class, Method, uninitializedThis) :-
    methodName(Method, '<init>'),
    classClassName(Class, ClassName),
    classDefiningLoader(Class, CurrentLoader),
    superclassChain(ClassName, CurrentLoader, Chain),
    Chain \= [].

instanceMethodInitialThisType(Class, Method, class(ClassName, L)) :-
```

Are you a developer? Try out the [HTML to PDF API](#)

```
        methodName(Method, MethodName),
        MethodName \= '<init>',
        classDefiningLoader(Class, L),
        classClassName(Class, ClassName).
```

We now compute whether the merged stream for a method is type correct, using the method's initial type state:

- If we have a stack map frame and an incoming type state, the type state must be assignable to the one in the stack map. We may then proceed to type check the rest of the stream with the type state given in the stack map.

```
mergedCodeIsTypeSafe(Environment, [stackMap(Offset, MapFrame) | MoreCode],
                     frame(Locals, OperandStack, Flags)) :-
    frameIsAssignable(frame(Locals, OperandStack, Flags), MapFrame),
    mergedCodeIsTypeSafe(Environment, MoreCode, MapFrame).
```

- A merged code stream is type safe relative to an incoming type state T if it begins with an instruction I that is type safe relative to T, and I *satisfies its exception handlers* (see below), and the tail of the stream is type safe given the type state following that execution of I.

  NextStackFrame indicates what falls through to the following instruction. For an unconditional branch instruction, it will have the special value afterGoto. ExceptionStackFrame indicates what is passed to exception handlers.

```
mergedCodeIsTypeSafe(Environment, [instruction(Offset, Parse) | MoreCode],
                     frame(Locals, OperandStack, Flags)) :-
    instructionIsTypeSafe(Parse, Environment, Offset,
                          frame(Locals, OperandStack, Flags),
                          NextStackFrame, ExceptionStackFrame),
    instructionSatisfiesHandlers(Environment, Offset, ExceptionStackFrame),
    mergedCodeIsTypeSafe(Environment, MoreCode, NextStackFrame).
```

- After an unconditional branch (indicated by an incoming type state of afterGoto), if we have a stack map giving the type state for the following instructions, we can proceed and

type check them using the type state provided by the stack map.

```
mergedCodeIsTypeSafe(Environment, [stackMap(Offset, MapFrame) | MoreCode],
                     afterGoto) :-
    mergedCodeIsTypeSafe(Environment, MoreCode, MapFrame).
```

- It is illegal to have code after an unconditional branch without a stack map frame being provided for it.

```
mergedCodeIsTypeSafe(_Environment, [instruction(_, _) | _MoreCode],
                     afterGoto) :-
    write_ln('No stack frame after unconditional branch'),
    fail.
```

- If we have an unconditional branch at the end of the code, stop.

```
mergedCodeIsTypeSafe(_Environment, [endOfCode(Offset)],
                     afterGoto).
```

Branching to a target is type safe if the target has an associated stack frame, Frame, and the current stack frame, StackFrame, is assignable to Frame.

```
targetIsTypeSafe(Environment, StackFrame, Target) :-
    offsetStackFrame(Environment, Target, Frame),
    frameIsAssignable(StackFrame, Frame).
```

An instruction *satisfies its exception handlers* if it satisfies every exception handler that is applicable to the instruction.

```
instructionSatisfiesHandlers(Environment, Offset, ExceptionStackFrame) :-
    exceptionHandlers(Environment, Handlers),
    sublist(isApplicableHandler(Offset), Handlers, ApplicableHandlers),
    checklist(instructionSatisfiesHandler(Environment, ExceptionStackFrame),
              ApplicableHandlers).
```

An exception handler is *applicable* to an instruction if the offset of the instruction is greater or equal to the start of the handler's range and less than the end of the handler's range.

```
isApplicableHandler(Offset, handler(Start, End, _Target, _ClassName)) :-
    Offset >= Start,
    Offset < End.
```

An instruction *satisfies* an exception handler if its incoming type state is StackFrame, and the handler's target (the initial instruction of the handler code) is type safe assuming an incoming type state T. The type state T is derived from StackFrame by replacing the operand stack with a stack whose sole element is the handler's exception class.

```
instructionSatisfiesHandler(Environment, StackFrame, Handler) :-
    Handler = handler(_, _, Target, _),
    currentClassLoader(Environment, CurrentLoader),
    handlerExceptionClass(Handler, ExceptionClass, CurrentLoader),
    /* The stack consists of just the exception. */
    StackFrame = frame(Locals, _, Flags),
    ExcStackFrame = frame(Locals, [ ExceptionClass ], Flags),
    operandStackHasLegalLength(Environment, ExcStackFrame),
    targetIsTypeSafe(Environment, ExcStackFrame, Target).
```

### 4.10.1.7. Type Checking Load and Store Instructions

All load instructions are variations on a common pattern, varying the type of the value that the instruction loads.

Loading a value of type Type from local variable Index is type safe, if the type of that local variable is ActualType, ActualType is assignable to Type, and pushing ActualType onto the incoming operand stack is a valid type transition ([§4.10.1.4](#)) that yields a new type state NextStackFrame. After execution of the load instruction, the type state will be NextStackFrame.

```
loadIsTypeSafe(Environment, Index, Type, StackFrame, NextStackFrame) :-
```

```
        StackFrame = frame(Locals, _OperandStack, _Flags),
        nth0(Index, Locals, ActualType),
        isAssignable(ActualType, Type),
        validTypeTransition(Environment, [], ActualType, StackFrame,
                                       NextStackFrame).
```

All store instructions are variations on a common pattern, varying the type of the value that the instruction stores.

In general, a store instruction is type safe if the local variable it references is of a type that is a supertype of Type, and the top of the operand stack is of a subtype of Type, where Type is the type the instruction is designed to store.

More precisely, the store is type safe if one can pop a type ActualType that "matches" Type (that is, is a subtype of Type) off the operand stack ([§4.10.1.4](#)), and then legally assign that type the local variable $L_{Index}$.

```
  storeIsTypeSafe(_Environment, Index, Type,
                  frame(Locals, OperandStack, Flags),
                  frame(NextLocals, NextOperandStack, Flags)) :-
      popMatchingType(OperandStack, Type, NextOperandStack, ActualType),
      modifyLocalVariable(Index, ActualType, Locals, NextLocals).
```

Given local variables Locals, modifying Index to have type Type results in the local variable list NewLocals. The modifications are somewhat involved, because some values (and their corresponding types) occupy two local variables. Hence, modifying $L_N$ may require modifying $L_{N+1}$ (because the type will occupy both the N and N+1 slots) or $L_{N-1}$ (because local N used to be the upper half of the two word value/type starting at local N-1, and so local N-1 must be invalidated), or both. This is described further below. We start at $L_0$ and count up.

```
  modifyLocalVariable(Index, Type, Locals, NewLocals) :-
      modifyLocalVariable(0, Index, Type, Locals, NewLocals).
```

Given LocalsRest, the suffix of the local variable list starting at index I, modifying local variable Index to have type Type results in the local variable list suffix NextLocalsRest.

If $I$ < $Index-1$, just copy the input to the output and recurse forward. If $I$ = $Index-1$, the type of local $I$ may change. This can occur if $L_I$ has a type of size 2. Once we set $L_{I+1}$ to the new type (and the corresponding value), the type/value of $L_I$ will be invalidated, as its upper half will be trashed. Then we recurse forward.

```
modifyLocalVariable(I, Index, Type,
                    [Locals1 | LocalsRest],
                    [Locals1 | NextLocalsRest] ) :-
    I < Index - 1,
    I1 is I + 1,
    modifyLocalVariable(I1, Index, Type, LocalsRest, NextLocalsRest).

modifyLocalVariable(I, Index, Type,
                    [Locals1 | LocalsRest],
                    [NextLocals1 | NextLocalsRest] ) :-
    I =:= Index - 1,
    modifyPreIndexVariable(Locals1, NextLocals1),
    modifyLocalVariable(Index, Index, Type, LocalsRest, NextLocalsRest).
```

When we find the variable, and it only occupies one word, we change it to Type and we're done. When we find the variable, and it occupies two words, we change its type to Type and the next word to `top`.

```
modifyLocalVariable(Index, Index, Type,
                    [_ | LocalsRest], [Type | LocalsRest]) :-
    sizeOf(Type, 1).

modifyLocalVariable(Index, Index, Type,
                    [_, _ | LocalsRest], [Type, top | LocalsRest]) :-
    sizeOf(Type, 2).
```

We refer to a local whose index immediately precedes a local whose type will be modified as a *pre-index variable*. The future type of a pre-index variable of type `InputType` is `Result`. If the type, `Type`, of the pre-index local is of size 1, it doesn't change. If the type of the pre-index local, `Type`, is 2, we need to mark the lower half of its two word value as unusable, by setting its type

to `top`.

```
modifyPreIndexVariable(Type, Type) :- sizeOf(Type, 1).
modifyPreIndexVariable(Type, top) :- sizeOf(Type, 2).
```

### 4.10.1.8. Type Checking for `protected` Members

All instructions that access members must contend with the rules concerning `protected` members. This section describes the `protected` check that corresponds to JLS §6.6.2.1.

The `protected` check applies only to `protected` members of superclasses of the current class. `protected` members in other classes will be caught by the access checking done at resolution (§5.4.4). There are four cases:

- If the name of a class is not the name of any superclass, it cannot be a superclass, and so it can safely be ignored.

```
passesProtectedCheck(Environment, MemberClassName, MemberName,
                     MemberDescriptor, StackFrame) :-
    thisClass(Environment, class(CurrentClassName, CurrentLoader)),
    superclassChain(CurrentClassName, CurrentLoader, Chain),
    notMember(class(MemberClassName, _), Chain).
```

- If the `MemberClassName` is the same as the name of a superclass, the class being resolved may indeed be a superclass. In this case, if no superclass named `MemberClassName` in a different run-time package has a `protected` member named `MemberName` with descriptor `MemberDescriptor`, the `protected` check does not apply.

  *This is because the actual class being resolved will either be one of these superclasses, in which case we know that it is either in the same run-time package, and the access is legal; or the member in question is not `protected` and the check does not apply; or it will be a subclass, in which case the check would succeed anyway; or it will be some other class in the same run-time package, in which case the access is legal and the check need not take place; or the verifier need not flag this as a problem, since it will be caught anyway because resolution will per force fail.*

```
passesProtectedCheck(Environment, MemberClassName, MemberName,
                     MemberDescriptor, StackFrame) :-
    thisClass(Environment, class(CurrentClassName, CurrentLoader)),
    superclassChain(CurrentClassName, CurrentLoader, Chain),
    member(class(MemberClassName, _), Chain),
    classesInOtherPkgWithProtectedMember(
      class(CurrentClassName, CurrentLoader),
      MemberName, MemberDescriptor, MemberClassName, Chain, []).
```

- If there does exist a `protected` superclass member in a different run-time package, then load `MemberClassName`; if the member in question is not `protected`, the check does not apply. (Using a superclass member that is not `protected` is trivially correct.)

```
passesProtectedCheck(Environment, MemberClassName, MemberName,
                     MemberDescriptor,
                     frame(_Locals, [Target | Rest], _Flags)) :-
    thisClass(Environment, class(CurrentClassName, CurrentLoader)),
    superclassChain(CurrentClassName, CurrentLoader, Chain),
    member(class(MemberClassName, _), Chain),
    classesInOtherPkgWithProtectedMember(
      class(CurrentClassName, CurrentLoader),
      MemberName, MemberDescriptor, MemberClassName, Chain, List),
    List /= [],
    loadedClass(MemberClassName, CurrentLoader, ReferencedClass),
    isNotProtected(ReferencedClass, MemberName, MemberDescriptor).
```

- Otherwise, use of a member of an object of type `Target` requires that `Target` be assignable to the type of the current class.

```
passesProtectedCheck(Environment, MemberClassName, MemberName,
                     MemberDescriptor,
                     frame(_Locals, [Target | Rest], _Flags)) :-
    thisClass(Environment, class(CurrentClassName, CurrentLoader)),
    superclassChain(CurrentClassName, CurrentLoader, Chain),
    member(class(MemberClassName, _), Chain),
    classesInOtherPkgWithProtectedMember(
```

```
            class(CurrentClassName, CurrentLoader),
          MemberName, MemberDescriptor, MemberClassName, Chain, List),
        List /= [],
        loadedClass(MemberClassName, CurrentLoader, ReferencedClass),
        isProtected(ReferencedClass, MemberName, MemberDescriptor),
        isAssignable(Target, class(CurrentClassName, CurrentLoader)).
```

The predicate `classesInOtherPkgWithProtectedMember(Class, MemberName, MemberDescriptor, MemberClassName, Chain, List)` is true if `List` is the set of classes in `Chain` with name `MemberClassName` that are in a different run-time package than `Class` which have a `protected` member named `MemberName` with descriptor `MemberDescriptor`.

```
classesInOtherPkgWithProtectedMember(_, _, _, _, [], []).

classesInOtherPkgWithProtectedMember(Class, MemberName,
                                     MemberDescriptor, MemberClassName,
                                     [class(MemberClassName, L) | Tail],
                                     [class(MemberClassName, L) | T]) :-
    differentRuntimePackage(Class, class(MemberClassName, L)),
    loadedClass(MemberClassName, L, Super),
    isProtected(Super, MemberName, MemberDescriptor),
    classesInOtherPkgWithProtectedMember(
      Class, MemberName, MemberDescriptor, MemberClassName, Tail, T).

classesInOtherPkgWithProtectedMember(Class, MemberName,
                                     MemberDescriptor, MemberClassName,
                                     [class(MemberClassName, L) | Tail],
                                     T) :-
    differentRuntimePackage(Class, class(MemberClassName, L)),
    loadedClass(MemberClassName, L, Super),
    isNotProtected(Super, MemberName, MemberDescriptor),
    classesInOtherPkgWithProtectedMember(
      Class, MemberName, MemberDescriptor, MemberClassName, Tail, T).

classesInOtherPkgWithProtectedMember(Class, MemberName,
                                     MemberDescriptor, MemberClassName,
```

```
                              [class(MemberClassName, L) | Tail],
                          T] :-
    sameRuntimePackage(Class, class(MemberClassName, L)),
    classesInOtherPkgWithProtectedMember(
      Class, MemberName, MemberDescriptor, MemberClassName, Tail, T).
```

```
sameRuntimePackage(Class1, Class2) :-
    classDefiningLoader(Class1, L),
    classDefiningLoader(Class2, L),
    samePackageName(Class1, Class2).

differentRuntimePackage(Class1, Class2) :-
    classDefiningLoader(Class1, L1),
    classDefiningLoader(Class2, L2),
    L1 \= L2.

differentRuntimePackage(Class1, Class2) :-
    differentPackageName(Class1, Class2).
```

### 4.10.1.9. Type Checking Instructions

In general, the type rule for an instruction is given relative to an environment `Environment` that defines the class and method in which the instruction occurs (§4.10.1.1), and the offset `Offset` within the method at which the instruction occurs. The rule states that if the incoming type state `StackFrame` fulfills certain requirements, then:

- The instruction is type safe.

- It is provable that the type state after the instruction completes normally has a particular form given by `NextStackFrame`, and that the type state after the instruction completes abruptly is given by `ExceptionStackFrame`.

  The type state after an instruction completes abruptly is the same as the incoming type state, except that the operand stack is empty.

```
exceptionStackFrame(StackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, _OperandStack, Flags),
    ExceptionStackFrame = frame(Locals, [], Flags).
```

Many instructions have type rules that are completely isomorphic to the rules for other instructions. If an instruction b1 is isomorphic to another instruction b2, then the type rule for b1 is the same as the type rule for b2.

```
instructionIsTypeSafe(Instruction, Environment, Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    instructionHasEquivalentTypeRule(Instruction, IsomorphicInstruction),
    instructionIsTypeSafe(IsomorphicInstruction, Environment, Offset,
                          StackFrame, NextStackFrame,
                          ExceptionStackFrame).
```

The English language description of each rule is intended to be readable, intuitive, and concise. As such, the description avoids repeating all the contextual assumptions given above. In particular:

- The description does not explicitly mention the environment.

- When the description speaks of the operand stack or local variables in the following, it is referring to the operand stack and local variable components of a type state: either the incoming type state or the outgoing one.

- The type state after the instruction completes abruptly is almost always identical to the incoming type state. The description only discusses the type state after the instruction completes abruptly when that is not the case.

- The description speaks of popping and pushing types onto the operand stack, and does not explicitly discuss issues of stack underflow or overflow. The description assumes these operations can be completed successfully, but the Prolog clauses for operand stack manipulation ensure that the necessary checks are made.

- The description discusses only the manipulation of logical types. In practice, some types take more than one word. The description abstracts from these representation details, but the Prolog clauses that manipulate data do not.

Any ambiguities can be resolved by referring to the formal Prolog clauses.

### *aaload*

An *aaload* instruction is type safe iff one can validly replace types matching `int` and an array type with component type `ComponentType` where `ComponentType` is a subtype of `Object`, with `ComponentType` yielding the outgoing type state.

```
instructionIsTypeSafe(aaload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    nth1OperandStackIs(2, StackFrame, ArrayType),
    arrayComponentType(ArrayType, ComponentType),
    isBootstrapLoader(BL),
    validTypeTransition(Environment,
                        [int, arrayOf(class('java/lang/Object', BL))],
                        ComponentType, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The component type of an array of X is X. We define the component type of `null` to be `null`.

```
arrayComponentType(arrayOf(X), X).
arrayComponentType(null, null).
```

### *aastore*

An *aastore* instruction is type safe iff one can validly pop types matching `Object`, `int`, and an array of `Object` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(aastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    isBootstrapLoader(BL),
    canPop(StackFrame,
           [class('java/lang/Object', BL),
            int,
```

```
                arrayOf(class('java/lang/Object', BL))],
            NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### aconst_null

An *aconst_null* instruction is type safe if one can validly push the type `null` onto the incoming
operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(aconst_null, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [], null, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### aload

An *aload* instruction with operand `Index` is type safe and yields an outgoing type state
`NextStackFrame`, if a load instruction with operand `Index` and type `reference` is type safe
and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(aload(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    loadIsTypeSafe(Environment, Index, reference, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### aload_<n>

The instructions *aload_<n>*, for $0 \leq n \leq 3$, are type safe iff the equivalent *aload* instruction is
type safe.

```
instructionHasEquivalentTypeRule(aload_0, aload(0)).
instructionHasEquivalentTypeRule(aload_1, aload(1)).
instructionHasEquivalentTypeRule(aload_2, aload(2)).
```

```
instructionHasEquivalentTypeRule(aload_3, aload(3)).
```

### *anewarray*

An *anewarray* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting either a class type or an array type, and one can legally replace a type matching int on the incoming operand stack with an array with component type CP yielding the outgoing type state.

```
instructionIsTypeSafe(anewarray(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    (CP = class(_, _) ; CP = arrayOf(_)),
    validTypeTransition(Environment, [int], arrayOf(CP),
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *areturn*

An *areturn* instruction is type safe iff the enclosing method has a declared return type, ReturnType, that is a reference type, and one can validly pop a type matching ReturnType off the incoming operand stack.

```
instructionIsTypeSafe(areturn, Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    thisMethodReturnType(Environment, ReturnType),
    isAssignable(ReturnType, reference),
    canPop(StackFrame, [ReturnType], _PoppedStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *arraylength*

An *arraylength* instruction is type safe iff one can validly replace an array type on the incoming operand stack with the type int yielding the outgoing type state.

```
instructionIsTypeSafe(arraylength, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    nth1OperandStackIs(1, StackFrame, ArrayType),
    arrayComponentType(ArrayType, _),
    validTypeTransition(Environment, [top], int, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *astore*

An *astore* instruction with operand `Index` is type safe and yields an outgoing type state
`NextStackFrame`, if a store instruction with operand `Index` and type `reference` is type safe
and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(astore(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    storeIsTypeSafe(Environment, Index, reference, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *astore_<n>*

The instructions *astore_<n>*, for $0 \leq n \leq 3$, are type safe iff the equivalent *astore* instruction is
type safe.

```
instructionHasEquivalentTypeRule(astore_0, astore(0)).
instructionHasEquivalentTypeRule(astore_1, astore(1)).
instructionHasEquivalentTypeRule(astore_2, astore(2)).
instructionHasEquivalentTypeRule(astore_3, astore(3)).
```

### *athrow*

An *athrow* instruction is type safe iff the top of the operand stack matches `Throwable`.

```
instructionIsTypeSafe(athrow, _Environment, _Offset, StackFrame,
```

```
                    afterGoto, ExceptionStackFrame) :-
    isBootstrapLoader(BL),
    canPop(StackFrame, [class('java/lang/Throwable', BL)], _PoppedStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *baload*

A *baload* instruction is type safe iff one can validly replace types matching `int` and a small array type on the incoming operand stack with `int` yielding the outgoing type state.

```
instructionIsTypeSafe(baload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :
    nth1OperandStackIs(2, StackFrame, ArrayType),
    isSmallArray(ArrayType),
    validTypeTransition(Environment, [int, top], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

An array type is a *small array type* if it is an array of `byte`, an array of `boolean`, or a subtype thereof (`null`).

```
isSmallArray(arrayOf(byte)).
isSmallArray(arrayOf(boolean)).
isSmallArray(null).
```

### *bastore*

A *bastore* instruction is type safe iff one can validly pop types matching `int`, `int` and a small array type off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(bastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    nth1OperandStackIs(3, StackFrame, ArrayType),
    isSmallArray(ArrayType),
```

```
        canPop(StackFrame, [int, int, top], NextStackFrame),
        exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

***bipush***

A *bipush* instruction is type safe iff the equivalent *sipush* instruction is type safe.

```
  instructionHasEquivalentTypeRule(bipush(Value), sipush(Value)).
```

***caload***

A *caload* instruction is type safe iff one can validly replace types matching `int` and array of `char` on the incoming operand stack with `int` yielding the outgoing type state.

```
  instructionIsTypeSafe(caload, Environment, _Offset, StackFrame,
                        NextStackFrame, ExceptionStackFrame) :-
      validTypeTransition(Environment, [int, arrayOf(char)], int,
                          StackFrame, NextStackFrame),
      exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

***castore***

A *castore* instruction is type safe iff one can validly pop types matching `int`, `int` and array of `char` off the incoming operand stack yielding the outgoing type state.

```
  instructionIsTypeSafe(castore, _Environment, _Offset, StackFrame,
                        NextStackFrame, ExceptionStackFrame) :-
      canPop(StackFrame, [int, int, arrayOf(char)], NextStackFrame),
      exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

***checkcast***

Are you a developer? Try out the HTML to PDF API

A *checkcast* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting either a class or an array, and one can validly replace the type `Object` on top of the incoming operand stack with the type denoted by CP yielding the outgoing type state.

```
instructionIsTypeSafe(checkcast(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    (CP = class(_, _) ; CP = arrayOf(_)),
    isBootstrapLoader(BL),
    validTypeTransition(Environment, [class('java/lang/Object', BL)], CP,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### d2f

A *d2f* instruction is type safe if one can validly pop `double` off the incoming operand stack and replace it with `float`, yielding the outgoing type state.

```
instructionIsTypeSafe(d2f, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [double], float,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### d2i

A *d2i* instruction is type safe if one can validly pop `double` off the incoming operand stack and replace it with `int`, yielding the outgoing type state.

```
instructionIsTypeSafe(d2i, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [double], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### d2l

A *d2l* instruction is type safe if one can validly pop `double` off the incoming operand stack and replace it with `long`, yielding the outgoing type state.

```
instructionIsTypeSafe(d2l, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [double], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### dadd

A *dadd* instruction is type safe iff one can validly replace types matching `double` and `double` on the incoming operand stack with `double` yielding the outgoing type state.

```
instructionIsTypeSafe(dadd, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [double, double], double,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### daload

A *daload* instruction is type safe iff one can validly replace types matching `int` and array of `double` on the incoming operand stack with `double` yielding the outgoing type state.

```
instructionIsTypeSafe(daload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, arrayOf(double)], double,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### dastore

A *dastore* instruction is type safe iff one can validly pop types matching double, int and array of double off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(dastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [double, int, arrayOf(double)], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### dcmp<op>

A *dcmpg* instruction is type safe iff one can validly replace types matching double and double on the incoming operand stack with int yielding the outgoing type state.

```
instructionIsTypeSafe(dcmpg, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [double, double], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *dcmpl* instruction is type safe iff the equivalent *dcmpg* instruction is type safe.

```
instructionHasEquivalentTypeRule(dcmpl, dcmpg).
```

### dconst_<d>

A *dconst_0* instruction is type safe if one can validly push the type double onto the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(dconst_0, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [], double, StackFrame, NextStackFrame),
```

```
                    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *dconst_1* instruction is type safe iff the equivalent *dconst_0* instruction is type safe.

```
instructionHasEquivalentTypeRule(dconst_1, dconst_0).
```

### ddiv

A *ddiv* instruction is type safe iff the equivalent *dadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(ddiv, dadd).
```

### dload

A *dload* instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `double` is type safe and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(dload(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    loadIsTypeSafe(Environment, Index, double, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### dload_<n>

The instructions *dload_<n>*, for $0 \leq n \leq 3$, are typesafe iff the equivalent *dload* instruction is type safe.

```
instructionHasEquivalentTypeRule(dload_0, dload(0)).
instructionHasEquivalentTypeRule(dload_1, dload(1)).
instructionHasEquivalentTypeRule(dload_2, dload(2)).
instructionHasEquivalentTypeRule(dload_3, dload(3)).
```

### dmul

A *dmul* instruction is type safe iff the equivalent *dadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(dmul, dadd).
```

### dneg

A *dneg* instruction is type safe iff there is a type matching `double` on the incoming operand stack. The *dneg* instruction does not alter the type state.

```
instructionIsTypeSafe(dneg, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [double], double,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### drem

A *drem* instruction is type safe iff the equivalent *dadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(drem, dadd).
```

### dreturn

A *dreturn* instruction is type safe if the enclosing method has a declared return type of `double`, and one can validly pop a type matching `double` off the incoming operand stack.

```
instructionIsTypeSafe(dreturn, Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    thisMethodReturnType(Environment, double),
```

```
        canPop(StackFrame, [double], _PoppedStackFrame),
        exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *dstore*

A *dstore* instruction with operand `Index` is type safe and yields an outgoing type state
`NextStackFrame`, if a store instruction with operand `Index` and type `double` is type safe and
yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(dstore(Index), Environment, _Offset, StackFrame,
                        NextStackFrame, ExceptionStackFrame) :-
    storeIsTypeSafe(Environment, Index, double, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *dstore_<n>*

The instructions *dstore_<n>*, for 0 ≤ *n* ≤ 3, are type safe iff the equivalent *dstore* instruction is
type safe.

```
instructionHasEquivalentTypeRule(dstore_0, dstore(0)).
instructionHasEquivalentTypeRule(dstore_1, dstore(1)).
instructionHasEquivalentTypeRule(dstore_2, dstore(2)).
instructionHasEquivalentTypeRule(dstore_3, dstore(3)).
```

### *dsub*

A *dsub* instruction is type safe iff the equivalent *dadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(dsub, dadd).
```

### *dup*

A *dup* instruction is type safe iff one can validly replace a category 1 type, Type, with the types Type, Type, yielding the outgoing type state.

```
instructionIsTypeSafe(dup, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
    popCategory1(InputOperandStack, Type, _),
    canSafelyPush(Environment, InputOperandStack, Type, OutputOperandStack),
    NextStackFrame = frame(Locals, OutputOperandStack, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### dup_x1

A *dup_x1* instruction is type safe iff one can validly replace two category 1 types, Type1, and Type2, on the incoming operand stack with the types Type1, Type2, Type1, yielding the outgoing type state.

```
instructionIsTypeSafe(dup_x1, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
    popCategory1(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Rest),
    canSafelyPushList(Environment, Rest, [Type1, Type2, Type1],
                      OutputOperandStack),
    NextStackFrame = frame(Locals, OutputOperandStack, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### dup_x2

A *dup_x2* instruction is type safe iff it is a *type safe form* of the *dup_x2* instruction.

```
instructionIsTypeSafe(dup_x2, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
```

```
        dup_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack),
        NextStackFrame = frame(Locals, OutputOperandStack, Flags),
        exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *dup_x2* instruction is a *type safe form* of the *dup_x2* instruction iff it is a *type safe form 1 dup_x2* instruction or a *type safe form 2 dup_x2* instruction.

```
dup_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).
```

A *dup_x2* instruction is a *type safe form 1 dup_x2* instruction iff one can validly replace three category 1 types, Type1, Type2, Type3 on the incoming operand stack with the types Type1, Type2, Type3, Type1, yielding the outgoing type state.

```
dup_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory1(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Stack2),
    popCategory1(Stack2, Type3, Rest),
    canSafelyPushList(Environment, Rest, [Type1, Type3, Type2, Type1],
                      OutputOperandStack).
```

A *dup_x2* instruction is a *type safe form 2 dup_x2* instruction iff one can validly replace a category 1 type, Type1, and a category 2 type, Type2, on the incoming operand stack with the types Type1, Type2, Type1, yielding the outgoing type state.

```
dup_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory1(InputOperandStack, Type1, Stack1),
    popCategory2(Stack1, Type2, Rest),
    canSafelyPushList(Environment, Rest, [Type1, Type2, Type1],
                      OutputOperandStack).
```

**dup2**

.

A *dup2* instruction is type safe iff it is a *type safe form* of the *dup2* instruction.

```
instructionIsTypeSafe(dup2, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
    dup2SomeFormIsTypeSafe(Environment,InputOperandStack, OutputOperandStack),
    NextStackFrame = frame(Locals, OutputOperandStack, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *dup2* instruction is a *type safe form* of the *dup2* instruction iff it is a *type safe form 1 dup2*
instruction or a *type safe form 2 dup2* instruction.

```
dup2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2Form1IsTypeSafe(Environment,InputOperandStack, OutputOperandStack).

dup2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2Form2IsTypeSafe(Environment,InputOperandStack, OutputOperandStack).
```

A *dup2* instruction is a *type safe form 1 dup2* instruction iff one can validly replace two category
1 types, `Type1` and `Type2` on the incoming operand stack with the types `Type1`, `Type2`,
`Type1`, `Type2`, yielding the outgoing type state.

```
dup2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack):-
    popCategory1(InputOperandStack, Type1, TempStack),
    popCategory1(TempStack, Type2, _),
    canSafelyPushList(Environment, InputOperandStack, [Type1, Type2],
                      OutputOperandStack).
```

A *dup2* instruction is a *type safe form 2 dup2* instruction iff one can validly replace a category 2
type, `Type` on the incoming operand stack with the types `Type`, `Type`, yielding the outgoing
type state.

```
dup2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack):-
    popCategory2(InputOperandStack, Type, _),
```

Are you a developer? Try out the HTML to PDF API

```
        canSafelyPush(Environment, InputOperandStack, Type, OutputOperandStack).
```

### dup2_x1

A *dup2_x1* instruction is type safe iff it is a *type safe form* of the *dup2_x1* instruction.

```
instructionIsTypeSafe(dup2_x1, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
    dup2_x1SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack),
    NextStackFrame = frame(Locals, OutputOperandStack, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *dup2_x1* instruction is a *type safe form* of the *dup2_x1* instruction iff it is a *type safe form 1 dup2_x1* instruction or a *type safe form 2 dup_x2* instruction.

```
dup2_x1SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2_x1Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x1SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2_x1Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).
```

A *dup2_x1* instruction is a *type safe form 1 dup2_x1* instruction iff one can validly replace three category 1 types, Type1, Type2, Type3, on the incoming operand stack with the types Type1, Type2, Type3, Type1, Type2, yielding the outgoing type state.

```
dup2_x1Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory1(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Stack2),
    popCategory1(Stack2, Type3, Rest),
    canSafelyPushList(Environment, Rest, [Type2, Type1, Type3, Type2, Type1],
                      OutputOperandStack).
```

A *dup2_x1* instruction is a *type safe form 2 dup2_x1* instruction iff one can validly replace a

category 2 type, Type1, and a category 1 type, Type2, on the incoming operand stack with the types Type1, Type2, Type1, yielding the outgoing type state.

```
dup2_x1Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory2(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Rest),
    canSafelyPushList(Environment, Rest, [Type1, Type2, Type1],
                      OutputOperandStack).
```

### dup2_x2

A *dup2_x2* instruction is type safe iff it is a *type safe form* of the *dup2_x2* instruction.

```
instructionIsTypeSafe(dup2_x2, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
    dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack),
    NextStackFrame = frame(Locals, OutputOperandStack, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *dup2_x2* instruction is a *type safe form* of the *dup2_x2* instruction iff one of the following holds:

- it is a *type safe form 1 dup2_x2* instruction.

- it is a *type safe form 2 dup2_x2* instruction.

- it is a *type safe form 3 dup2_x2* instruction.

- it is a *type safe form 4 dup2_x2* instruction.

```
dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).
```

```
dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2_x2Form3IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    dup2_x2Form4IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).
```

A *dup2_x2* instruction is a *type safe form 1 dup2_x2* instruction iff one can validly replace four
category 1 types, Type1, Type2, Type3, Type4, on the incoming operand stack with the types
Type1, Type2, Type3, Type4, Type1, Type2, yielding the outgoing type state.

```
dup2_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory1(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Stack2),
    popCategory1(Stack2, Type3, Stack3),
    popCategory1(Stack3, Type4, Rest),
    canSafelyPushList(Environment, Rest,
                      [Type2, Type1, Type4, Type3, Type2, Type1],
                      OutputOperandStack).
```

A *dup2_x2* instruction is a *type safe form 2 dup2_x2* instruction iff one can validly replace a
category 2 type, Type1, and two category 1 types, Type2, Type3, on the incoming operand
stack with the types Type1, Type2, Type3, Type1, yielding the outgoing type state.

```
dup2_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory2(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Stack2),
    popCategory1(Stack2, Type3, Rest),
    canSafelyPushList(Environment, Rest,
                      [Type1, Type3, Type2, Type1],
                      OutputOperandStack).
```

A *dup2_x2* instruction is a *type safe form 3 dup2_x2* instruction iff one can validly replace two
category 1 types, Type1, Type2, and a category 2 type, Type3, on the incoming operand stack
with the types Type1, Type2, Type3, Type1, Type2, yielding the outgoing type state.

```
dup2_x2Form3IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory1(InputOperandStack, Type1, Stack1),
    popCategory1(Stack1, Type2, Stack2),
    popCategory2(Stack2, Type3, Rest),
    canSafelyPushList(Environment, Rest,
                      [Type2, Type1, Type3, Type2, Type1],
                      OutputOperandStack).
```

A *dup2_x2* instruction is a *type safe form 4 dup2_x2* instruction iff one can validly replace two
category 2 types, Type1, Type2, on the incoming operand stack with the types Type1, Type2,
Type1, yielding the outgoing type state.

```
dup2_x2Form4IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-
    popCategory2(InputOperandStack, Type1, Stack1),
    popCategory2(Stack1, Type2, Rest),
    canSafelyPushList(Environment, Rest, [Type1, Type2, Type1],
                      OutputOperandStack).
```

### f2d

An *f2d* instruction is type safe if one can validly pop float off the incoming operand stack and
replace it with double, yielding the outgoing type state.

```
instructionIsTypeSafe(f2d, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [float], double,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### f2i

An *f2i* instruction is type safe if one can validly pop float off the incoming operand stack and
replace it with int, yielding the outgoing type state.

```
instructionIsTypeSafe(f2i, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [float], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### f2l

An *f2l* instruction is type safe if one can validly pop `float` off the incoming operand stack and replace it with `long`, yielding the outgoing type state.

```
instructionIsTypeSafe(f2l, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [float], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### fadd

An *fadd* instruction is type safe iff one can validly replace types matching `float` and `float` on the incoming operand stack with `float` yielding the outgoing type state.

```
instructionIsTypeSafe(fadd, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [float, float], float,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### faload

An *faload* instruction is type safe iff one can validly replace types matching `int` and array of `float` on the incoming operand stack with `float` yielding the outgoing type state.

```
instructionIsTypeSafe(faload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, arrayOf(float)], float,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *fastore*

An *fastore* instruction is type safe iff one can validly pop types matching `float`, `int` and array of `float` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(fastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [float, int, arrayOf(float)], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *fcmp<op>*

An *fcmpg* instruction is type safe iff one can validly replace types matching `float` and `float` on the incoming operand stack with `int` yielding the outgoing type state.

```
instructionIsTypeSafe(fcmpg, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [float, float], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

An *fcmpl* instruction is type safe iff the equivalent *fcmpg* instruction is type safe.

```
instructionHasEquivalentTypeRule(fcmpl, fcmpg).
```

### *fconst_<f>*

An *fconst_0* instruction is type safe if one can validly push the type `float` onto the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(fconst_0, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [], float, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The rules for the other variants of *fconst* are equivalent.

```
instructionHasEquivalentTypeRule(fconst_1, fconst_0).
instructionHasEquivalentTypeRule(fconst_2, fconst_0).
```

### *fdiv*

An *fdiv* instruction is type safe iff the equivalent *fadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(fdiv, fadd).
```

### *fload*

An *fload* instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `float` is type safe and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(fload(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    loadIsTypeSafe(Environment, Index, float, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *fload_<n>*

Are you a developer? Try out the [HTML to PDF API](#)

The instructions *fload_<n>*, for 0 ≤ *n* ≤ 3, are typesafe iff the equivalent *fload* instruction is type safe.

```
instructionHasEquivalentTypeRule(fload_0, fload(0)).
instructionHasEquivalentTypeRule(fload_1, fload(1)).
instructionHasEquivalentTypeRule(fload_2, fload(2)).
instructionHasEquivalentTypeRule(fload_3, fload(3)).
```

### *fmul*

An *fmul* instruction is type safe iff the equivalent *fadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(fmul, fadd).
```

### *fneg*

An *fneg* instruction is type safe iff there is a type matching `float` on the incoming operand stack. The *fneg* instruction does not alter the type state.

```
instructionIsTypeSafe(fneg, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [float], float,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *frem*

An *frem* instruction is type safe iff the equivalent *fadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(frem, fadd).
```

### *freturn*

Are you a developer? Try out the HTML to PDF API

An *freturn* instruction is type safe if the enclosing method has a declared return type of `float`, and one can validly pop a type matching `float` off the incoming operand stack.

```
instructionIsTypeSafe(freturn, Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    thisMethodReturnType(Environment, float),
    canPop(StackFrame, [float], _PoppedStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *fstore*

An *fstore* instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a store instruction with operand `Index` and type `float` is type safe and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(fstore(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    storeIsTypeSafe(Environment, Index, float, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *fstore_<n>*

The instructions *fstore_<n>*, for $0 \leq n \leq 3$, are typesafe iff the equivalent *fstore* instruction is type safe.

```
instructionHasEquivalentTypeRule(fstore_0, fstore(0)).
instructionHasEquivalentTypeRule(fstore_1, fstore(1)).
instructionHasEquivalentTypeRule(fstore_2, fstore(2)).
instructionHasEquivalentTypeRule(fstore_3, fstore(3)).
```

### *fsub*

An *fsub* instruction is type safe iff the equivalent *fadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(fsub, fadd).
```

### getfield

A *getfield* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting a field whose declared type is `FieldType`, declared in a class `FieldClass`, and one can validly replace a type matching `FieldClass` with type `FieldType` on the incoming operand stack yielding the outgoing type state. `FieldClass` must not be an array type. `protected` fields are subject to additional checks ([§4.10.1.8](#)).

```
instructionIsTypeSafe(getfield(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = field(FieldClass, FieldName, FieldDescriptor),
    parseFieldDescriptor(FieldDescriptor, FieldType),
    passesProtectedCheck(Environment, FieldClass, FieldName,
                         FieldDescriptor, StackFrame),
    validTypeTransition(Environment, [class(FieldClass)], FieldType,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### getstatic

A *getstatic* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting a field whose declared type is `FieldType`, and one can validly push `FieldType` on the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(getstatic(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = field(_FieldClass, _FieldName, FieldDescriptor),
    parseFieldDescriptor(FieldDescriptor, FieldType),
    validTypeTransition(Environment, [], FieldType,
                        StackFrame, NextStackFrame),
```

```
                    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### goto

A *goto* instruction is type safe iff its target operand is a valid branch target.

```
instructionIsTypeSafe(goto(Target), Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    targetIsTypeSafe(Environment, StackFrame, Target),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### goto_w

A *goto_w* instruction is type safe iff the equivalent *goto* instruction is type safe.

```
instructionHasEquivalentTypeRule(goto_w(Target), goto(Target)).
```

### i2b

An *i2b* instruction is type safe iff the equivalent *ineg* instruction is type safe.

```
instructionHasEquivalentTypeRule(i2b, ineg).
```

### i2c

An *i2c* instruction is type safe iff the equivalent *ineg* instruction is type safe.

```
instructionHasEquivalentTypeRule(i2c, ineg).
```

### i2d

An *i2d* instruction is type safe if one can validly pop int off the incoming operand stack and replace it with double, yielding the outgoing type state.

```
instructionIsTypeSafe(i2d, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int], double,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

*i2f*

An *i2f* instruction is type safe if one can validly pop int off the incoming operand stack and replace it with float, yielding the outgoing type state.

```
instructionIsTypeSafe(i2f, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int], float,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

*i2l*

An *i2l* instruction is type safe if one can validly pop int off the incoming operand stack and replace it with long, yielding the outgoing type state.

```
instructionIsTypeSafe(i2l, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

*i2s*

An *i2s* instruction is type safe iff the equivalent *ineg* instruction is type safe.

```
instructionHasEquivalentTypeRule(i2s, ineg).
```

### *iadd*

An *iadd* instruction is type safe iff one can validly replace types matching `int` and `int` on the incoming operand stack with `int` yielding the outgoing type state.

```
instructionIsTypeSafe(iadd, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, int], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *iaload*

An *iaload* instruction is type safe iff one can validly replace types matching `int` and array of `int` on the incoming operand stack with `int` yielding the outgoing type state.

```
instructionIsTypeSafe(iaload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, arrayOf(int)], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *iand*

An *iand* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(iand, iadd).
```

### iastore

An *iastore* instruction is type safe iff one can validly pop types matching `int`, `int` and array of `int` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(iastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [int, int, arrayOf(int)], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### if_acmp<cond>

An *if_acmpeq* instruction is type safe iff one can validly pop types matching `reference` and `reference` on the incoming operand stack yielding the outgoing type state `NextStackFrame`, and the operand of the instruction, `Target`, is a valid branch target assuming an incoming type state of `NextStackFrame`.

```
instructionIsTypeSafe(if_acmpeq(Target), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [reference, reference], NextStackFrame),
    targetIsTypeSafe(Environment, NextStackFrame, Target),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The rule for *if_acmpne* is identical.

```
instructionHasEquivalentTypeRule(if_acmpne(Target), if_acmpeq(Target)).
```

### if_icmp<cond>

An *if_icmpeq* instruction is type safe iff one can validly pop types matching `int` and `int` on the incoming operand stack yielding the outgoing type state `NextStackFrame`, and the operand of the instruction, `Target`, is a valid branch target assuming an incoming type state of `NextStackFrame`.

```
instructionIsTypeSafe(if_icmpeq(Target), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [int, int], NextStackFrame),
    targetIsTypeSafe(Environment, NextStackFrame, Target),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The rules for all other variants of the *if_icmp<cond>* instruction are identical.

```
instructionHasEquivalentTypeRule(if_icmpge(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmpgt(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmple(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmplt(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmpne(Target), if_icmpeq(Target)).
```

### if<cond>

An *ifeq* instruction is type safe iff one can validly pop a type matching `int` off the incoming operand stack yielding the outgoing type state `NextStackFrame`, and the operand of the instruction, `Target`, is a valid branch target assuming an incoming type state of `NextStackFrame`.

```
instructionIsTypeSafe(ifeq(Target), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [int], NextStackFrame),
    targetIsTypeSafe(Environment, NextStackFrame, Target),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The rules for all other variations of the *if<cond>* instruction are identical.

```
instructionHasEquivalentTypeRule(ifge(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(ifgt(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(ifle(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(iflt(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(ifne(Target), ifeq(Target)).
```

### ifnonnull

An *ifnonnull* instruction is type safe iff one can validly pop a type matching `reference` off the incoming operand stack yielding the outgoing type state `NextStackFrame`, and the operand of the instruction, `Target`, is a valid branch target assuming an incoming type state of `NextStackFrame`.

```
instructionIsTypeSafe(ifnonnull(Target), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [reference], NextStackFrame),
    targetIsTypeSafe(Environment, NextStackFrame, Target),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### ifnull

An *ifnull* instruction is type safe iff the equivalent *ifnonnull* instruction is type safe.

```
instructionHasEquivalentTypeRule(ifnull(Target), ifnonnull(Target)).
```

### iinc

An *iinc* instruction with first operand `Index` is type safe iff $L_{Index}$ has type `int`. The *iinc* instruction does not change the type state.

```
instructionIsTypeSafe(iinc(Index, _Value), _Environment, _Offset,
                      StackFrame, StackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, _OperandStack, _Flags),
    nth0(Index, Locals, int),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### iload

An *iload* instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `int` is type safe and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(iload(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    loadIsTypeSafe(Environment, Index, int, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### iload_<n>

The instructions *iload_<n>*, for 0 ≤ *n* ≤ 3, are typesafe iff the equivalent *iload* instruction is type safe.

```
instructionHasEquivalentTypeRule(iload_0, iload(0)).
instructionHasEquivalentTypeRule(iload_1, iload(1)).
instructionHasEquivalentTypeRule(iload_2, iload(2)).
instructionHasEquivalentTypeRule(iload_3, iload(3)).
```

### imul

An *imul* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(imul, iadd).
```

### ineg

An *ineg* instruction is type safe iff there is a type matching `int` on the incoming operand stack. The *ineg* instruction does not alter the type state.

```
instructionIsTypeSafe(ineg, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
```

```
              validTypeTransition(Environment, [int], int, StackFrame, NextStackFrame),
              exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *instanceof*

An *instanceof* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting either a class or an array, and one can validly replace the type `Object` on top of the incoming operand stack with type `int` yielding the outgoing type state.

```
instructionIsTypeSafe(instanceof(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    (CP = class(_, _) ; CP = arrayOf(_)),
    isBootstrapLoader(BL),
    validTypeTransition(Environment, [class('java/lang/Object'), BL], int,
                        StackFrame,NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *invokedynamic*

An *invokedynamic* instruction is type safe iff all of the following conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting an dynamic call site with name `CallSiteName` with descriptor `Descriptor`.

- `CallSiteName` is not `<init>`.

- `CallSiteName` is not `<clinit>`.

- One can validly replace types matching the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`, yielding the outgoing type state.

```
instructionIsTypeSafe(invokedynamic(CP,0,0), Environment, _Offset,
                      StackFrame, NextStackFrame, ExceptionStackFrame) :-
    CP = dmethod(CallSiteName, Descriptor),
    CallSiteName \= '<init>',
```

```
        CallSiteName \= '<clinit>',
        parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
        reverse(OperandArgList, StackArgList),
        validTypeTransition(Environment, StackArgList, ReturnType,
                            StackFrame, NextStackFrame),
        exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *invokeinterface*

An *invokeinterface* instruction is type safe iff all of the following conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting an interface method named MethodName with descriptor Descriptor that is a member of an interface MethodIntfName.

- MethodName is not <init>.

- MethodName is not <clinit>.

- Its second operand, Count, is a valid count operand (see below).

- One can validly replace types matching the type MethodIntfName and the argument types given in Descriptor on the incoming operand stack with the return type given in Descriptor, yielding the outgoing type state.

```
 instructionIsTypeSafe(invokeinterface(CP, Count, 0), Environment, _Offset,
                        StackFrame, NextStackFrame, ExceptionStackFrame) :-
    CP = imethod(MethodIntfName, MethodName, Descriptor),
    MethodName \= '<init>',
    MethodName \= '<clinit>',
    parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
    currentClassLoader(Environment, L),
    reverse([class(MethodIntfName, L) | OperandArgList], StackArgList),
    canPop(StackFrame, StackArgList, TempFrame),
    validTypeTransition(Environment, [], ReturnType, TempFrame, NextStackFrame),
    countIsValid(Count, StackFrame, TempFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The `Count` operand of an *invokeinterface* instruction is valid if it equals the size of the arguments to the instruction. This is equal to the difference between the size of `InputFrame` and `OutputFrame`.

```
countIsValid(Count, InputFrame, OutputFrame) :-
    InputFrame = frame(_Locals1, OperandStack1, _Flags1),
    OutputFrame = frame(_Locals2, OperandStack2, _Flags2),
    length(OperandStack1, Length1),
    length(OperandStack2, Length2),
    Count =:= Length1 - Length2.
```

***invokespecial***

An *invokespecial* instruction is type safe iff all of the following conditions hold:

- Its first operand, `CP`, refers to a constant pool entry denoting a method named `MethodName` with descriptor `Descriptor` that is a member of a class `MethodClassName`.

- Either:

    - `MethodName` is not `<init>`.

    - `MethodName` is not `<clinit>`.

    - One can validly replace types matching the current class and the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`, yielding the outgoing type state.

    - One can validly replace types matching the class `MethodClassName` and the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`.

```
instructionIsTypeSafe(invokespecial(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = method(MethodClassName, MethodName, Descriptor),
    MethodName \= '<init>',
```

```
        MethodName \= '<clinit>',
        parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
        thisClass(Environment, CurrentClass),
        reverse([CurrentClass | OperandArgList], StackArgList),
        validTypeTransition(Environment, StackArgList, ReturnType,
                            StackFrame, NextStackFrame),
        currentClassLoader(Environment, L),
        reverse([class(MethodClassName, L) | OperandArgList], StackArgList2),
        validTypeTransition(Environment, StackArgList2, ReturnType,
                            StackFrame, _ResultStackFrame),
        isAssignable(class(CurrentClassName, L), class(MethodClassName, L)).
        exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

- Or:

  - MethodName is `<init>`.

  - `Descriptor` specifies a `void` return type.

  - One can validly pop types matching the argument types given in `Descriptor` and an
    uninitialized type, `UninitializedArg`, off the incoming operand stack, yielding
    `OperandStack`.

  - The outgoing type state is derived from the incoming type state by first replacing the
    incoming operand stack with `OperandStack` and then replacing all instances of
    `UninitializedArg` with the type of instance being initialized.

```
instructionIsTypeSafe(invokespecial(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = method(MethodClassName, '<init>', Descriptor),
    parseMethodDescriptor(Descriptor, OperandArgList, void),
    reverse(OperandArgList, StackArgList),
    canPop(StackFrame, StackArgList, TempFrame),
    TempFrame = frame(Locals, FullOperandStack, Flags),
    FullOperandStack = [UninitializedArg | OperandStack],
    currentClassLoader(Environment, CurrentLoader),
    rewrittenUninitializedType(UninitializedArg, Environment,
                               class(MethodClassName, CurrentLoader), This),
```

```
        rewrittenInitializationFlags(UninitializedArg, Flags, NextFlags),
        substitute(UninitializedArg, This, OperandStack, NextOperandStack),
        substitute(UninitializedArg, This, Locals, NextLocals),
        NextStackFrame = frame(NextLocals, NextOperandStack, NextFlags),
        ExceptionStackFrame = frame(NextLocals, [], Flags),
        passesProtectedCheck(Environment, MethodClassName, '<init>',
                             Descriptor, NextStackFrame).

rewrittenUninitializedType(uninitializedThis, Environment,
                           _MethodClass, This) :-
    thisClass(Environment, This).

rewrittenUninitializedType(uninitialized(Address), Environment,
                           MethodClass, MethodClass) :-
    allInstructions(Environment, Instructions),
    member(instruction(Address, new(MethodClass)), Instructions).

substitute(_Old, _New, [], []).
substitute(Old, New, [Old | FromRest], [New | ToRest]) :-
    substitute(Old, New, FromRest, ToRest).
substitute(Old, New, [From1 | FromRest], [From1 | ToRest]) :-
    From1 \= Old,
    substitute(Old, New, FromRest, ToRest).
```

To compute what type the uninitialized argument's type needs to be rewritten to, there are two cases:

- If we are initializing an object within its constructor, its type is initially `uninitializedThis`. This type will be rewritten to the type of the class of the `<init>` method.

- The second case arises from initialization of an object created by *new*. The uninitialized arg type is rewritten to `MethodClass`, the type of the method holder of `<init>`. We check whether there really is a *new* instruction at `Address`.

```
rewrittenInitializationFlags(uninitializedThis, _Flags, []).
rewrittenInitializationFlags(uninitialized(_), Flags, Flags).
```

*The rule for invokespecial of an `<init>` method is the sole motivation for passing back a distinct exception stack frame. The concern is that invokespecial can cause a superclass `<init>` method to be invoked, and that invocation could fail, leaving `this` uninitialized. This situation cannot be created using source code in the Java programming language, but can be created by programming in bytecode directly.*

*The original frame holds an uninitialized object in a local and has flag uninitializedThis. Normal termination of invokespecial initializes the uninitialized object and turns off the uninitializedThis flag. But if the invocation of an `<init>` method throws an exception, the uninitialized object might be left in a partially initialized state, and needs to be made permanently unusable. This is represented by an exception frame containing the broken object (the new value of the local) and the uninitializedThis flag (the old flag). There is no way to get from an apparently-initialized object bearing the uninitializedThis flag to a properly initialized object, so the object is permanently unusable. If not for this case, the exception stack frame could be the same as the input stack frame.*

**invokestatic**

An *invokestatic* instruction is type safe iff all of the following conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting a method named `MethodName` with descriptor `Descriptor`.

- `MethodName` is not `<init>`.

- `MethodName` is not `<clinit>`.

- One can validly replace types matching the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`, yielding the outgoing type state.

```
instructionIsTypeSafe(invokestatic(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = method(_MethodClassName, MethodName, Descriptor),
    MethodName \= '<init>',
    MethodName \= '<clinit>',
    parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
    reverse(OperandArgList, StackArgList),
    validTypeTransition(Environment, StackArgList, ReturnType,
                        StackFrame, NextStackFrame),
```

```
        exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

***invokevirtual***

An *invokevirtual* instruction is type safe iff all of the following conditions hold:

- Its first operand, `CP`, refers to a constant pool entry denoting a method named `MethodName` with descriptor `Descriptor` that is a member of a class `MethodClassName`.

- `MethodName` is not `<init>`.

- `MethodName` is not `<clinit>`.

- One can validly replace types matching the class `MethodClassName` and the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`, yielding the outgoing type state.

- If the method is `protected`, the usage conforms to the special rules governing access to `protected` members ([§4.10.1.8](#)).

```
instructionIsTypeSafe(invokevirtual(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = method(MethodClassName, MethodName, Descriptor),
    MethodName \= '<init>',
    MethodName \= '<clinit>',
    parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
    reverse(OperandArgList, ArgList),
    currentClassLoader(Environment, L),
    reverse([class(MethodClassName, L) | OperandArgList], StackArgList),
    validTypeTransition(Environment, StackArgList, ReturnType,
                        StackFrame, NextStackFrame),
    canPop(StackFrame, ArgList, PoppedFrame),
    passesProtectedCheck(Environment, MethodClassName, MethodName,
                         Descriptor, PoppedFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### ior

An *ior* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(ior, iadd).
```

### irem

An *irem* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(irem, iadd).
```

### ireturn

An *ireturn* instruction is type safe if the enclosing method has a declared return type of `int`, and one can validly pop a type matching `int` off the incoming operand stack.

```
instructionIsTypeSafe(ireturn, Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    thisMethodReturnType(Environment, int),
    canPop(StackFrame, [int], _PoppedStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### ishl

An *ishl* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(ishl, iadd).
```

### ishr

An *ishr* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(ishr, iadd).
```

### *istore*

An *istore* instruction with operand Index is type safe and yields an outgoing type state NextStackFrame, if a store instruction with operand Index and type int is type safe and yields an outgoing type state NextStackFrame.

```
instructionIsTypeSafe(istore(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    storeIsTypeSafe(Environment, Index, int, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *istore_<n>*

The instructions *istore_<n>*, for $0 \leq n \leq 3$, are type safe iff the equivalent *istore* instruction is type safe.

```
instructionHasEquivalentTypeRule(istore_0, istore(0)).
instructionHasEquivalentTypeRule(istore_1, istore(1)).
instructionHasEquivalentTypeRule(istore_2, istore(2)).
instructionHasEquivalentTypeRule(istore_3, istore(3)).
```

### *isub*

An *isub* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(isub, iadd).
```

### *iushr*

An *iushr* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(iushr, iadd).
```

**ixor**

An *ixor* instruction is type safe iff the equivalent *iadd* instruction is type safe.

```
instructionHasEquivalentTypeRule(ixor, iadd).
```

**l2d**

An *l2d* instruction is type safe if one can validly pop `long` off the incoming operand stack and replace it with `double`, yielding the outgoing type state.

```
instructionIsTypeSafe(l2d, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long], double,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

**l2f**

An *l2f* instruction is type safe if one can validly pop `long` off the incoming operand stack and replace it with `float`, yielding the outgoing type state.

```
instructionIsTypeSafe(l2f, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long], float,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### l2i

An *l2i* instruction is type safe if one can validly pop `long` off the incoming operand stack and replace it with `int`, yielding the outgoing type state.

```
instructionIsTypeSafe(l2i, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### ladd

An *ladd* instruction is type safe iff one can validly replace types matching `long` and `long` on the incoming operand stack with `long` yielding the outgoing type state.

```
instructionIsTypeSafe(ladd, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long, long], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### laload

An *laload* instruction is type safe iff one can validly replace types matching `int` and array of `long` on the incoming operand stack with `long` yielding the outgoing type state.

```
instructionIsTypeSafe(laload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, arrayOf(long)], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### land

An *land* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(land, ladd).
```

### lastore

An *lastore* instruction is type safe iff one can validly pop types matching `long`, `int` and array of `long` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(lastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [long, int, arrayOf(long)], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### lcmp

A *lcmp* instruction is type safe iff one can validly replace types matching `long` and `long` on the incoming operand stack with `int` yielding the outgoing type state.

```
instructionIsTypeSafe(lcmp, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long, long], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### lconst_<l>

An *lconst_0* instruction is type safe if one can validly push the type `long` onto the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(lconst_0, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [], long, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

An *lconst_1* instruction is type safe iff the equivalent *lconst_0* instruction is type safe.

```
instructionHasEquivalentTypeRule(lconst_1, lconst_0).
```

### *ldc*

An *ldc* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting an entity of type Type, where Type is either int, float, String, Class, java.lang.invoke.MethodType, or java.lang.invoke.MethodHandle, and one can validly push Type onto the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(ldc(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    functor(CP, Tag, _),
    isBootstrapLoader(BL),
    member([Tag, Type], [
        [int, int],
        [float, float],
        [string, class('java/lang/String', BL)],
        [classConst, class('java/lang/Class', BL)],
        [methodTypeConst, class('java/lang/invoke/MethodType', BL)],
        [methodHandleConst, class('java/lang/invoke/MethodHandle', BL)],
    ]),
    validTypeTransition(Environment, [], Type, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *ldc_w*

An *ldc_w* instruction is type safe iff the equivalent *ldc* instruction is type safe.

```
instructionHasEquivalentTypeRule(ldc_w(CP), ldc(CP))
```

### ldc2_w

An *ldc2_w* instruction with operand CP is type safe iff CP refers to a constant pool entry
denoting an entity of type Tag, where Tag is either long or double, and one can validly push
Tag onto the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(ldc2_w(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    functor(CP, Tag, _),
    member(Tag, [long, double]),
    validTypeTransition(Environment, [], Tag, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### ldiv

An *ldiv* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(ldiv, ladd).
```

### lload

An *lload* instruction with operand Index is type safe and yields an outgoing type state
NextStackFrame, if a load instruction with operand Index and type long is type safe and
yields an outgoing type state NextStackFrame.

```
instructionIsTypeSafe(lload(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    loadIsTypeSafe(Environment, Index, long, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### lload_&lt;n&gt;

The instructions *lload_&lt;n&gt;*, for 0 ≤ *n* ≤ 3, are type safe iff the equivalent *lload* instruction is type safe.

```
instructionHasEquivalentTypeRule(lload_0, lload(0)).
instructionHasEquivalentTypeRule(lload_1, lload(1)).
instructionHasEquivalentTypeRule(lload_2, lload(2)).
instructionHasEquivalentTypeRule(lload_3, lload(3)).
```

### lmul

An *lmul* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(lmul, ladd).
```

### lneg

An *lneg* instruction is type safe iff there is a type matching `long` on the incoming operand stack. The *lneg* instruction does not alter the type state.

```
instructionIsTypeSafe(lneg, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### lookupswitch

A *lookupswitch* instruction is type safe if its keys are sorted, one can validly pop `int` off the incoming operand stack yielding a new type state `BranchStackFrame`, and all of the instruction's targets are valid branch targets assuming `BranchStackFrame` as their incoming type state.

```
instructionIsTypeSafe(lookupswitch(Targets, Keys), Environment, _, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    sort(Keys, Keys),
    canPop(StackFrame, [int], BranchStackFrame),
    checklist(targetIsTypeSafe(Environment, BranchStackFrame), Targets),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *lor*

A *lor* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(lor, ladd).
```

### *lrem*

An *lrem* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(lrem, ladd).
```

### *lreturn*

An *lreturn* instruction is type safe if the enclosing method has a declared return type of `long`, and one can validly pop a type matching `long` off the incoming operand stack.

```
instructionIsTypeSafe(lreturn, Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    thisMethodReturnType(Environment, long),
    canPop(StackFrame, [long], _PoppedStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *lshl*

An *lshl* instruction is type safe if one can validly replace the types `int` and `long` on the incoming operand stack with the type `long` yielding the outgoing type state.

```
instructionIsTypeSafe(lshl, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, long], long,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

*lshr*

An *lshr* instruction is type safe iff the equivalent *lshl* instruction is type safe.

```
instructionHasEquivalentTypeRule(lshr, lshl).
```

*lstore*

An *lstore* instruction with operand Index is type safe and yields an outgoing type state `NextStackFrame`, if a store instruction with operand Index and type `long` is type safe and yields an outgoing type state `NextStackFrame`.

```
instructionIsTypeSafe(lstore(Index), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    storeIsTypeSafe(Environment, Index, long, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

*lstore_<n>*

The instructions *lstore_<n>*, for 0 ≤ *n* ≤ 3, are type safe iff the equivalent *lstore* instruction is type safe.

```
instructionHasEquivalentTypeRule(lstore_0, lstore(0)).
```

```
instructionHasEquivalentTypeRule(lstore_1, lstore(1)).
instructionHasEquivalentTypeRule(lstore_2, lstore(2)).
instructionHasEquivalentTypeRule(lstore_3, lstore(3)).
```

### *lsub*

An *lsub* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(lsub, ladd).
```

### *lushr*

An *lushr* instruction is type safe iff the equivalent *lshl* instruction is type safe.

```
instructionHasEquivalentTypeRule(lushr, lshl).
```

### *lxor*

An *lxor* instruction is type safe iff the equivalent *ladd* instruction is type safe.

```
instructionHasEquivalentTypeRule(lxor, ladd).
```

### *monitorenter*

A *monitorenter* instruction is type safe iff one can validly pop a type matching `reference` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(monitorenter, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [reference], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### monitorexit

A *monitorexit* instruction is type safe iff the equivalent *monitorenter* instruction is type safe.

```
instructionHasEquivalentTypeRule(monitorexit, monitorenter).
```

### multianewarray

A *multianewarray* instruction with operands `CP` and `Dim` is type safe iff `CP` refers to a constant pool entry denoting an array type whose dimension is greater or equal to `Dim`, `Dim` is strictly positive, and one can validly replace `Dim` `int` types on the incoming operand stack with the type denoted by `CP` yielding the outgoing type state.

```
instructionIsTypeSafe(multianewarray(CP, Dim), Environment, _Offset,
                      StackFrame, NextStackFrame, ExceptionStackFrame) :-
    CP = arrayOf(_),
    classDimension(CP, Dimension),
    Dimension >= Dim,
    Dim > 0,
    /* Make a list of Dim ints */
    findall(int, between(1, Dim, _), IntList),
    validTypeTransition(Environment, IntList, CP,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The dimension of an array type whose component type is also an array type is one more than the dimension of its component type.

```
classDimension(arrayOf(X), Dimension) :-
    classDimension(X, Dimension1),
    Dimension is Dimension1 + 1.

classDimension(_, Dimension) :-
    Dimension = 0.
```

### *new*

A *new* instruction with operand `CP` at offset `Offset` is type safe iff `CP` refers to a constant pool entry denoting a class type, the type `uninitialized(Offset)` does not appear in the incoming operand stack, and one can validly push `uninitialized(Offset)` onto the incoming operand stack and replace `uninitialized(Offset)` with `top` in the incoming local variables yielding the outgoing type state.

```
instructionIsTypeSafe(new(CP), Environment, Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, OperandStack, Flags),
    CP = class(_, _),
    NewItem = uninitialized(Offset),
    notMember(NewItem, OperandStack),
    substitute(NewItem, top, Locals, NewLocals),
    validTypeTransition(Environment, [], NewItem,
                        frame(NewLocals, OperandStack, Flags),
                        NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *newarray*

A *newarray* instruction with operand `TypeCode` is type safe iff `TypeCode` corresponds to the primitive type `ElementType`, and one can validly replace the type `int` on the incoming operand stack with the type 'array of `ElementType`', yielding the outgoing type state.

```
instructionIsTypeSafe(newarray(TypeCode), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    primitiveArrayInfo(TypeCode, _TypeChar, ElementType, _VerifierType),
    validTypeTransition(Environment, [int], arrayOf(ElementType),
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

The correspondence between type codes and primitive types is specified by the following

predicate:

```
primitiveArrayInfo(4,  0'Z, boolean, int).
primitiveArrayInfo(5,  0'C, char,    int).
primitiveArrayInfo(6,  0'F, float,   float).
primitiveArrayInfo(7,  0'D, double,  double).
primitiveArrayInfo(8,  0'B, byte,    int).
primitiveArrayInfo(9,  0'S, short,   int).
primitiveArrayInfo(10, 0'I, int,     int).
primitiveArrayInfo(11, 0'J, long,    long).
```

***nop***

A *nop* instruction is always type safe. The *nop* instruction does not affect the type state.

```
instructionIsTypeSafe(nop, _Environment, _Offset, StackFrame,
                      StackFrame, ExceptionStackFrame) :-
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

***pop***

A *pop* instruction is type safe iff one can validly pop a category 1 type off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(pop, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, [Type | Rest], Flags),
    Type \= top,
    sizeOf(Type, 1),
    NextStackFrame = frame(Locals, Rest, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

***pop2***

A *pop2* instruction is type safe iff it is a *type safe form* of the *pop2* instruction.

```
instructionIsTypeSafe(pop2, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, InputOperandStack, Flags),
    pop2SomeFormIsTypeSafe(InputOperandStack, OutputOperandStack),
    NextStackFrame = frame(Locals, OutputOperandStack, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A *pop2* instruction is a *type safe form* of the *pop2* instruction iff it is a *type safe form 1 pop2* instruction or a *type safe form 2 pop2* instruction.

```
pop2SomeFormIsTypeSafe(InputOperandStack, OutputOperandStack) :-
    pop2Form1IsTypeSafe(InputOperandStack, OutputOperandStack).

pop2SomeFormIsTypeSafe(InputOperandStack, OutputOperandStack) :-
    pop2Form2IsTypeSafe(InputOperandStack, OutputOperandStack).
```

A *pop2* instruction is a *type safe form 1 pop2* instruction iff one can validly pop two types of size 1 off the incoming operand stack yielding the outgoing type state.

```
pop2Form1IsTypeSafe([Type1, Type2 | Rest], Rest) :-
    sizeOf(Type1, 1),
    sizeOf(Type2, 1).
```

A *pop2* instruction is a *type safe form 2 pop2* instruction iff one can validly pop a type of size 2 off the incoming operand stack yielding the outgoing type state.

```
pop2Form2IsTypeSafe([top, Type | Rest], Rest) :- sizeOf(Type, 2).
```

***putfield***

A *putfield* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting a field whose declared type is `FieldType`, declared in a class `FieldClass`, and one can

Are you a developer? Try out the [HTML to PDF API](#)

validly pop types matching `FieldType` and `FieldClass` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(putfield(CP), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = field(FieldClass, FieldName, FieldDescriptor),
    parseFieldDescriptor(FieldDescriptor, FieldType),
    canPop(StackFrame, [FieldType], PoppedFrame),
    passesProtectedCheck(Environment, FieldClass, FieldName,
                         FieldDescriptor, PoppedFrame),
    currentClassLoader(Environment, CurrentLoader),
    canPop(StackFrame, [FieldType, class(FieldClass, CurrentLoader)],
           NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *putstatic*

A *putstatic* instruction with operand CP is type safe iff CP refers to a constant pool entry denoting a field whose declared type is `FieldType`, and one can validly pop a type matching `FieldType` off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(putstatic(CP), _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    CP = field(_FieldClass, _FieldName, FieldDescriptor),
    parseFieldDescriptor(FieldDescriptor, FieldType),
    canPop(StackFrame, [FieldType], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *return*

A *return* instruction is type safe if the enclosing method declares a `void` return type, and either:

- The enclosing method is not an `<init>` method, or

- this has already been completely initialized at the point where the instruction occurs.

```
instructionIsTypeSafe(return, Environment, _Offset, StackFrame,
                      afterGoto, ExceptionStackFrame) :-
    thisMethodReturnType(Environment, void),
    StackFrame = frame(_Locals, _OperandStack, Flags),
    notMember(flagThisUninit, Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *saload*

An *saload* instruction is type safe iff one can validly replace types matching int and array of short on the incoming operand stack with int yielding the outgoing type state.

```
instructionIsTypeSafe(saload, Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, arrayOf(short)], int,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *sastore*

An *sastore* instruction is type safe iff one can validly pop types matching int, int, and array of short off the incoming operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(sastore, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [int, int, arrayOf(short)], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *sipush*

An *sipush* instruction is type safe iff one can validly push the type int onto the incoming

operand stack yielding the outgoing type state.

```
instructionIsTypeSafe(sipush(_Value), Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [], int, StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *swap*

A *swap* instruction is type safe iff one can validly replace two category 1 types, `Type1` and `Type2`, on the incoming operand stack with the types `Type2` and `Type1` yielding the outgoing type state.

```
instructionIsTypeSafe(swap, _Environment, _Offset, StackFrame,
                      NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(_Locals, [Type1, Type2 | Rest], _Flags),
    sizeOf(Type1, 1),
    sizeOf(Type2, 1),
    NextStackFrame = frame(_Locals, [Type2, Type1 | Rest], _Flags),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

### *tableswitch*

A *tableswitch* instruction is type safe if its keys are sorted, one can validly pop `int` off the incoming operand stack yielding a new type state `BranchStackFrame`, and all of the instruction's targets are valid branch targets assuming `BranchStackFrame` as their incoming type state.

```
instructionIsTypeSafe(tableswitch(Targets, Keys), Environment, _Offset,
                      StackFrame, afterGoto, ExceptionStackFrame) :-
    sort(Keys, Keys),
    canPop(StackFrame, [int], BranchStackFrame),
    checklist(targetIsTypeSafe(Environment, BranchStackFrame), Targets),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

*wide*

The *wide* instructions follow the same rules as the instructions they widen.

```
instructionHasEquivalentTypeRule(wide(WidenedInstruction),
                                 WidenedInstruction).
```

## 4.10.2. Verification by Type Inference

A `class` file that does not contain a `StackMapTable` attribute (which necessarily has a version number of 49.0 or below) must be verified using type inference.

### 4.10.2.1. The Process of Verification by Type Inference

During linking, the verifier checks the code array of the Code attribute for each method of the `class` file by performing data-flow analysis on each method. The verifier ensures that at any given point in the program, no matter what code path is taken to reach that point, the following is true:

- The operand stack is always the same size and contains the same types of values.

- No local variable is accessed unless it is known to contain a value of an appropriate type.

- Methods are invoked with the appropriate arguments.

- Fields are assigned only using values of appropriate types.

- All opcodes have appropriate type arguments on the operand stack and in the local variable array.

- There is never an uninitialized class instance in a local variable in code protected by an exception handler. However, an uninitialized class instance may be on the operand stack in code protected by an exception handler. When an exception is thrown, the contents of the operand stack are discarded.

For efficiency reasons, certain tests that could in principle be performed by the verifier are delayed until the first time the code for the method is actually invoked. In so doing, the verifier avoids loading `class` files unless it has to.

*For example, if a method invokes another method that returns an instance of class A, and that instance is assigned only to a field of the same type, the verifier does not bother to check if the class A actually exists. However, if it is assigned to a field of the type B, the definitions of both A and B must be loaded in to ensure that A is a subclass of B.*

### 4.10.2.2. The Bytecode Verifier

The code for each method is verified independently. First, the bytes that make up the code are broken up into a sequence of instructions, and the index into the code array of the start of each instruction is placed in an array. The verifier then goes through the code a second time and parses the instructions. During this pass a data structure is built to hold information about each Java Virtual Machine instruction in the method. The operands, if any, of each instruction are checked to make sure they are valid. For instance:

- Branches must be within the bounds of the code array for the method.

- The targets of all control-flow instructions are each the start of an instruction. In the case of a *wide* instruction, the *wide* opcode is considered the start of the instruction, and the opcode giving the operation modified by that *wide* instruction is not considered to start an instruction. Branches into the middle of an instruction are disallowed.

- No instruction can access or modify a local variable at an index greater than or equal to the number of local variables that its method indicates it allocates.

- All references to the constant pool must be to an entry of the appropriate type. (For example, the instruction *getfield* must reference a field.)

- The code does not end in the middle of an instruction.

- Execution cannot fall off the end of the code.

- For each exception handler, the starting and ending point of code protected by the handler must be at the beginning of an instruction or, in the case of the ending point, immediately past the end of the code. The starting point must be before the ending point. The exception handler code must start at a valid instruction, and it must not start at an opcode being

modified by the *wide* instruction.

For each instruction of the method, the verifier records the contents of the operand stack and the contents of the local variable array prior to the execution of that instruction. For the operand stack, it needs to know the stack height and the type of each value on it. For each local variable, it needs to know either the type of the contents of that local variable or that the local variable contains an unusable or unknown value (it might be uninitialized). The bytecode verifier does not need to distinguish between the integral types (e.g., byte, short, char) when determining the value types on the operand stack.

Next, a data-flow analyzer is initialized. For the first instruction of the method, the local variables that represent parameters initially contain values of the types indicated by the method's type descriptor; the operand stack is empty. All other local variables contain an illegal value. For the other instructions, which have not been examined yet, no information is available regarding the operand stack or local variables.

Finally, the data-flow analyzer is run. For each instruction, a "changed" bit indicates whether this instruction needs to be looked at. Initially, the "changed" bit is set only for the first instruction. The data-flow analyzer executes the following loop:

1. Select a Java Virtual Machine instruction whose "changed" bit is set. If no instruction remains whose "changed" bit is set, the method has successfully been verified. Otherwise, turn off the "changed" bit of the selected instruction.

2. Model the effect of the instruction on the operand stack and local variable array by doing the following:

   - If the instruction uses values from the operand stack, ensure that there are a sufficient number of values on the stack and that the top values on the stack are of an appropriate type. Otherwise, verification fails.

   - If the instruction uses a local variable, ensure that the specified local variable contains a value of the appropriate type. Otherwise, verification fails.

   - If the instruction pushes values onto the operand stack, ensure that there is sufficient room on the operand stack for the new values. Add the indicated types to the top of the modeled operand stack.

   - If the instruction modifies a local variable, record that the local variable now contains the new type.

3. Determine the instructions that can follow the current instruction. Successor instructions can be one of the following:

- The next instruction, if the current instruction is not an unconditional control transfer instruction (for instance, *goto*, *return*, or *athrow*). Verification fails if it is possible to "fall off" the last instruction of the method.

- The target(s) of a conditional or unconditional branch or switch.

- Any exception handlers for this instruction.

4. Merge the state of the operand stack and local variable array at the end of the execution of the current instruction into each of the successor instructions.

In the special case of control transfer to an exception handler, the operand stack is set to contain a single object of the exception type indicated by the exception handler information. There must be sufficient room on the operand stack for this single value, as if an instruction had pushed it.

- If this is the first time the successor instruction has been visited, record that the operand stack and local variable values calculated in steps 2 and 3 are the state of the operand stack and local variable array prior to executing the successor instruction. Set the "changed" bit for the successor instruction.

- If the successor instruction has been seen before, merge the operand stack and local variable values calculated in steps 2 and 3 into the values already there. Set the "changed" bit if there is any modification to the values.

5. Continue at step 1.

To merge two operand stacks, the number of values on each stack must be identical. The types of values on the stacks must also be identical, except that differently typed `reference` values may appear at corresponding places on the two stacks. In this case, the merged operand stack contains a `reference` to an instance of the first common superclass of the two types. Such a `reference` type always exists because the type `Object` is a superclass of all class and interface types. If the operand stacks cannot be merged, verification of the method fails.

To merge two local variable array states, corresponding pairs of local variables are compared. If the two types are not identical, then unless both contain `reference` values, the verifier records that the local variable contains an unusable value. If both of the pair of local variables

contain `reference` values, the merged state contains a `reference` to an instance of the first common `superclass` of the two types.

If the data-flow analyzer runs on a method without reporting a verification failure, then the method has been successfully verified by the `class` file verifier.

Certain instructions and data types complicate the data-flow analyzer. We now examine each of these in more detail.

### 4.10.2.3. Values of Types long and double

Values of the `long` and `double` types are treated specially by the verification process.

Whenever a value of type `long` or `double` is moved into a local variable at index *n*, index *n*+1 is specially marked to indicate that it has been reserved by the value at index *n* and must not be used as a local variable index. Any value previously at index *n*+1 becomes unusable.

Whenever a value is moved to a local variable at index *n*, the index *n*-1 is examined to see if it is the index of a value of type `long` or `double`. If so, the local variable at index *n*-1 is changed to indicate that it now contains an unusable value. Since the local variable at index *n* has been overwritten, the local variable at index *n*-1 cannot represent a value of type `long` or `double`.

Dealing with values of types `long` or `double` on the operand stack is simpler; the verifier treats them as single values on the stack. For example, the verification code for the *dadd* opcode (add two `double` values) checks that the top two items on the stack are both of type `double`. When calculating operand stack length, values of type `long` and `double` have length two.

Untyped instructions that manipulate the operand stack must treat values of type `long` and `double` as atomic (indivisible). For example, the verifier reports a failure if the top value on the stack is a `double` and it encounters an instruction such as *pop* or *dup*. The instructions *pop2* or *dup2* must be used instead.

### 4.10.2.4. Instance Initialization Methods and Newly Created Objects

Creating a new class instance is a multistep process. The statement:

```
...
new myClass(i, j, k);
...
```

can be implemented by the following:

```
...
new #1             // Allocate uninitialized space for myClass
dup                // Duplicate object on the operand stack
iload_1            // Push i
iload_2            // Push j
iload_3            // Push k
invokespecial #5   // Invoke myClass.<init>
...
```

This instruction sequence leaves the newly created and initialized object on top of the operand stack. (Additional examples of compilation to the instruction set of the Java Virtual Machine are given in [§3](#).)

The instance initialization method ([§2.9](#)) for class `myClass` sees the new uninitialized object as its `this` argument in local variable 0. Before that method invokes another instance initialization method of `myClass` or its direct superclass on `this`, the only operation the method can perform on `this` is assigning fields declared within `myClass`.

When doing dataflow analysis on instance methods, the verifier initializes local variable 0 to contain an object of the current class, or, for instance initialization methods, local variable 0 contains a special type indicating an uninitialized object. After an appropriate instance initialization method is invoked (from the current class or the current superclass) on this object, all occurrences of this special type on the verifier's model of the operand stack and in the local variable array are replaced by the current class type. The verifier rejects code that uses the new object before it has been initialized or that initializes the object more than once. In addition, it ensures that every normal return of the method has invoked an instance initialization method either in the class of this method or in the direct superclass.

Similarly, a special type is created and pushed on the verifier's model of the operand stack as the result of the Java Virtual Machine instruction *new*. The special type indicates the instruction by which the class instance was created and the type of the uninitialized class instance

created. When an instance initialization method declared in the class of the uninitialized class instance is invoked on that class instance, all occurrences of the special type are replaced by the intended type of the class instance. This change in type may propagate to subsequent instructions as the dataflow analysis proceeds.

The instruction number needs to be stored as part of the special type, as there may be multiple not-yet-initialized instances of a class in existence on the operand stack at one time. For example, the Java Virtual Machine instruction sequence that implements:

```
new InputStream(new Foo(), new InputStream("foo"))
```

may have two uninitialized instances of `InputStream` on the operand stack at once. When an instance initialization method is invoked on a class instance, only those occurrences of the special type on the operand stack or in the local variable array that are the same object as the class instance are replaced.

A valid instruction sequence must not have an uninitialized object on the operand stack or in a local variable at the target of a backwards branch if the special type of the uninitialized object is merged with a special type other than itself, or in a local variable in code protected by an exception handler or a `finally` clause. Otherwise, a devious piece of code might fool the verifier into thinking it had initialized a class instance when it had, in fact, initialized a class instance created in a previous pass through a loop.

### 4.10.2.5. Exceptions and `finally`

To implement the `try-finally` construct, a compiler for the Java programming language that generates `class` files with version number 50.0 or below may use the exception-handling facilities together with two special instructions: *jsr* ("jump to subroutine") and *ret* ("return from subroutine"). The `finally` clause is compiled as a subroutine within the Java Virtual Machine code for its method, much like the code for an exception handler. When a *jsr* instruction that invokes the subroutine is executed, it pushes its return address, the address of the instruction after the *jsr* that is being executed, onto the operand stack as a value of type `returnAddress`. The code for the subroutine stores the return address in a local variable. At the end of the subroutine, a *ret* instruction fetches the return address from the local variable and transfers control to the instruction at the return address.

Control can be transferred to the `finally` clause (the `finally` subroutine can be invoked) in

several different ways. If the `try` clause completes normally, the `finally` subroutine is invoked via a *jsr* instruction before evaluating the next expression. A `break` or `continue` inside the `try` clause that transfers control outside the `try` clause executes a *jsr* to the code for the `finally` clause first. If the `try` clause executes a *return*, the compiled code does the following:

1. Saves the return value (if any) in a local variable.

2. Executes a *jsr* to the code for the `finally` clause.

3. Upon return from the `finally` clause, returns the value saved in the local variable.

The compiler sets up a special exception handler, which catches any exception thrown by the `try` clause. If an exception is thrown in the `try` clause, this exception handler does the following:

1. Saves the exception in a local variable.

2. Executes a *jsr* to the `finally` clause.

3. Upon return from the `finally` clause, rethrows the exception.

*For more information about the implementation of the `try-finally` construct, see [§3.13](#).*

The code for the `finally` clause presents a special problem to the verifier. Usually, if a particular instruction can be reached via multiple paths and a particular local variable contains incompatible values through those multiple paths, then the local variable becomes unusable. However, a `finally` clause might be called from several different places, yielding several different circumstances:

- The invocation from the exception handler may have a certain local variable that contains an exception.

- The invocation to implement *return* may have some local variable that contains the return value.

- The invocation from the bottom of the `try` clause may have an indeterminate value in that same local variable.

The code for the `finally` clause itself might pass verification, but after completing the updating all the successors of the *ret* instruction, the verifier would note that the local variable

that the exception handler expects to hold an exception, or that the return code expects to hold a return value, now contains an indeterminate value.

Verifying code that contains a `finally` clause is complicated. The basic idea is the following:

- Each instruction keeps track of the list of *jsr* targets needed to reach that instruction. For most code, this list is empty. For instructions inside code for the `finally` clause, it is of length one. For multiply nested `finally` code (extremely rare!), it may be longer than one.

- For each instruction and each *jsr* needed to reach that instruction, a bit vector is maintained of all local variables accessed or modified since the execution of the *jsr* instruction.

- When executing the *ret* instruction, which implements a return from a subroutine, there must be only one possible subroutine from which the instruction can be returning. Two different subroutines cannot "merge" their execution to a single *ret* instruction.

- To perform the data-flow analysis on a *ret* instruction, a special procedure is used. Since the verifier knows the subroutine from which the instruction must be returning, it can find all the *jsr* instructions that call the subroutine and merge the state of the operand stack and local variable array at the time of the *ret* instruction into the operand stack and local variable array of the instructions following the *jsr*. Merging uses a special set of values for local variables:

  - For any local variable that the bit vector (constructed above) indicates has been accessed or modified by the subroutine, use the type of the local variable at the time of the *ret*.

  - For other local variables, use the type of the local variable before the *jsr* instruction.

## 4.11. Limitations of the Java Virtual Machine

The following limitations of the Java Virtual Machine are implicit in the `class` file format:

- The per-class or per-interface constant pool is limited to 65535 entries by the 16-bit `constant_pool_count` field of the `ClassFile` structure ([§4.1](#)). This acts as an internal limit on the total complexity of a single class or interface.

- The number of fields that may be declared by a class or interface is limited to 65535 by the

size of the `fields_count` item of the `ClassFile` structure ([§4.1](#)).

Note that the value of the `fields_count` item of the `ClassFile` structure does not include fields that are inherited from superclasses or superinterfaces.

- The number of methods that may be declared by a class or interface is limited to 65535 by the size of the `methods_count` item of the `ClassFile` structure ([§4.1](#)).

Note that the value of the `methods_count` item of the `ClassFile` structure does not include methods that are inherited from superclasses or superinterfaces.

- The number of direct superinterfaces of a class or interface is limited to 65535 by the size of the `interfaces_count` item of the `ClassFile` structure ([§4.1](#)).

- The greatest number of local variables in the local variables array of a frame created upon invocation of a method ([§2.6](#)) is limited to 65535 by the size of the `max_locals` item of the Code attribute ([§4.7.3](#)) giving the code of the method, and by the 16-bit local variable indexing of the Java Virtual Machine instruction set.

Note that values of type `long` and `double` are each considered to reserve two local variables and contribute two units toward the `max_locals` value, so use of local variables of those types further reduces this limit.

- The size of an operand stack in a frame ([§2.6](#)) is limited to 65535 values by the `max_stack` field of the Code attribute ([§4.7.3](#)).

Note that values of type `long` and `double` are each considered to contribute two units toward the `max_stack` value, so use of values of these types on the operand stack further reduces this limit.

- The number of method parameters is limited to 255 by the definition of a method descriptor ([§4.3.3](#)), where the limit includes one unit for `this` in the case of instance or interface method invocations.

Note that a method descriptor is defined in terms of a notion of method parameter length in which a parameter of type `long` or `double` contributes two units to the length, so parameters of these types further reduce the limit.

- The length of field and method names, field and method descriptors, and other constant string values (including those referenced by `ConstantValue` ([§4.7.2](#)) attributes) is limited

to 65535 characters by the 16-bit unsigned `length` item of the `CONSTANT_Utf8_info` structure ([§4.4.7](#)).

Note that the limit is on the number of bytes in the encoding and not on the number of encoded characters. UTF-8 encodes some characters using two or three bytes. Thus, strings incorporating multibyte characters are further constrained.

- The number of dimensions in an array is limited to 255 by the size of the *dimensions* opcode of the *multianewarray* instruction and by the constraints imposed on the *multianewarray*, *anewarray*, and *newarray* instructions ([§4.9.1](#), [§4.9.2](#)).

[Legal Notice](#)